
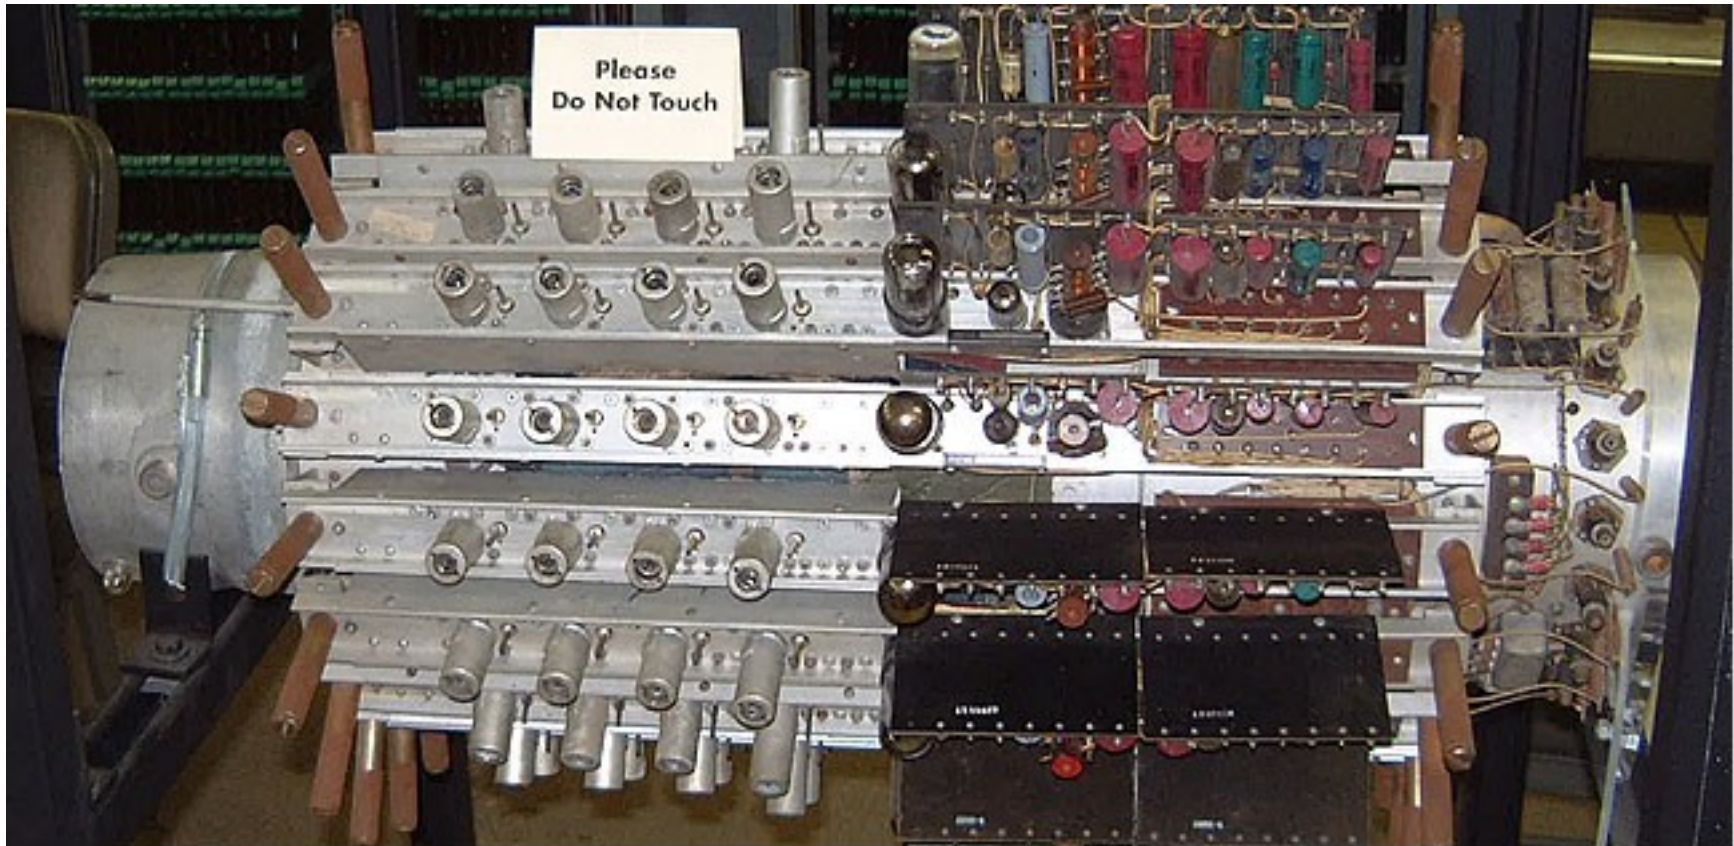
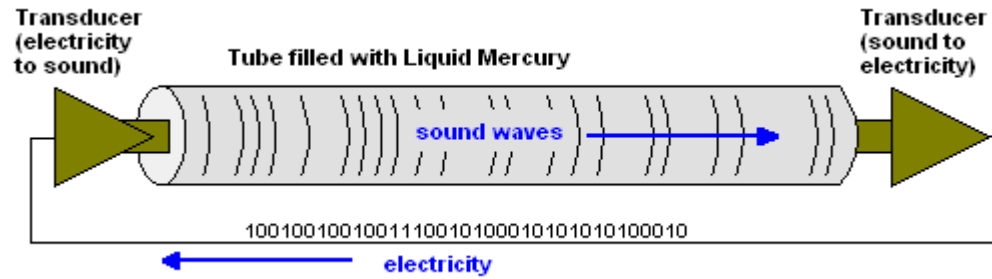


# Logic Design: Implementing Memory

Memory: an organism's ability to **store, retain and recall** information

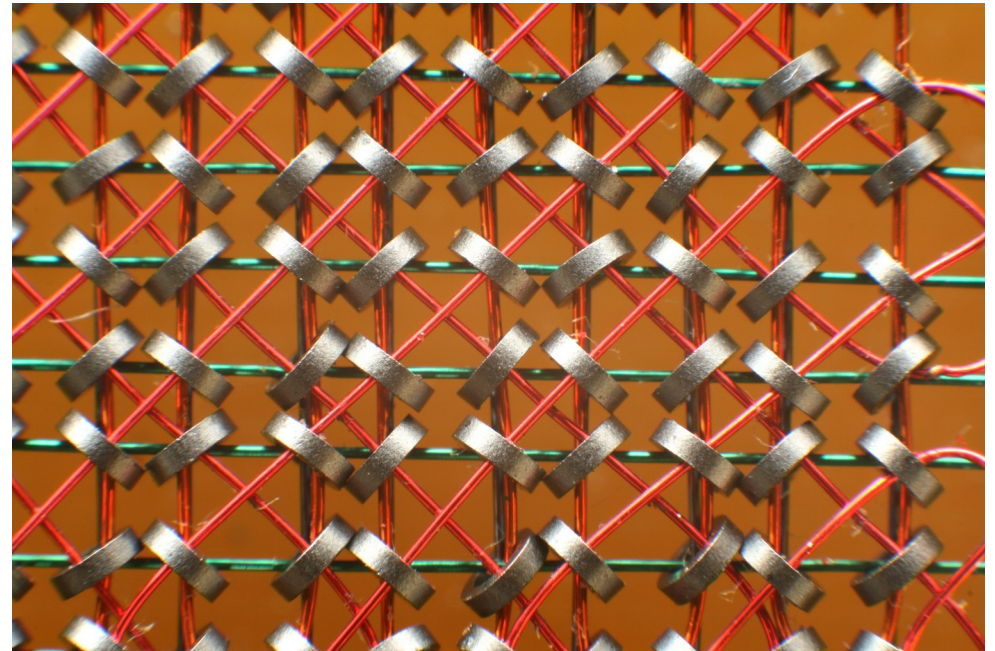
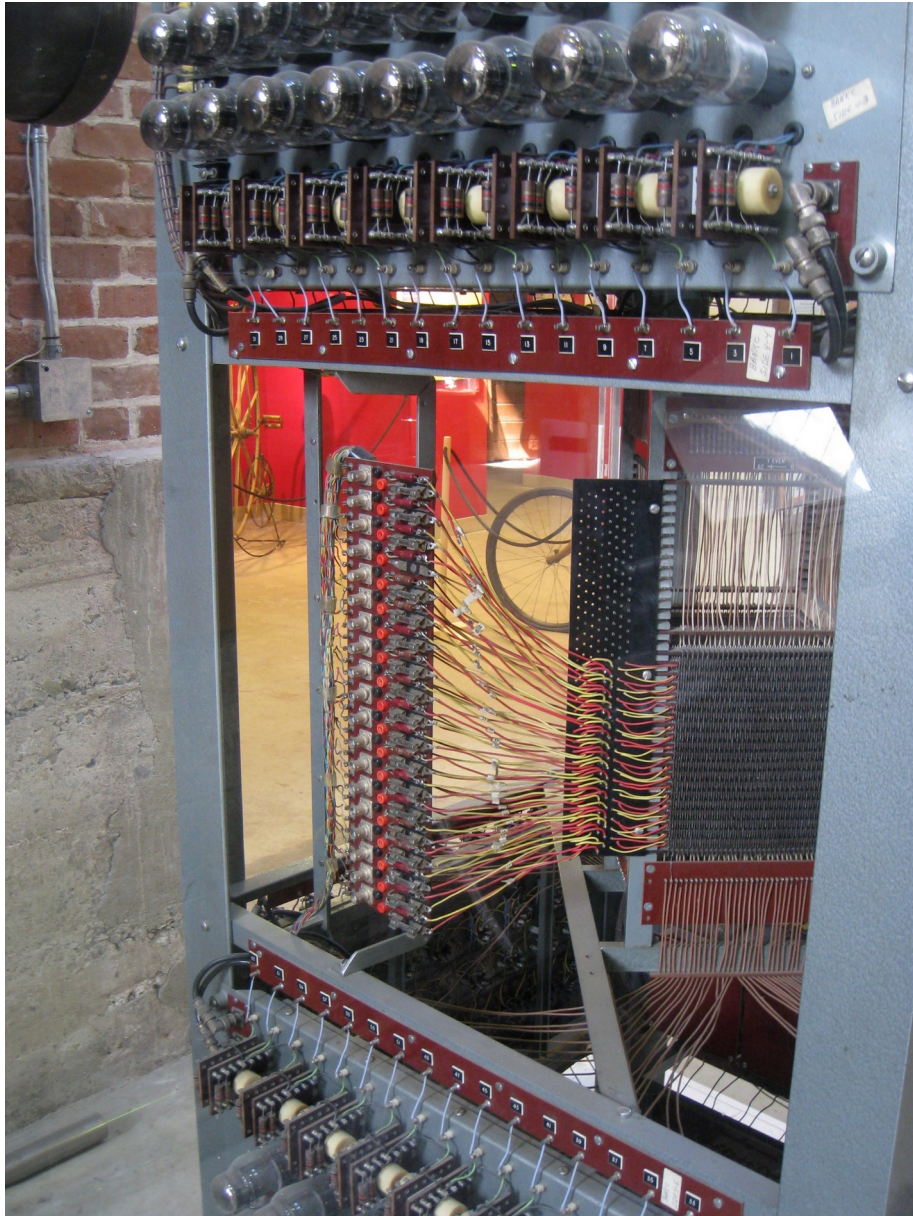


# Delay Line memory





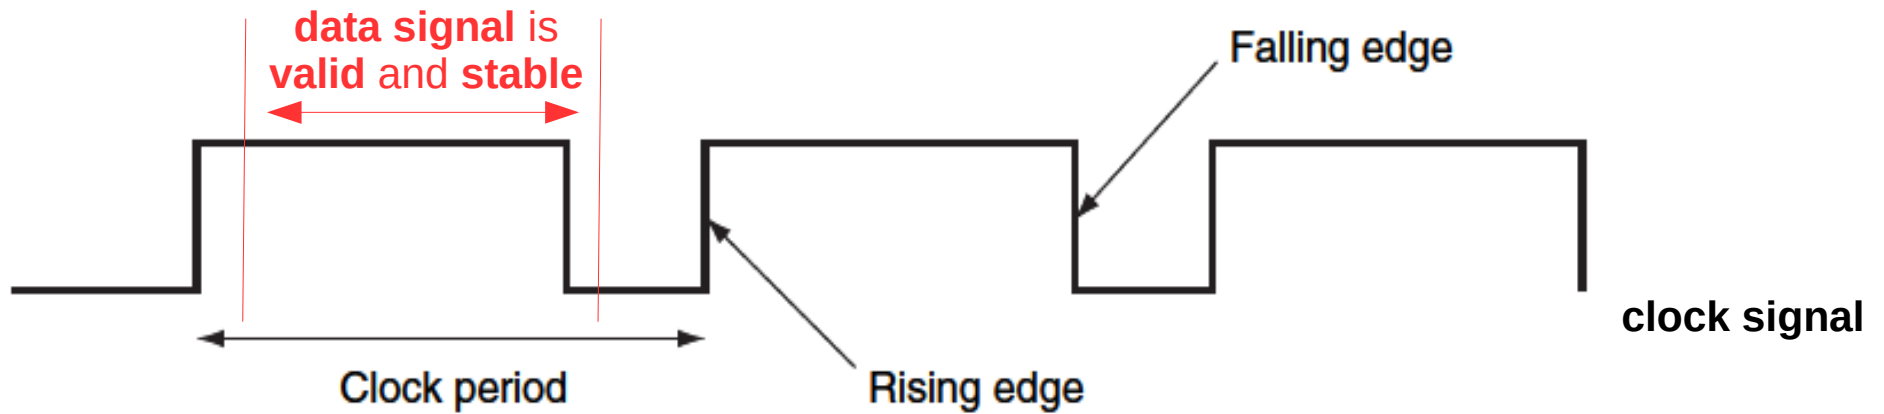
# Magnetic Core memory



# Modern memory bank



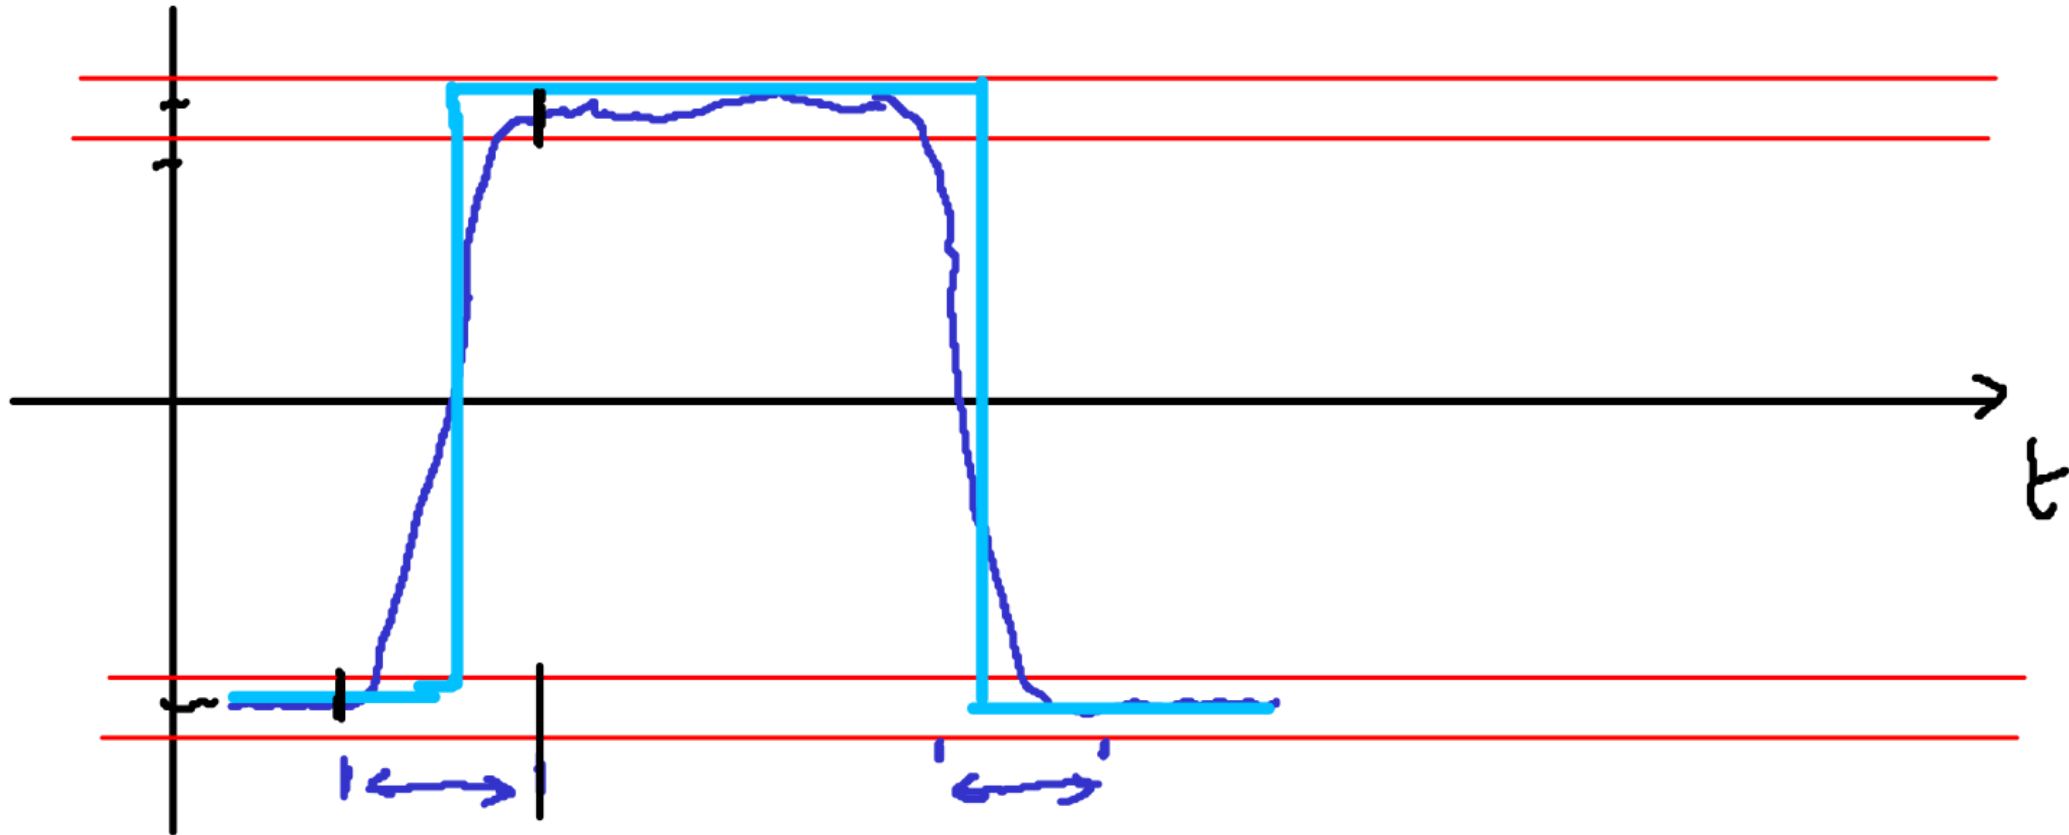
# Clock Signal (memory retains **data** over time)



Clocking **methodology**:  
determines **when** data is **valid** and **stable** relative to the **clock**

**Edge-triggered** clocking:  
all **state changes** occur on a **clock edge**

# Digital Signals (analog reality)



# State Elements: “memory”

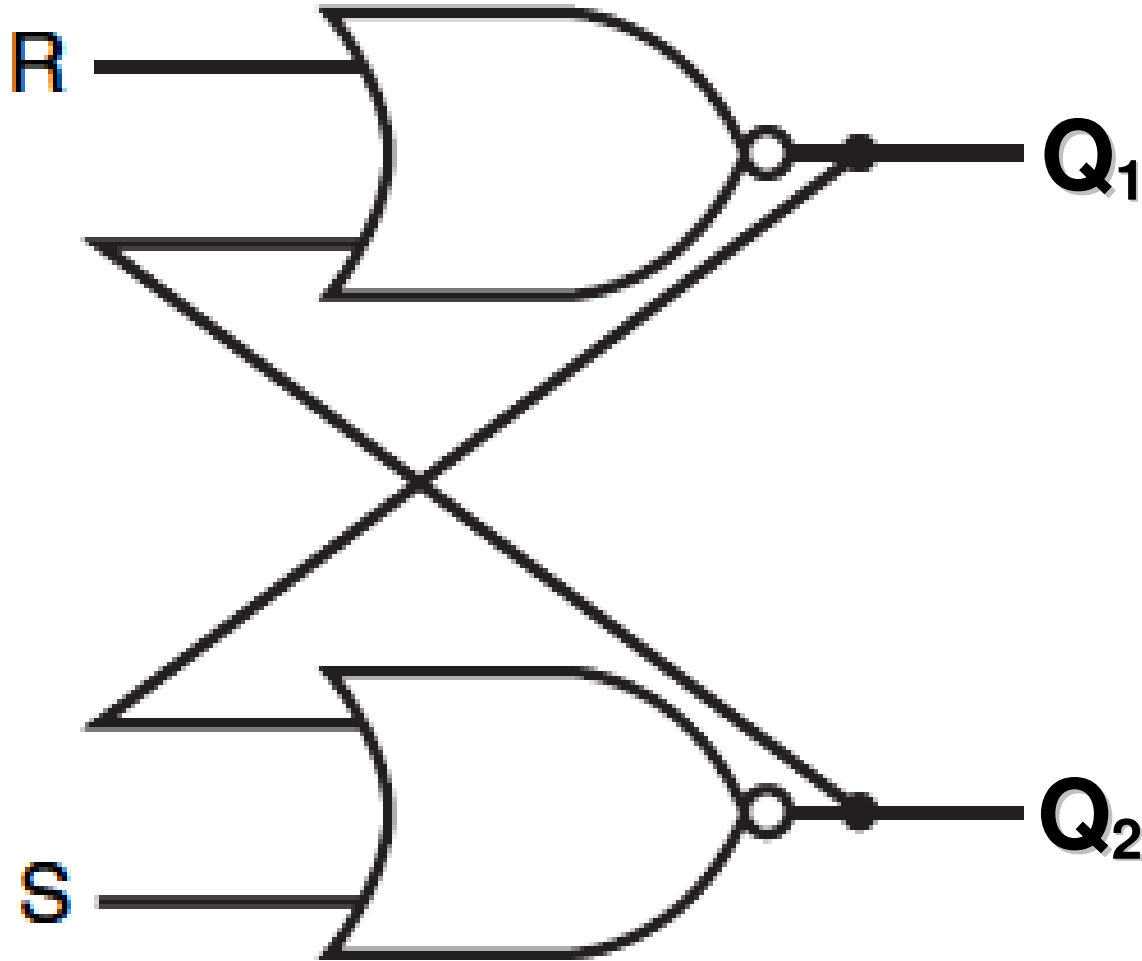
State Element → Combinational Logic → State Element  
(function/computation)



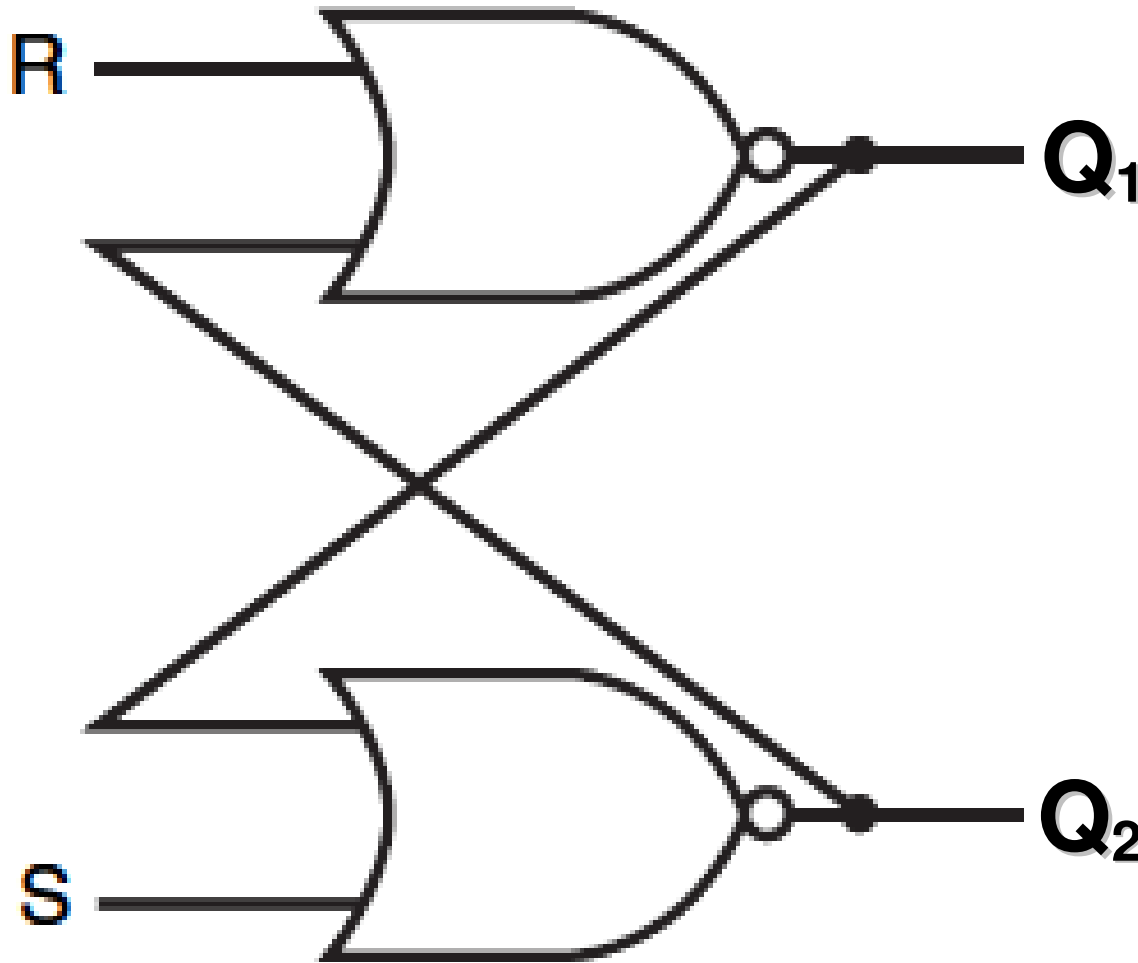
**Latch: store value**



# S-R Latch (unclocked): store and remember value



# S-R Latch (unclocked): store and remember value



$$Q_1 = \text{not } (R \text{ or } Q_2)$$

$$Q_2 = \text{not } (S \text{ or } Q_1)$$

→

$$Q_1 = \text{not } (R \text{ or } \text{not } (S \text{ or } Q_1))$$

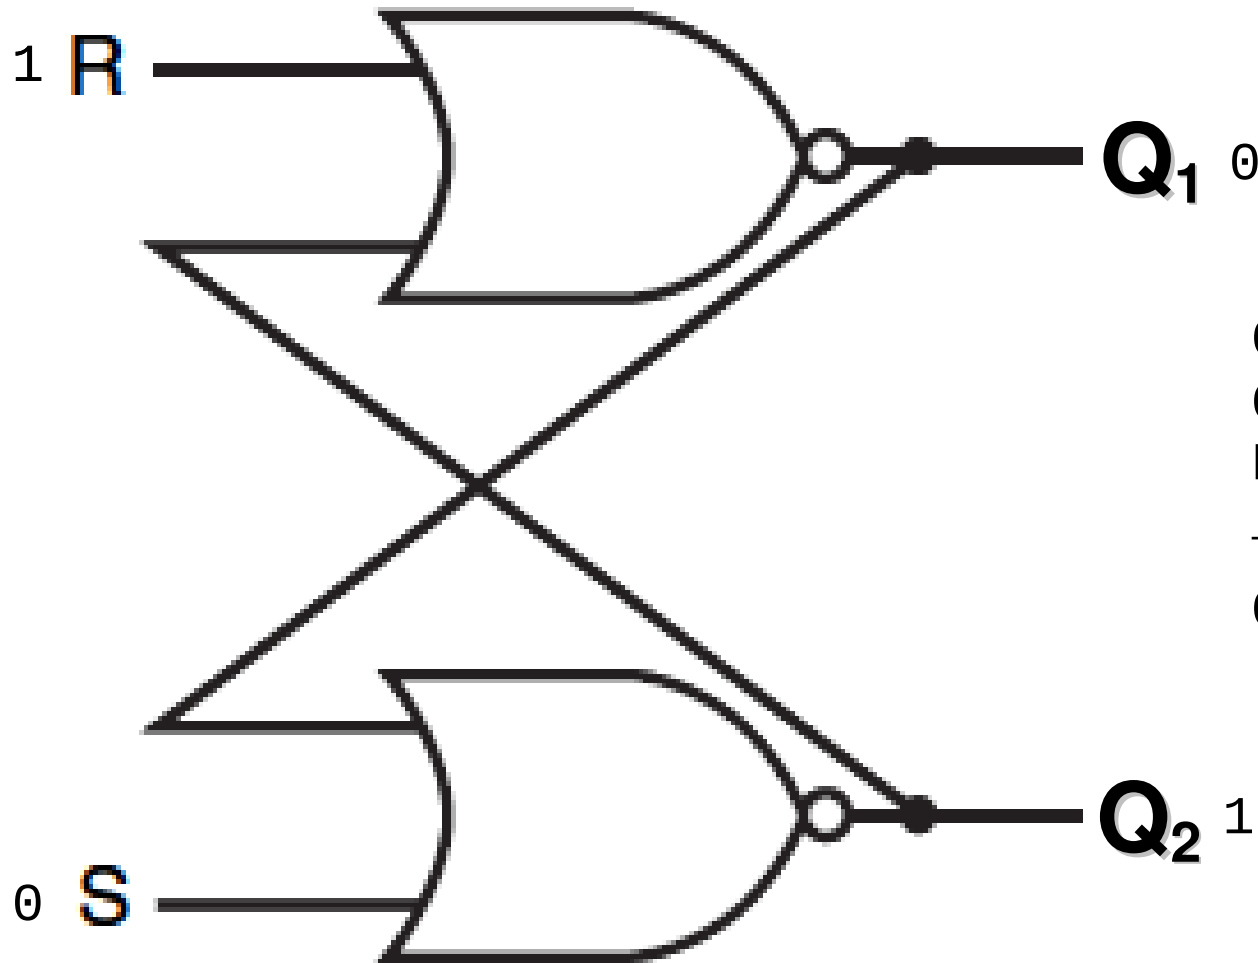
$$Q_2 = \text{not } (S \text{ or } Q_1)$$

→

$$Q_1 = \text{not } R \text{ and } (S \text{ or } Q_1)$$

$$Q_2 = \text{not } (S \text{ or } Q_1)$$

# S-R Latch set



$$Q_1 = \text{not } R \text{ and } (S \text{ or } Q_1)$$

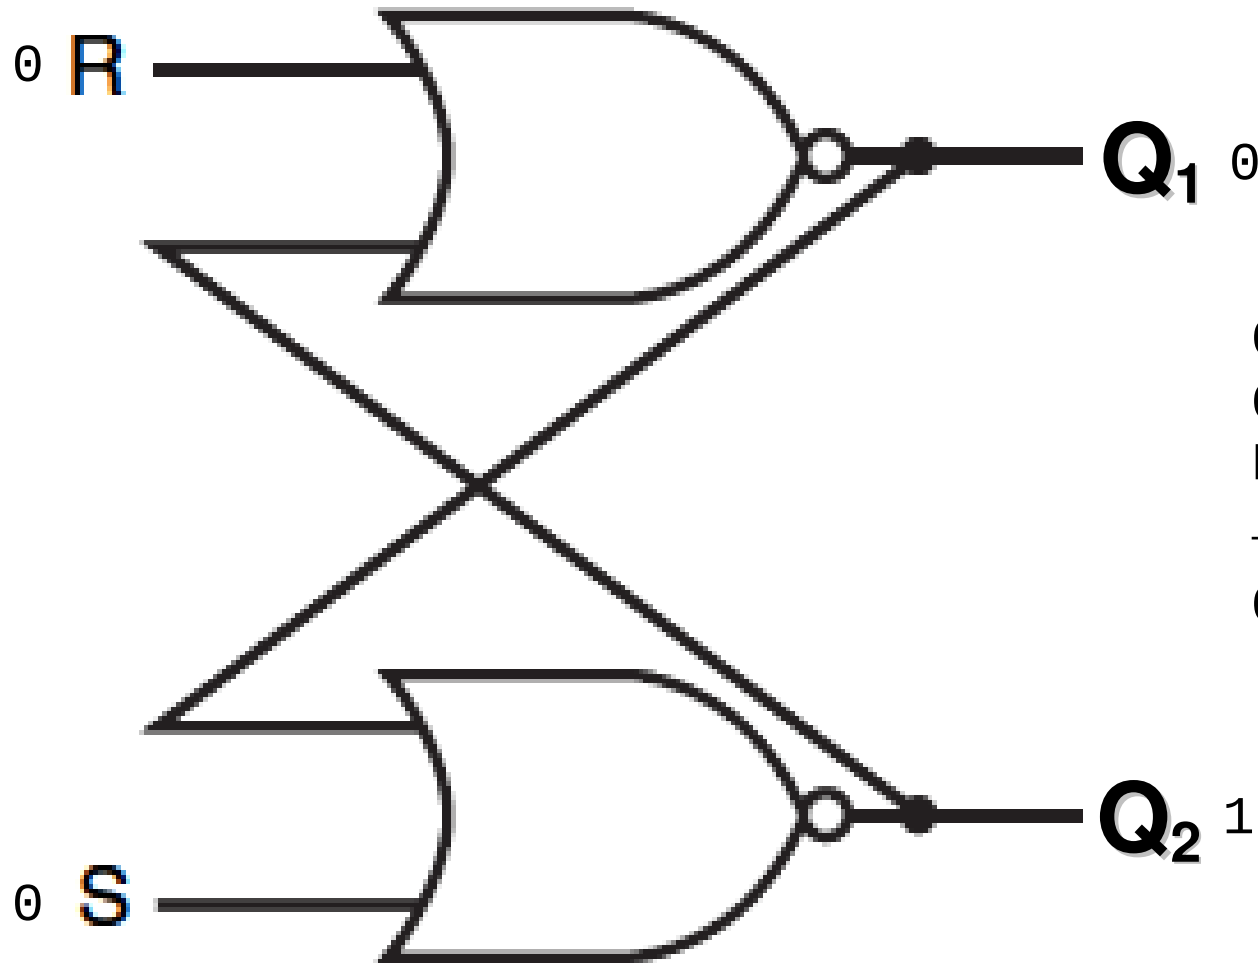
$$Q_2 = \text{not } (S \text{ or } Q_1)$$

$$R = 1, S = 0$$

→

$$Q_1 = 0, Q_2 = \text{not } Q_1 = 1$$

# S-R Latch remember output value



$$Q_1 = \text{not } R \text{ and } (S \text{ or } Q_1)$$

$$Q_2 = \text{not } (S \text{ or } Q_1)$$


$$R = 0, S = 0$$

→

$$Q_1 = Q_1, Q_2 = \text{not } Q_1$$



# S-R Latch set



$$Q_1 = \text{not } R \text{ and } (S \text{ or } Q_1)$$

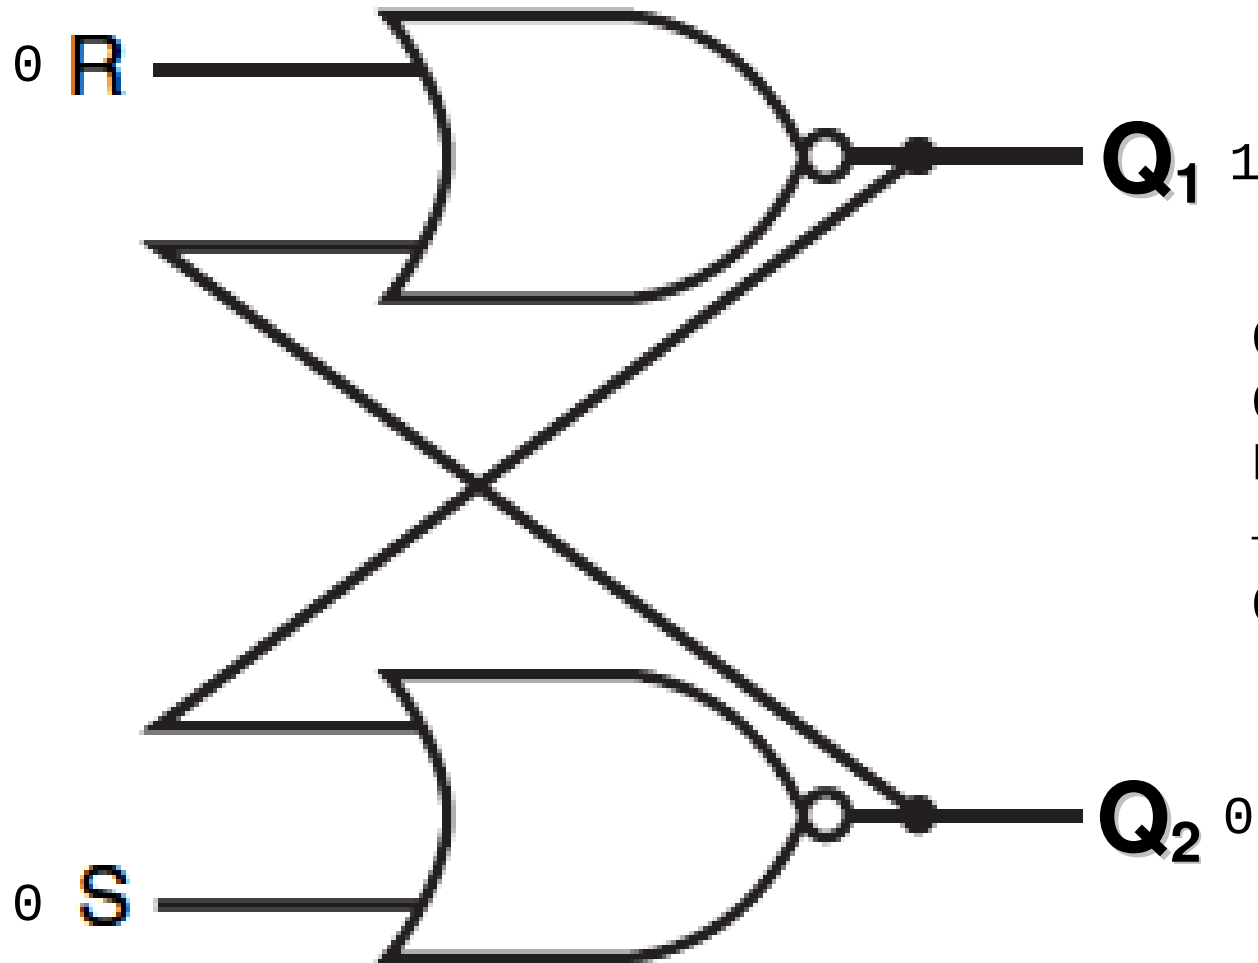
$$Q_2 = \text{not } (S \text{ or } Q_1)$$

$$R = 0, S = 1$$

→

$$Q_1 = 1, Q_2 = 0 = \text{not } Q_1$$

# S-R Latch remember output value



$$Q_1 = \text{not } R \text{ and } (S \text{ or } Q_1)$$

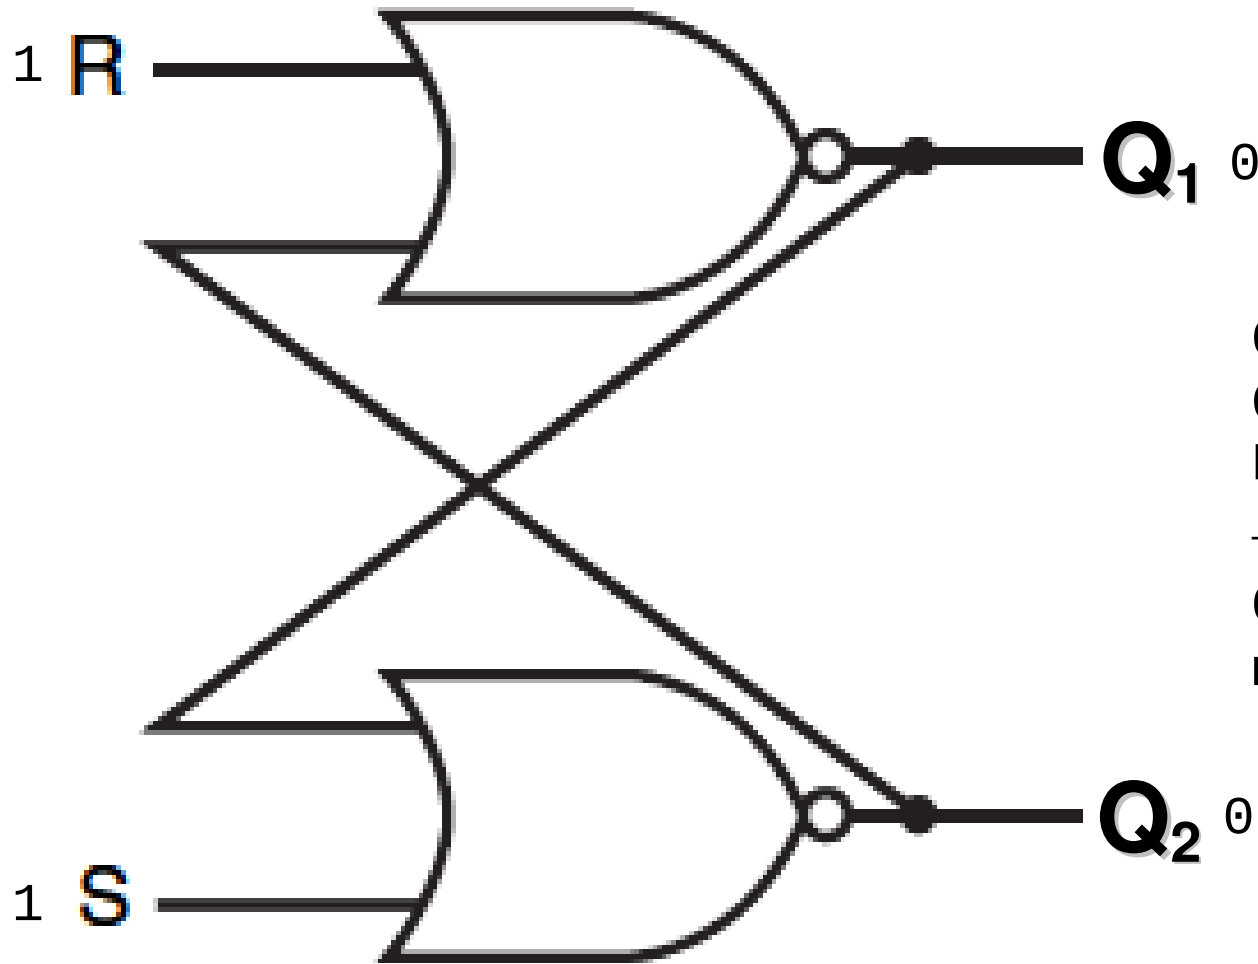
$$Q_2 = \text{not } (S \text{ or } Q_1)$$

$$R = 0, S = 0$$

→

$$Q_1 = Q_1, Q_2 = \text{not } Q_1$$

# S-R Latch — no longer $Q_2 = \text{not } Q_1$



$$Q_1 = \text{not } R \text{ and } (S \text{ or } Q_1)$$

$$Q_2 = \text{not } (S \text{ or } Q_1)$$

$$R = 1, S = 1$$

→

$$Q_1 = 0, Q_2 = 0$$

may become meta-stable  
(**analog** reality: delays)

Latch:

change **while** clock asserted

(clock connected to R or S)

Flip-Flop:


change **on** clock **edge**



# D latch ("transparent")

C: "Clock"

D: "Data"



$R = C \text{ and not } D$

$S = C \text{ and } D$

→

when  $C = 1$ : "open"

$R = \text{not } D$

$S = D$

→  $Q = S = D$

when  $C = 0$ : "closed"

$R = 0$


$S = 0$

→  $Q$  is not changed

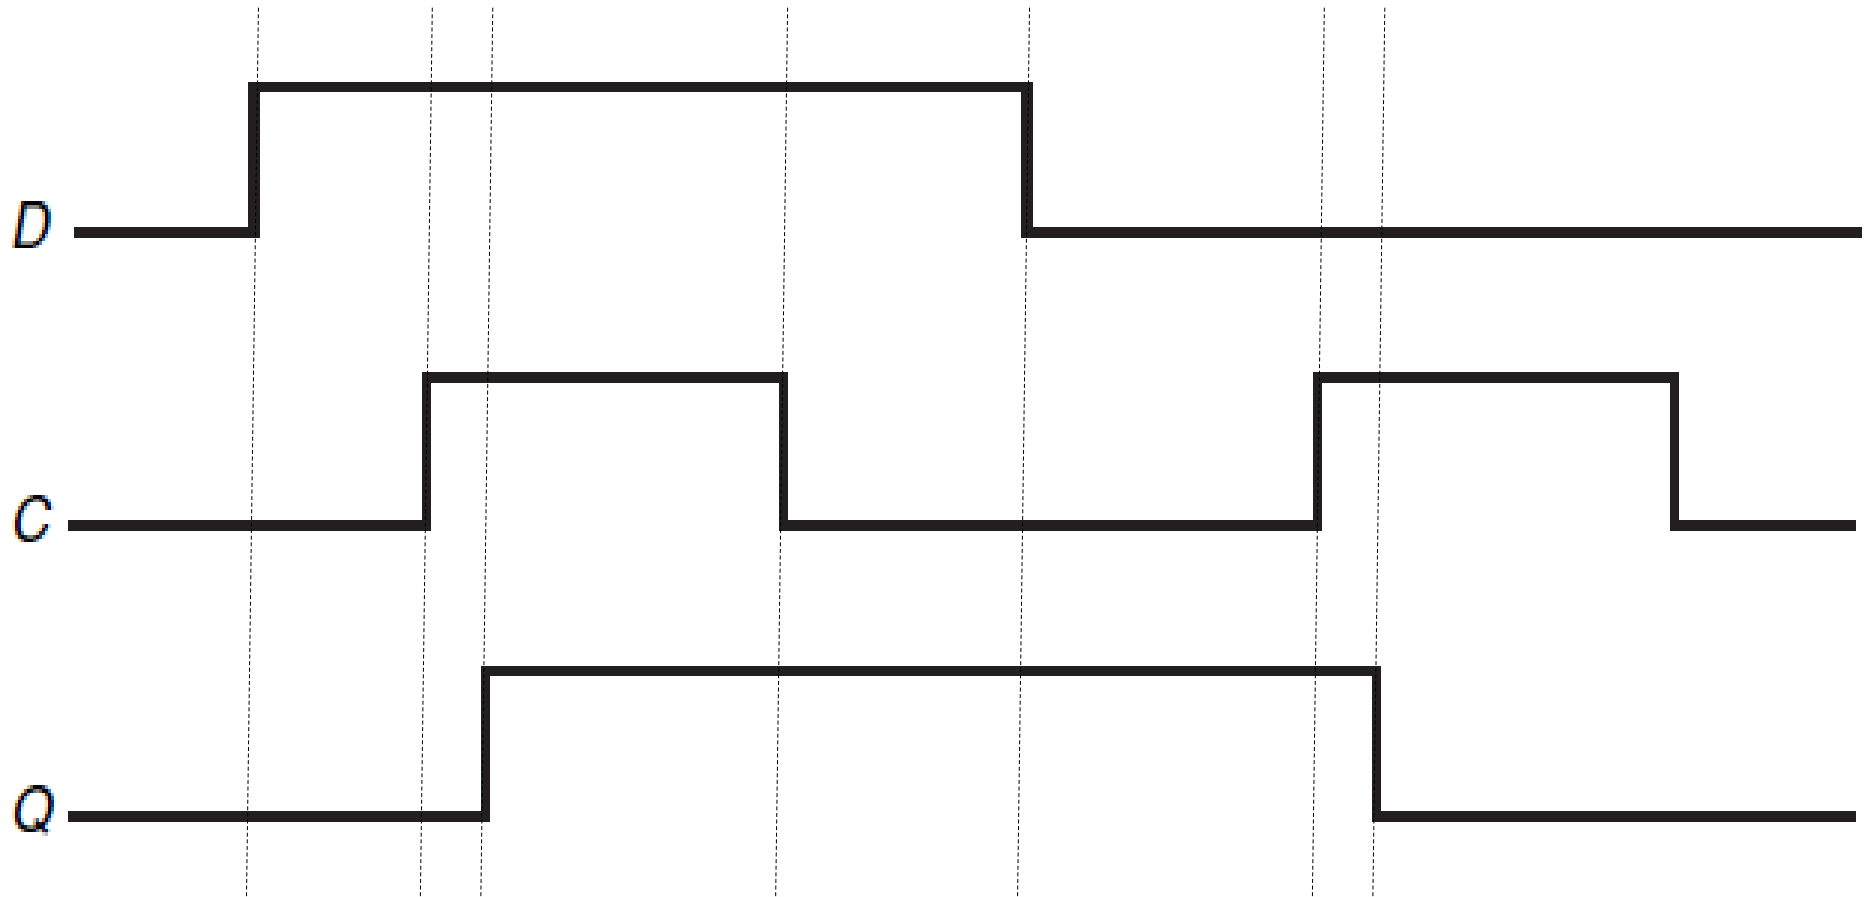
$R = S = 1$

not possible

# D latch operation

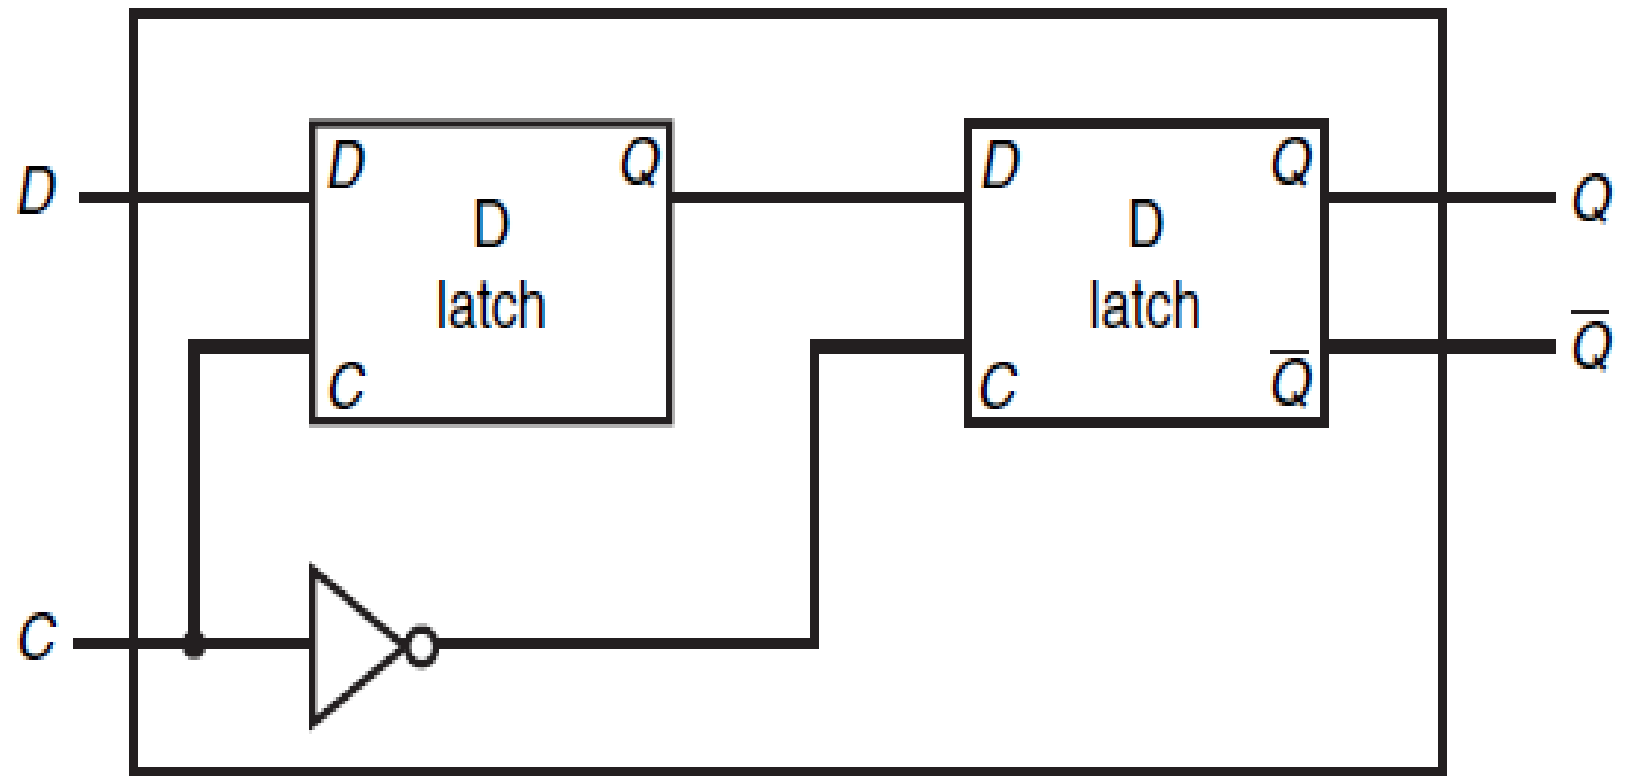


# D latch operation



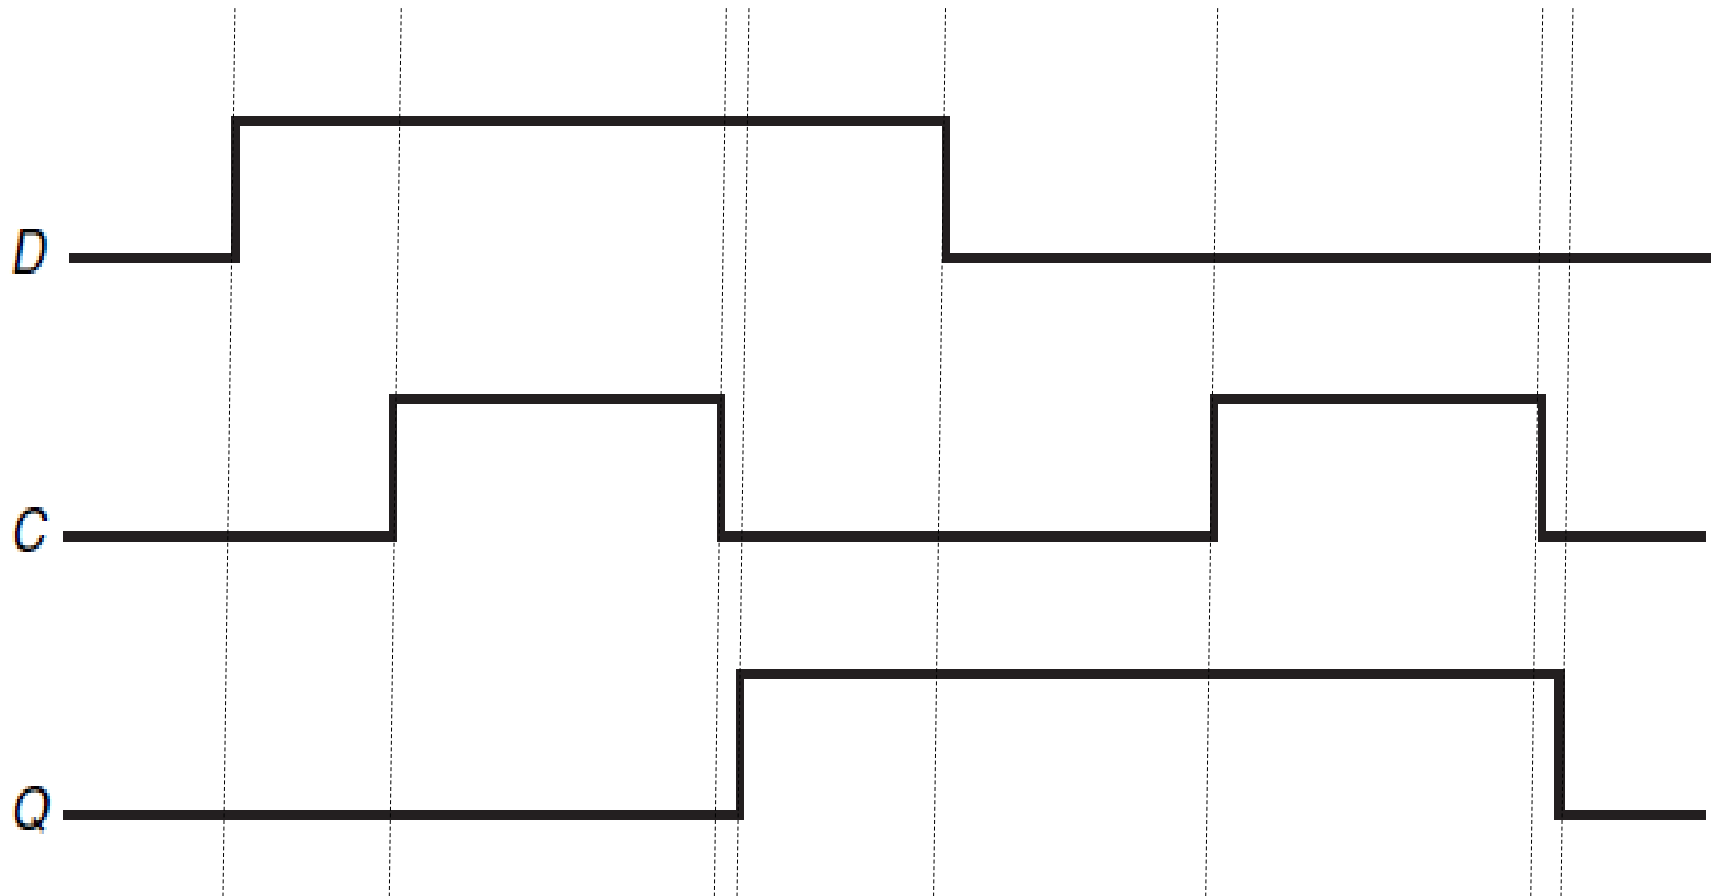
What if  $D$  changes during  $C$  asserted?

# D flip-flop (not “transparent”) with falling-edge trigger

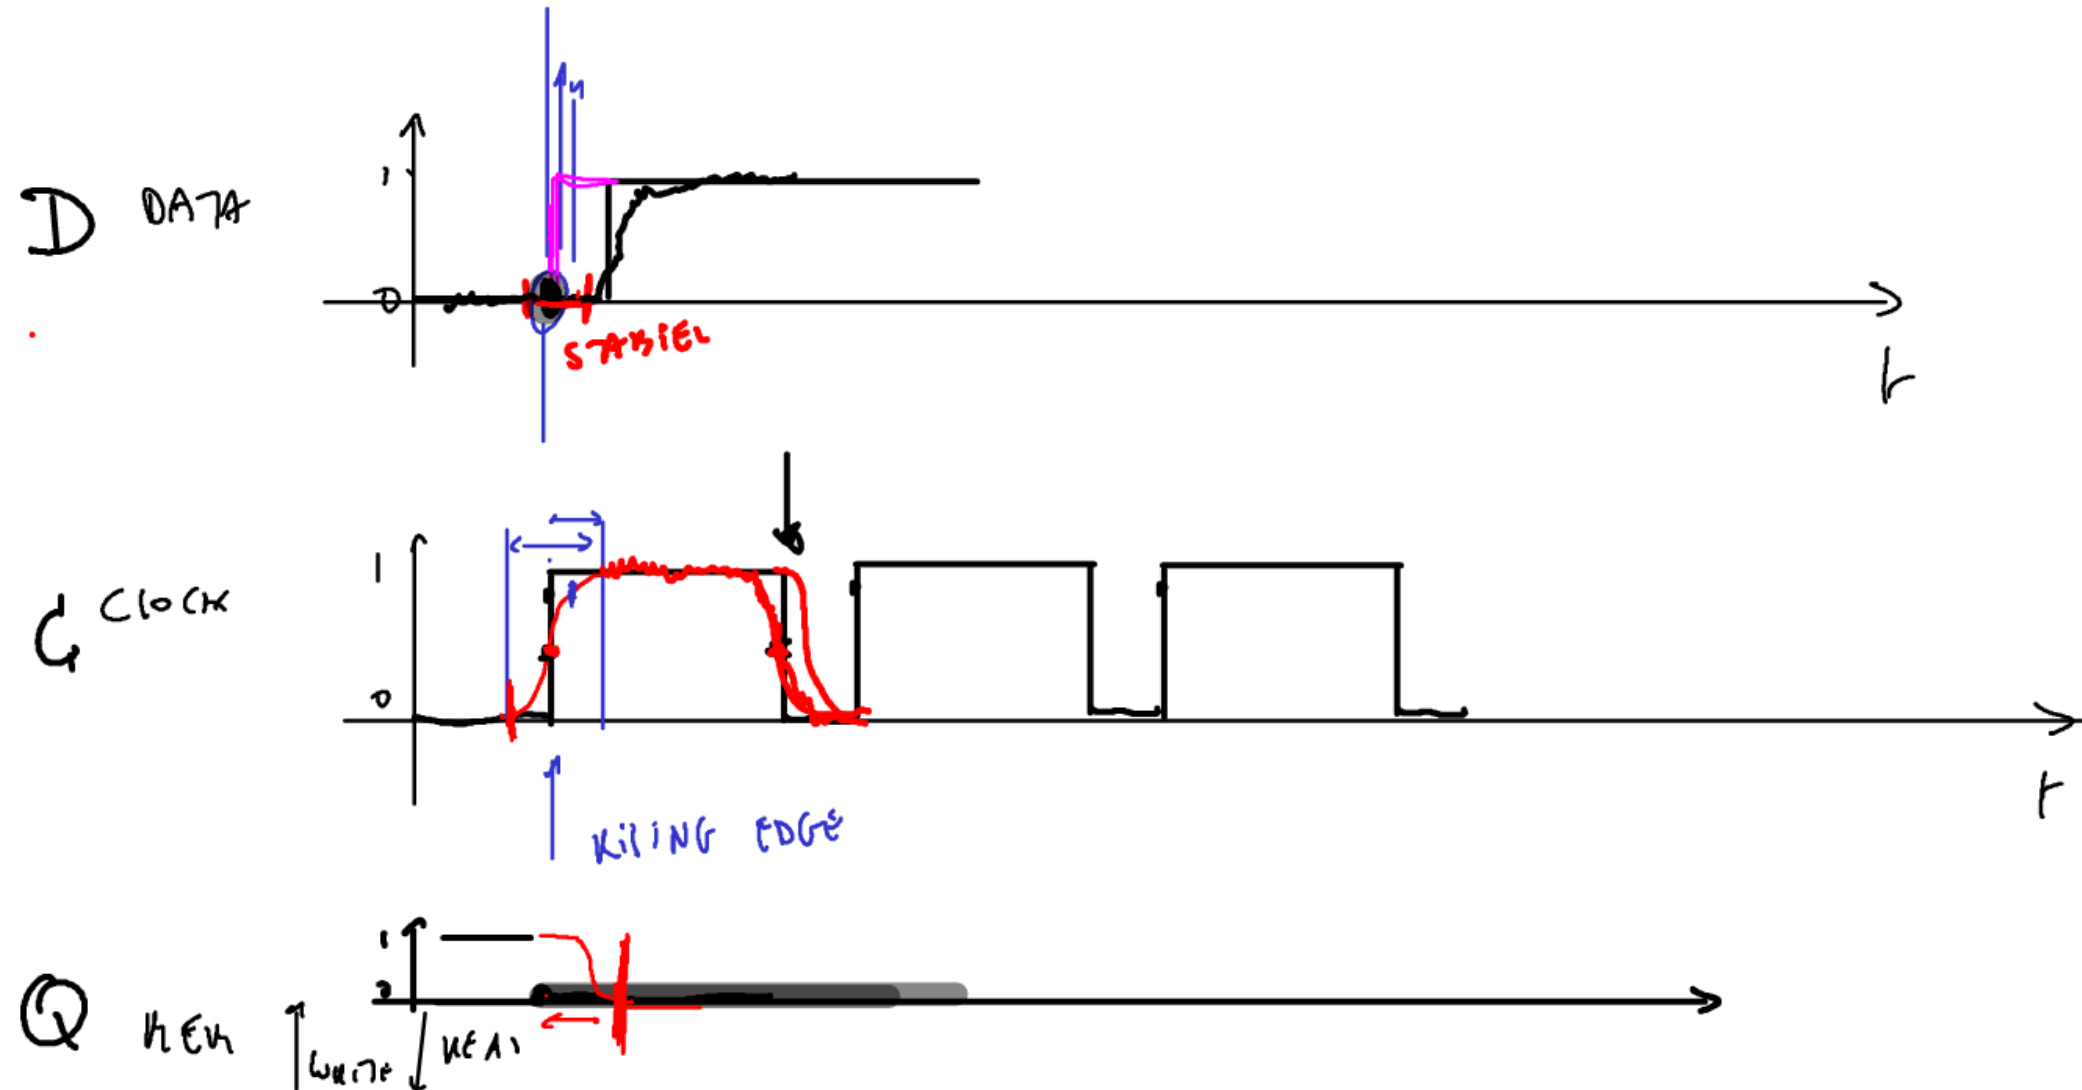




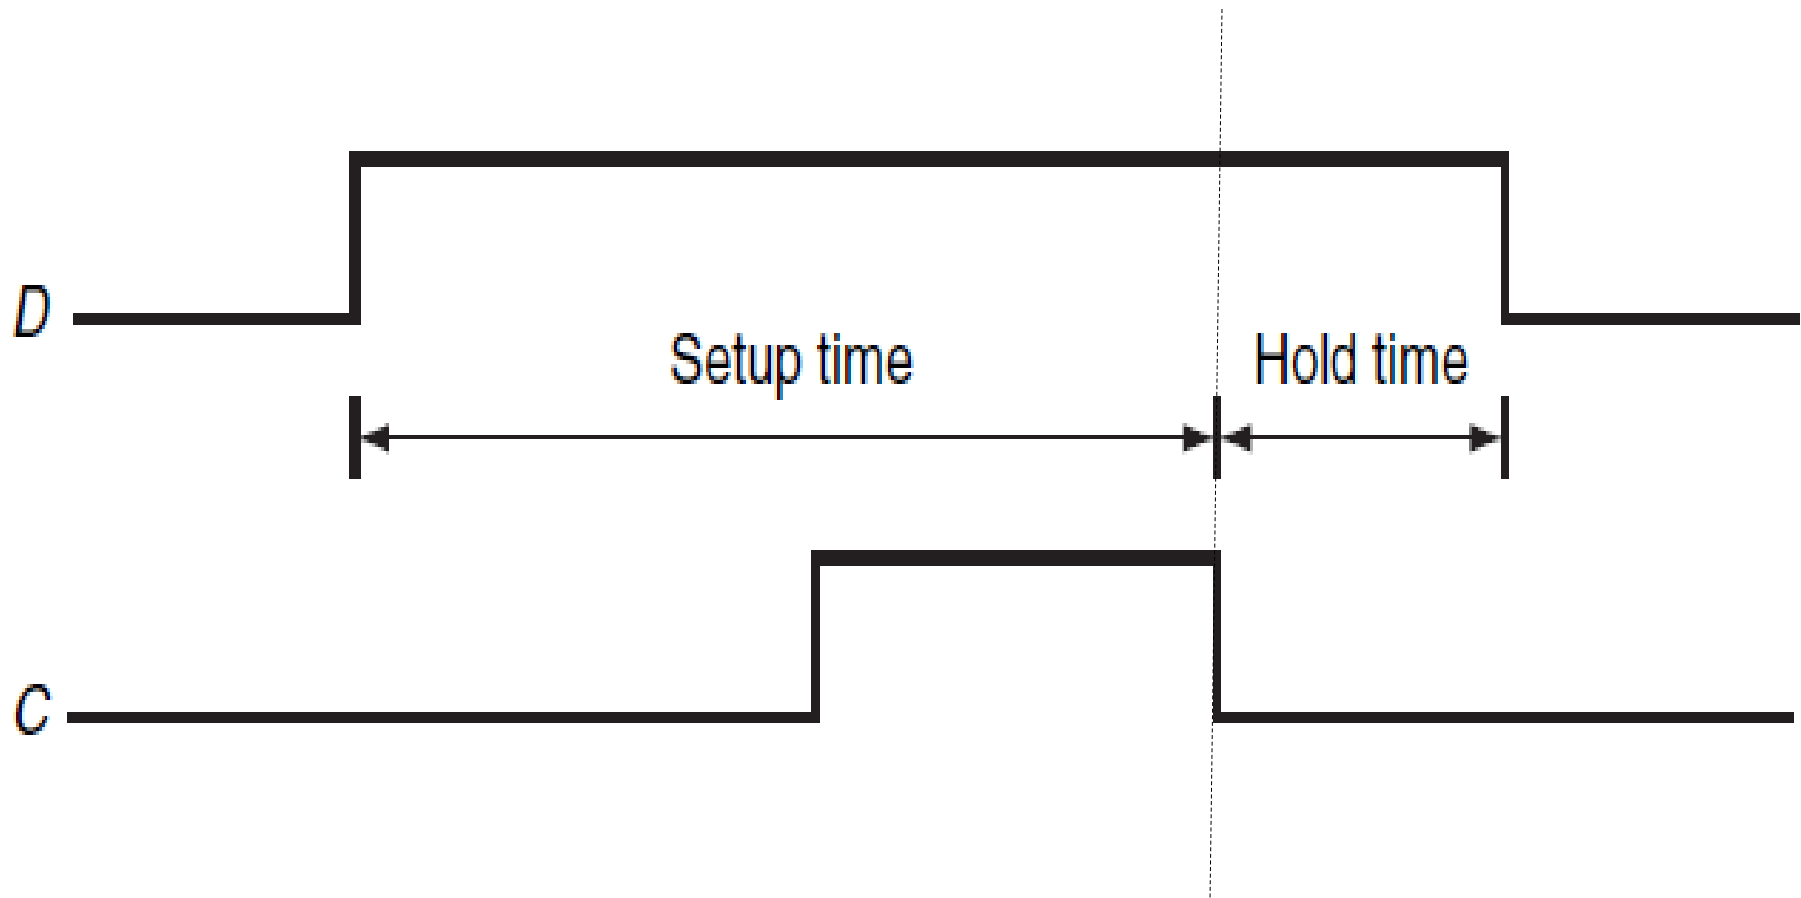
# D flip-flop with falling-edge trigger operation



# timing constraints: “leaky” abstraction (rising edge)



# D flip-flop (falling edge) timing constraints

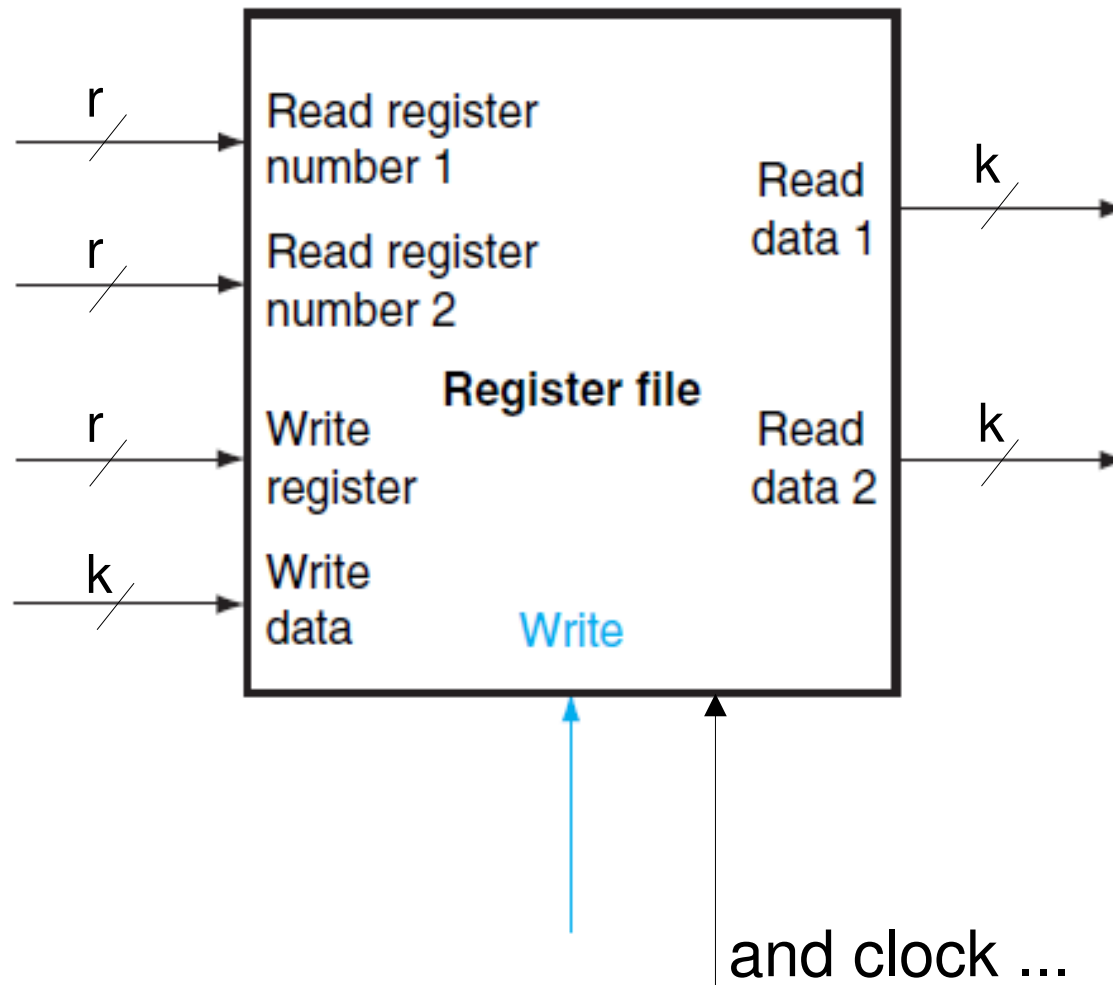


# Register File

k-bit registers  
(using k D flip-flops, see later: 4x2 SRAM)

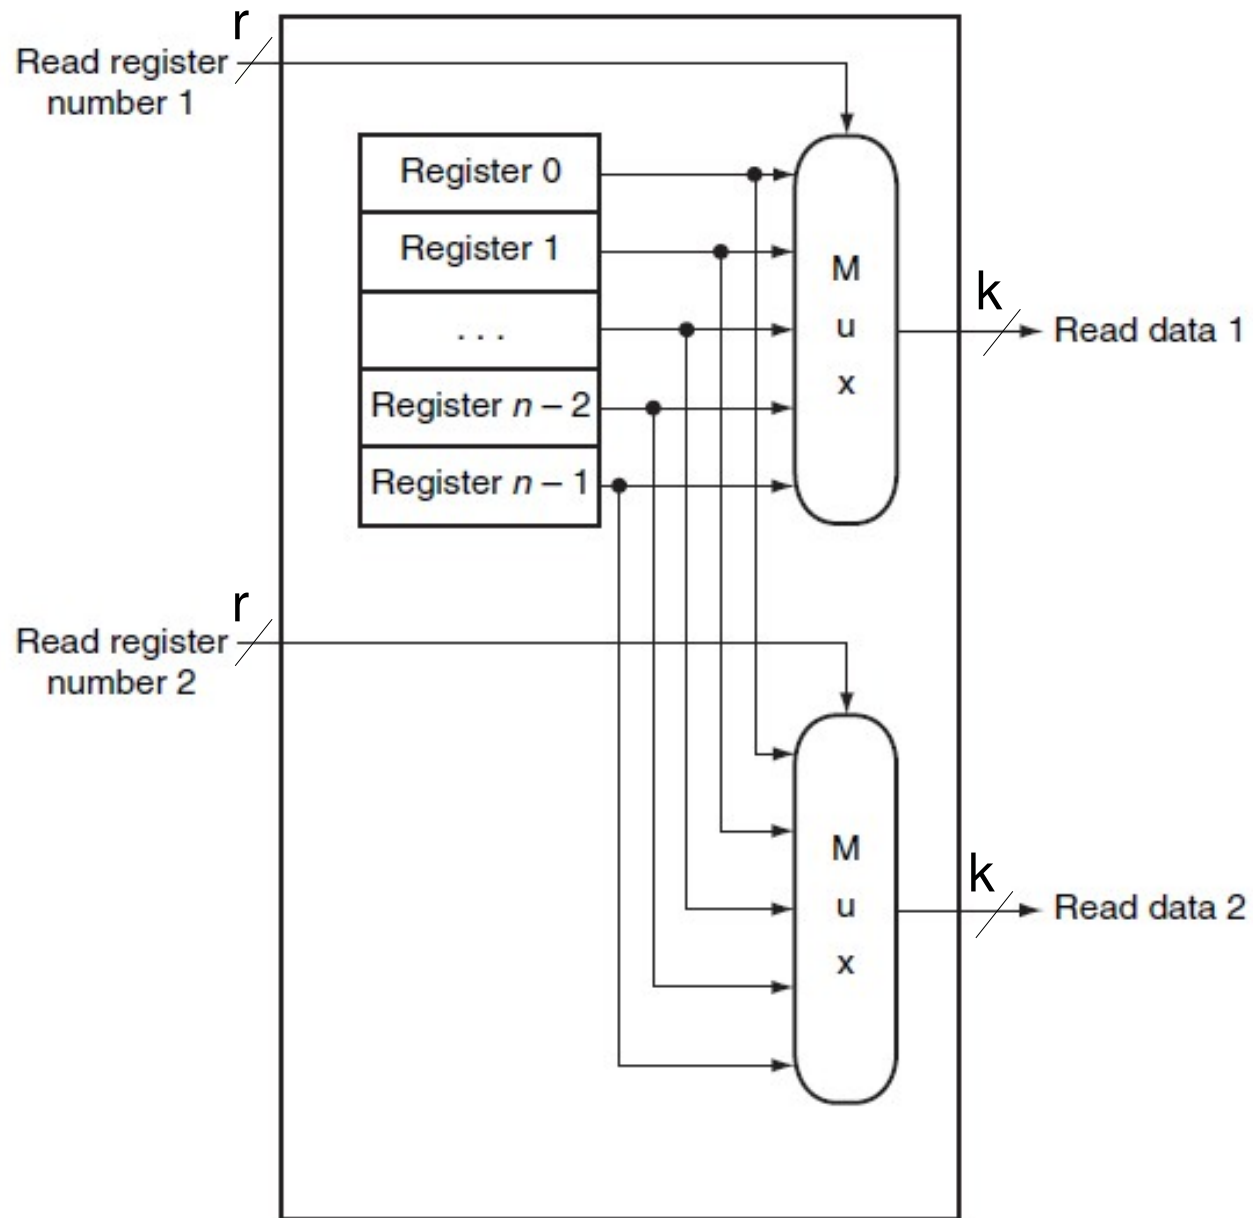
number of registers n

$$\rightarrow r = \lceil \log_2 n \rceil$$

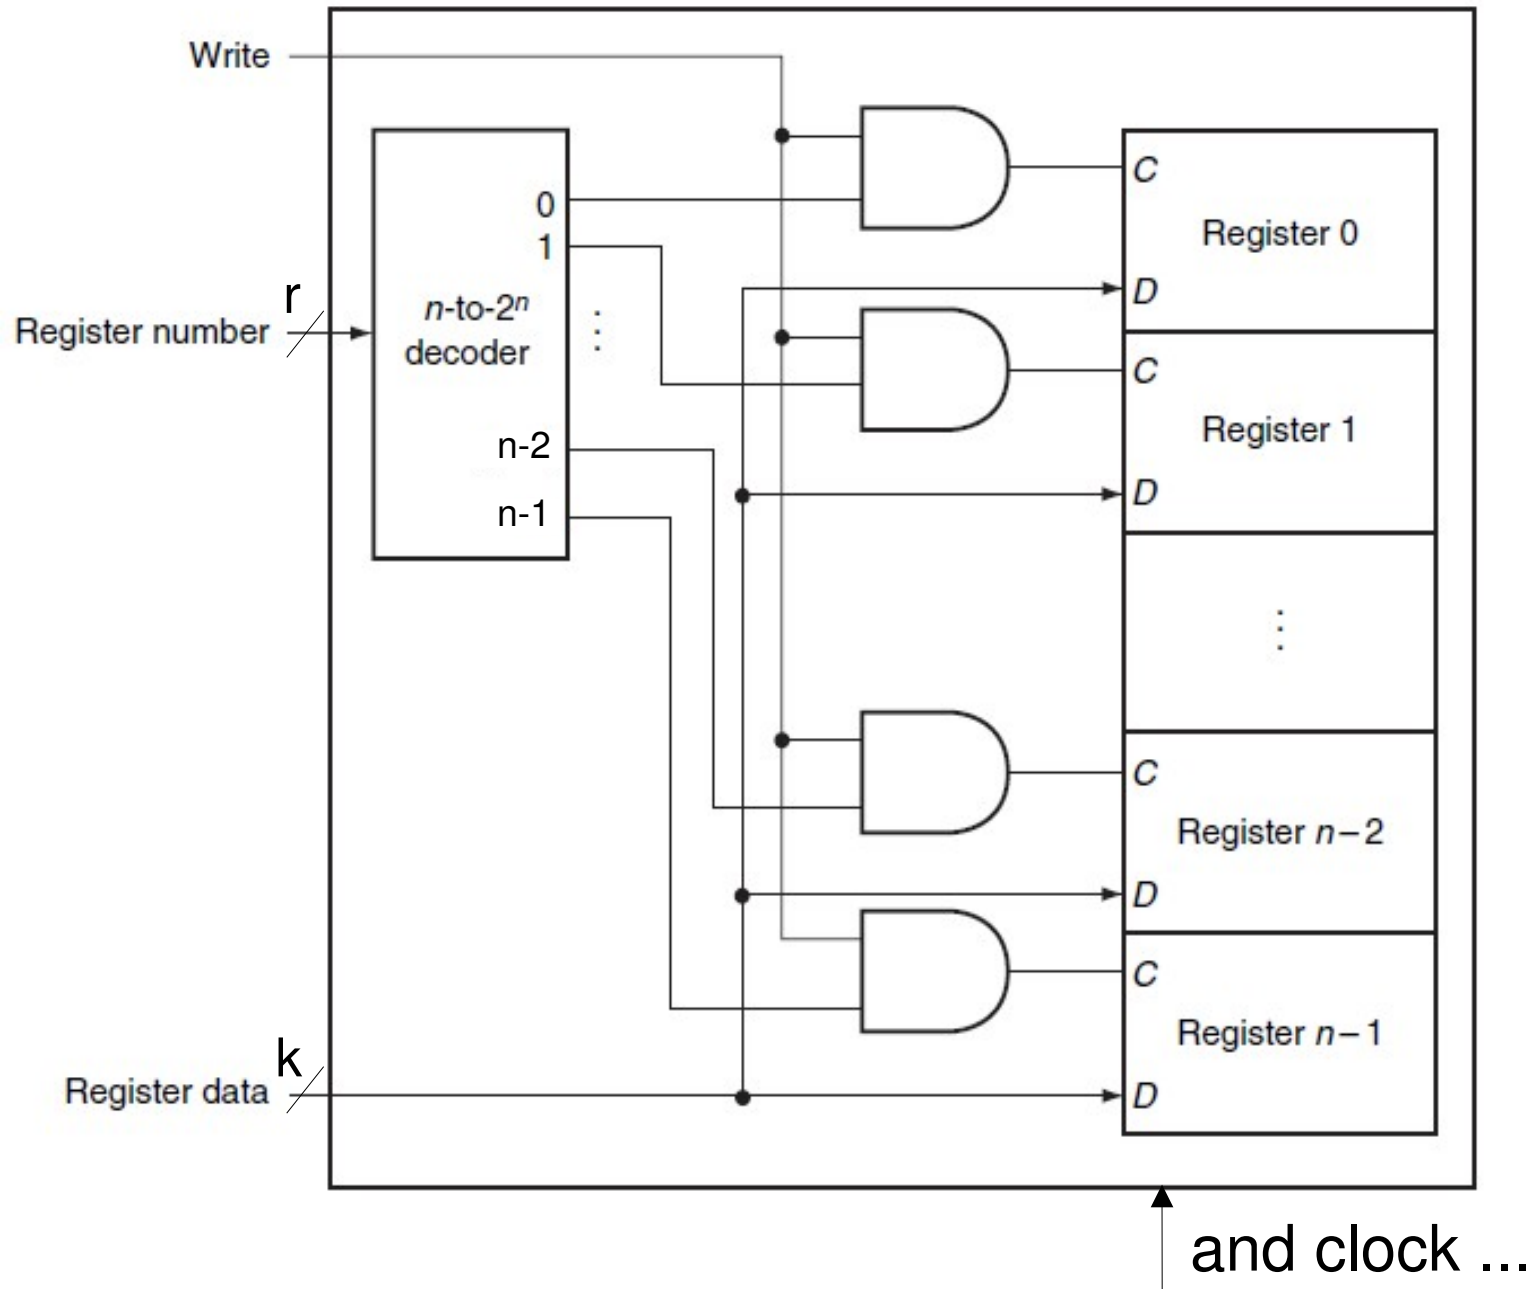




# Register File: read

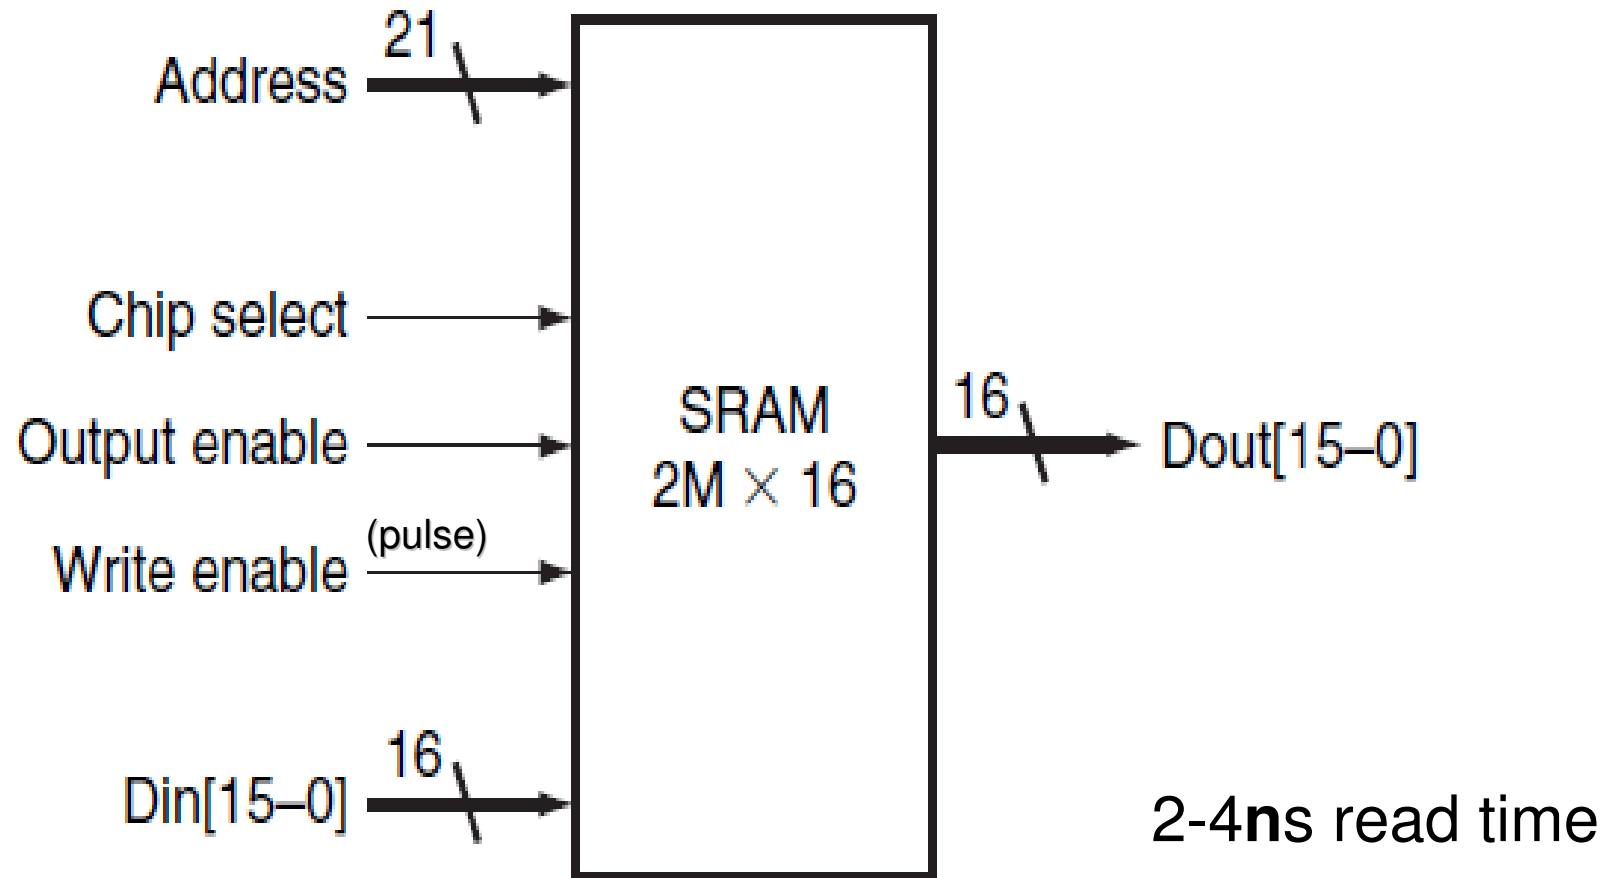


# Register File: write



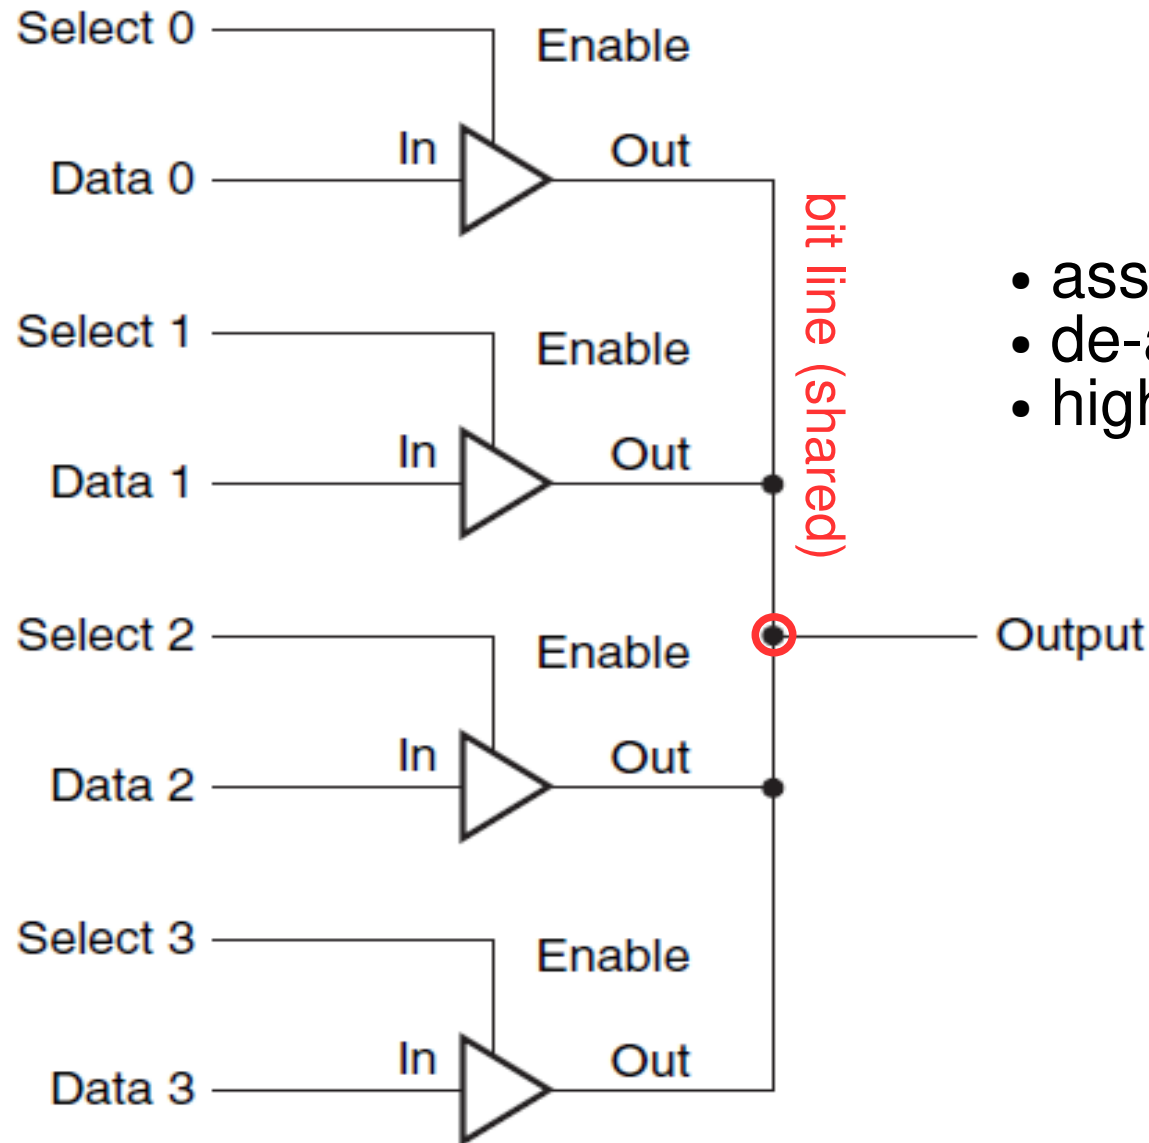
# Large Memory:

Static Random Access Memory (SRAM)  
vs. Dynamic Random Access Memory (DRAM)



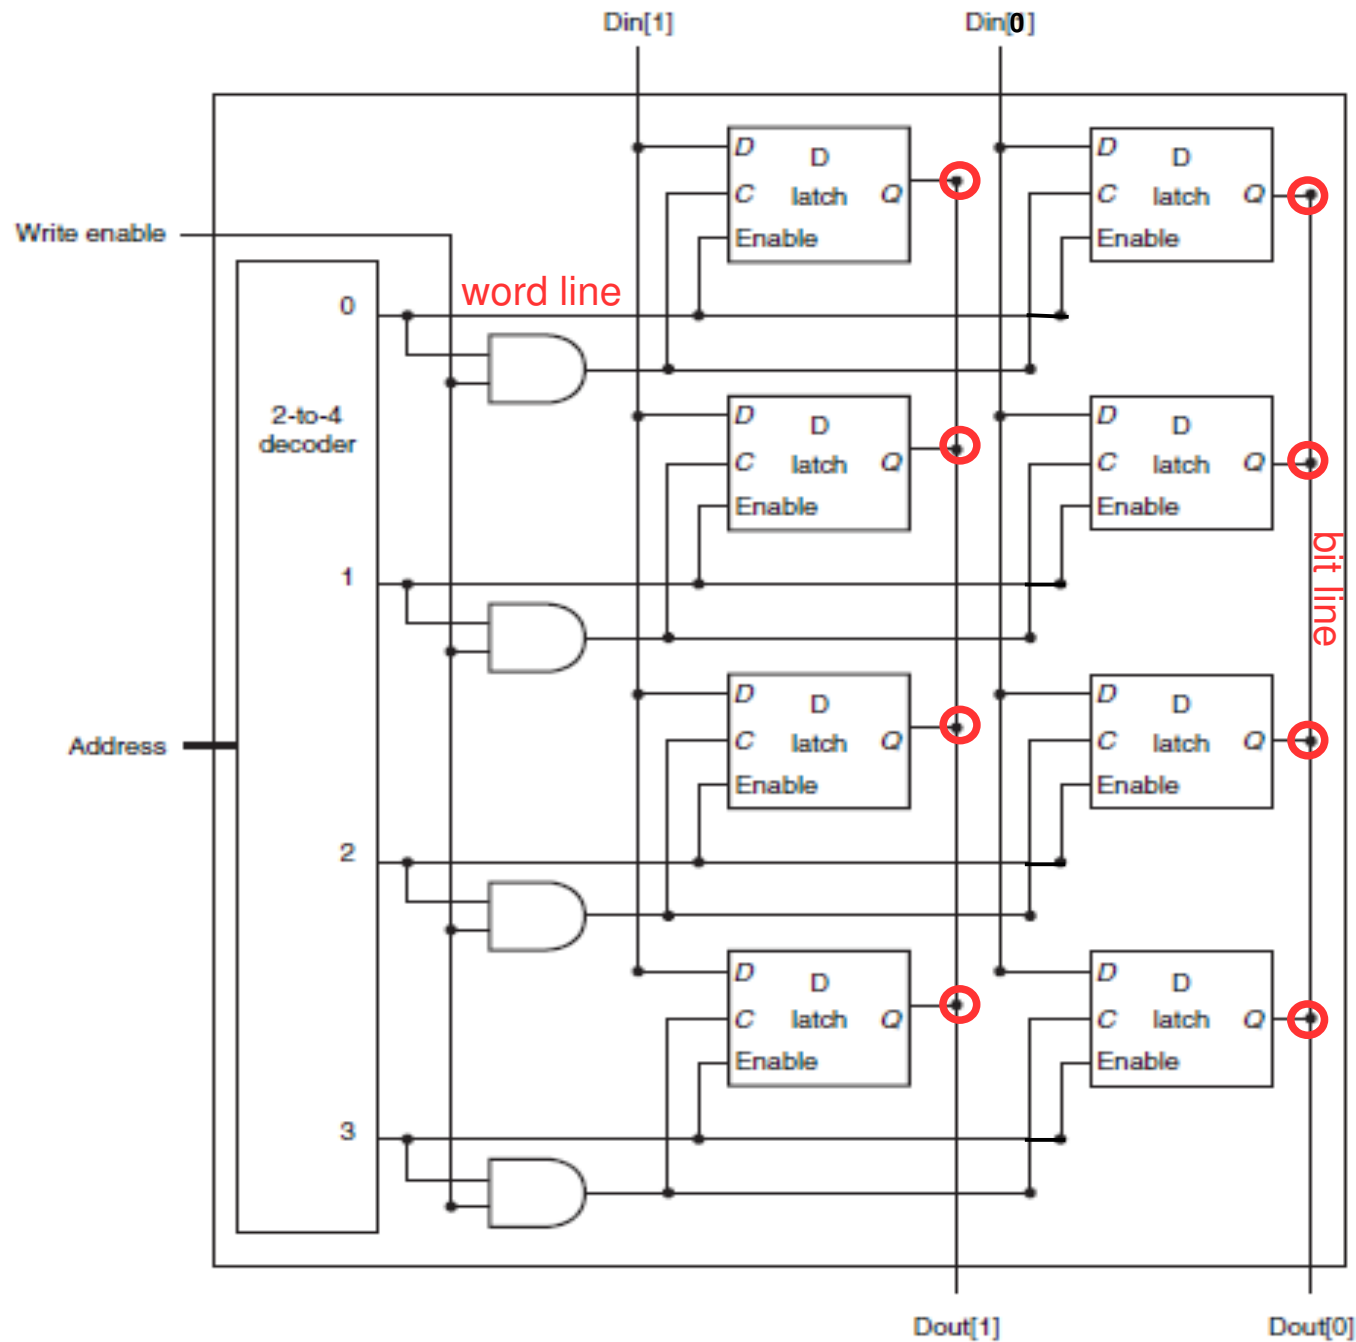
Random Access vs. Sequential Access: access time  
RAM vs. ROM (Read Only Memory)

# Three-state buffers (replace multiplexer)



- asserted
- de-asserted
- high-impedance

# 4x2 SRAM (output enable, chip select omitted)



No large multiplexer at output!

# SRAM array (read)

4Mx8

problem: need **huge** decoder !

solution: in analogy with carry lookahead adder where two levels were used

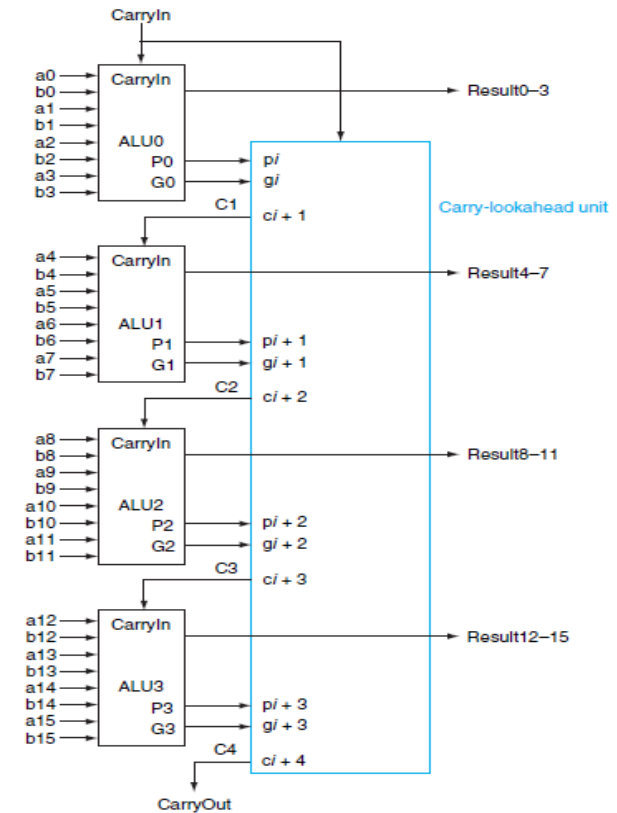
**use a two-level** decoding process

→ **slower**

→ but can use a reasonably-sized decoder

ripple adder tevel delay ==> carry look ahead

nu decoder te groot ==> 2 level decoding (volgende slide)





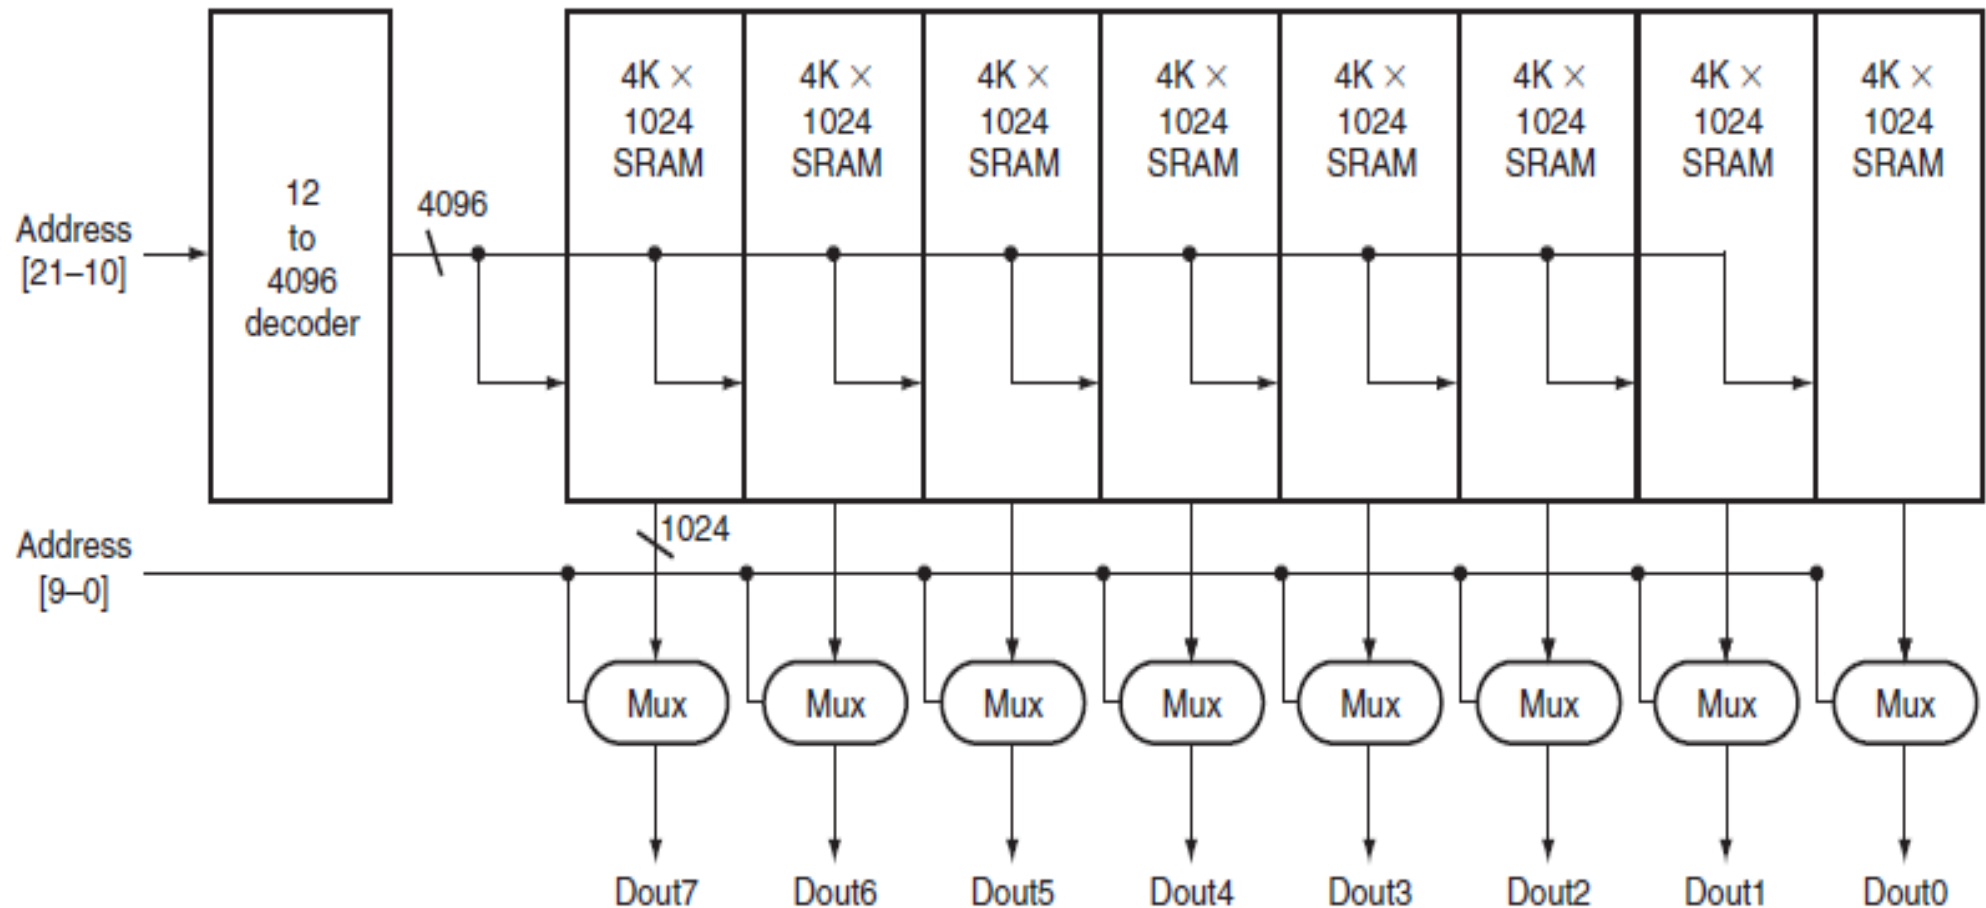
# SRAM array (read)

4Mx8

**2-level** decoding process  
to remove need for huge decoder

STATIC RAM !

verschillende geheugenbanken elk 1024 breed, allemaal flipflops  
we gaan in parallel info eruit halen



met mux alle 8 bits kiezen

Both register memory and cache RAM memory use **SRAM technology**: based on **flip-flops**.

### Register memory:

- + 1-level decoder and
- + no three-state buffers but multiplexer

### Cache memory

- + fast “buffer” between registers and main memory (data cache)
- + also instruction cache
- + different speeds and sizes: L1, L2, L3

snellste, iets minder snel, iets minder snek  
==> natuurlijk L1 dunder

# DRAM: Dynamic, use capacitor instead of flip-flop


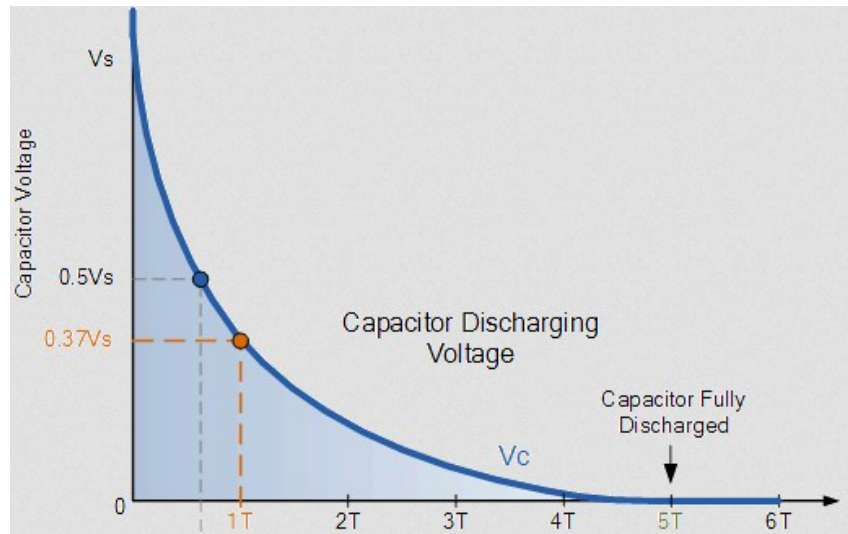
store information in capacitor →

**volatile** (even when powered), needs **refresh**

(andere is alleen volatile wanneer stroom uitvalt, hier ook refreshen wnr stroom er is)

only 1 transistor per bit (to access data) → denser and cheaper

dit is normale RAM van pc



maar 1 transistor  
ipv 24

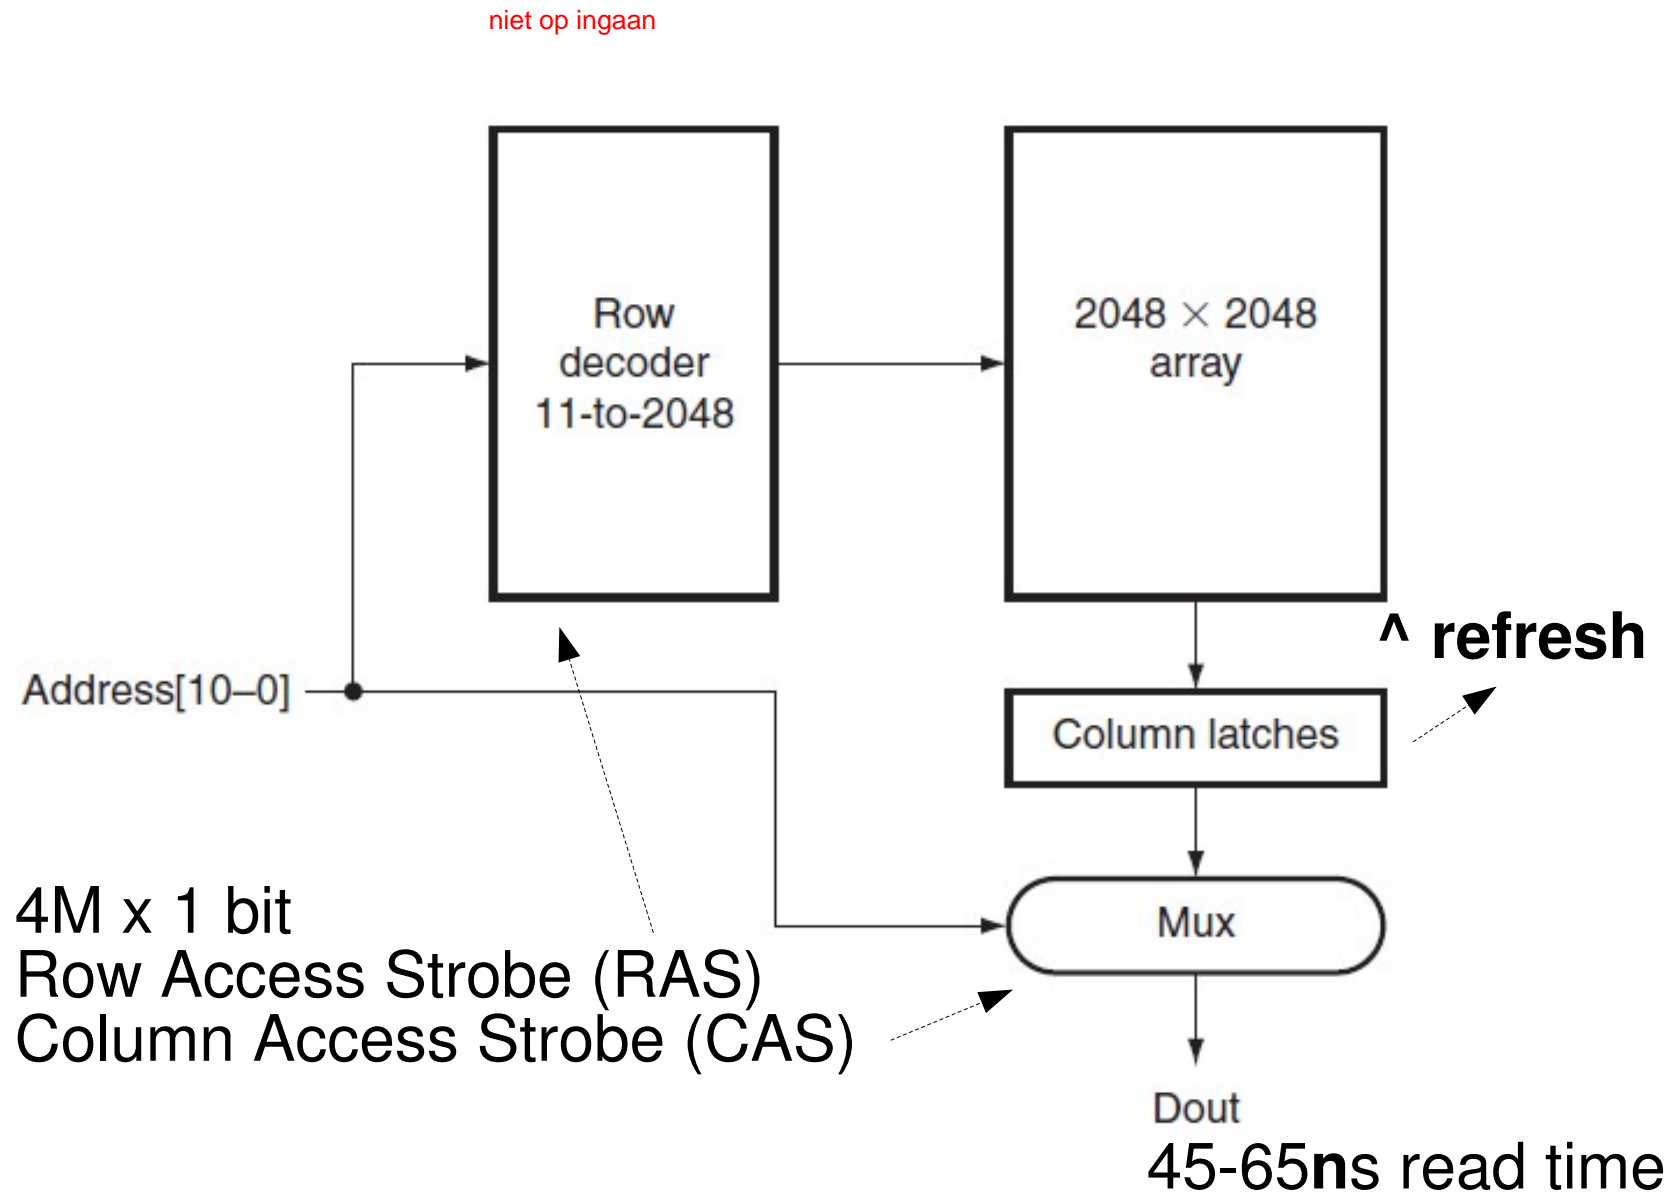
## read/write/refresh

(rate = **ms**, 64ms row refresh)

refresh by read and write-back

moet om de 64ms refreshen om te zeggen wat de waarde is

# DRAM array



Asynchronous DRAM (ADRAM) vs.


**Synchronous RAM**

**SDRAM** (for general memory) and **SSRAM** (for cache memory):

“burst” of data from a series of **sequential addresses**  
(use **clock**, not address)

Synchronized (on rising clock edge) with datapath

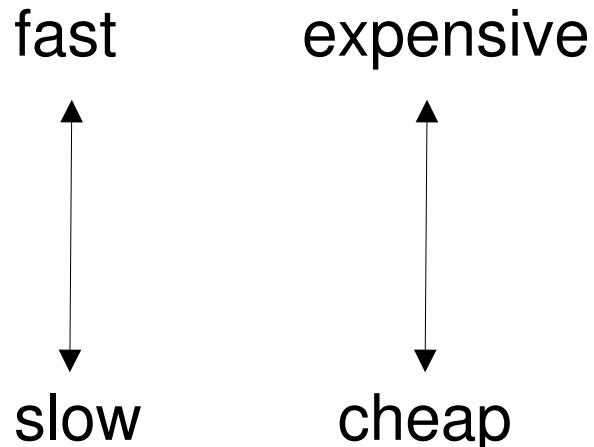
## SDR vs. DDR



“double pumping”: transfer data on both rising **and** falling edge

DDR, DDR2, DDR3, DDR4, DDR5

# Memory Technology

- Static RAM (SRAM)  
0.5ns – 2.5ns, \$2000 – \$5000 per GiB
  - caching  Dynamic RAM (DRAM)  
50ns – 70ns, \$20 – \$75 per GiB
  - virtueel geheugen  Magnetic disk  
5ms – 20ms, \$0.20 – \$2 per GiB
- 
- fast                      expensive
- slow                      cheap

## Ideal memory:

- Access time of SRAM
- Capacity and cost/GiB of magnetic disk

register spilling, hetgene da we ni meer nodig hebbe in STATiC  
==> register spillin gnaar DYNAMIC

# Principle of **Locality**

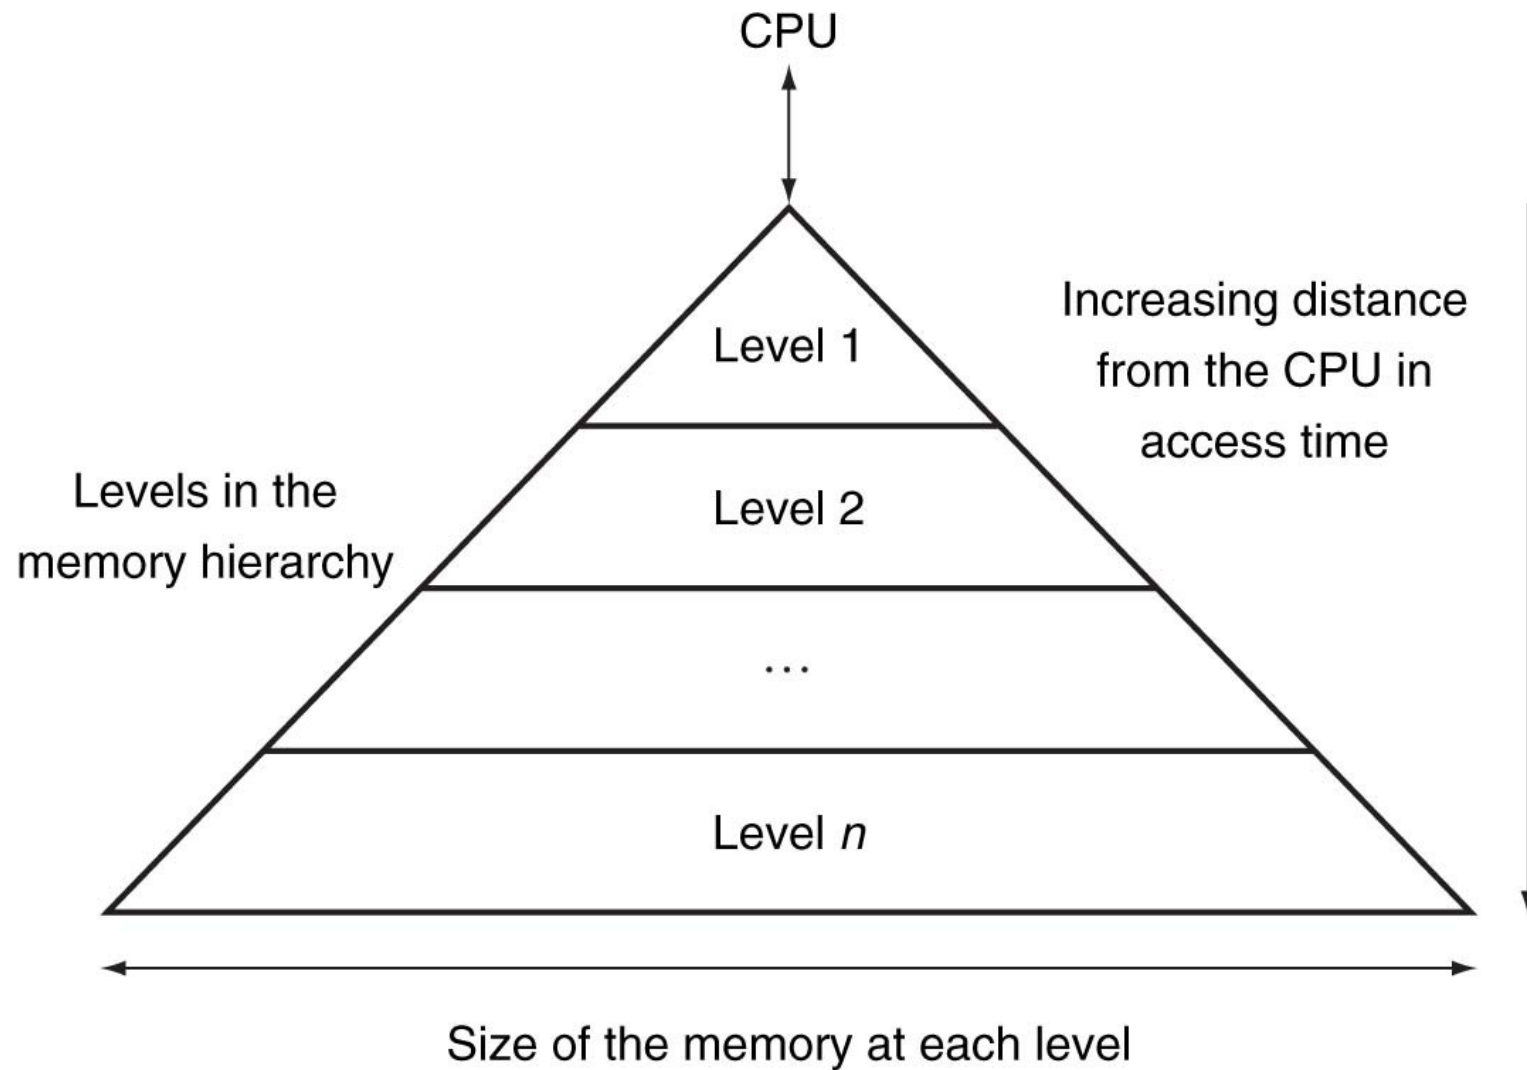
- Programs access a **small proportion** of their address space (instruction/data) at any time
- **Temporal** locality
  - Items accessed recently are likely to be accessed again soon
  - e.g., instructions in a loop, induction variables
- **Spatial** locality
  - Items near those accessed recently are likely to be accessed soon
  - E.g., sequential instruction access, array data

# Using Locality: **Memory Hierarchy**


- Store everything on disk
- Copy recently accessed (and nearby) items from disk to smaller DRAM memory
  - Main memory
- Copy more recently accessed (and nearby) items from DRAM to smaller SRAM memory
  - Cache memory attached to CPU



# Memory Hierarchy Levels



# Memory Hierarchy Levels



- **Block** (aka line): unit of copying
  - May be multiple words
- If accessed data is present in upper level
  - **Hit**: access satisfied by upper level
    - **Hit ratio**: hits/accesses
- If accessed data is absent
  - **Miss**: block copied from lower level
    - Time taken: **miss penalty**
    - **Miss ratio**: misses/accesses  
= 1 – hit ratio
  - Then accessed data supplied from upper level


probeer iets te vinden, is er niet  
naar volgend niveau

more about memory hierarchy algorithms in the Operating Systems course

- between cpu/registers and main memory (“cache”)
- between main memory and disk memory (“virtual memory”)

# Errors (in information transmission/storage)

→ encoding/decoding (hence the name “codec”, e.g., H.264)



# Error Detection Code

(in information transmission/storage)

1 bit parity: detect 1 bit error  
= distance-2 code

even/odd parity

# Error Correction Code (ECC)

Data Word	Code bits	Data	Code bits
0000	000	1000	111
0001	011	1001	100
0010	101	1010	010
0011	110	1011	001
0100	110	1100	001
0101	101	1101	010
0110	011	1110	100
0111	000	1111	111

0110 011 with 1 bit data error gives one of :

1110 011, 0010 011, 0100 011, 0111 011

manieren waarop het fout kan gaan  
1 bit error

ECCode 011 appears in table for


0110 (**edit distance** 1) ← most likely

0001 (edit distance 3, 2, 2, 2 respectively)

in een server wil je echt geen fouten, dan ook meer chips  
8 chips + 1 chip voor error correction bij te houden

# ECC RAM

aka “server class” memory




# Non-ECC RAM

# Finite State Machines (DFA, FSA, FSM)

DETERMINISTIC  
DFA

$$FSA = \langle S, s_i \in S, T \subseteq S' \times S' \times L, Acc \rangle$$



$$S = \{ s_1, s_2, s_3 \}$$

$$s_1 = s_1$$

$$T = \{ ((s_1, s_2), "a"), ((s_2, s_2), "b"), ((s_2, s_3), "c") \}$$

$$L = \{ "a", "b", "c" \}$$

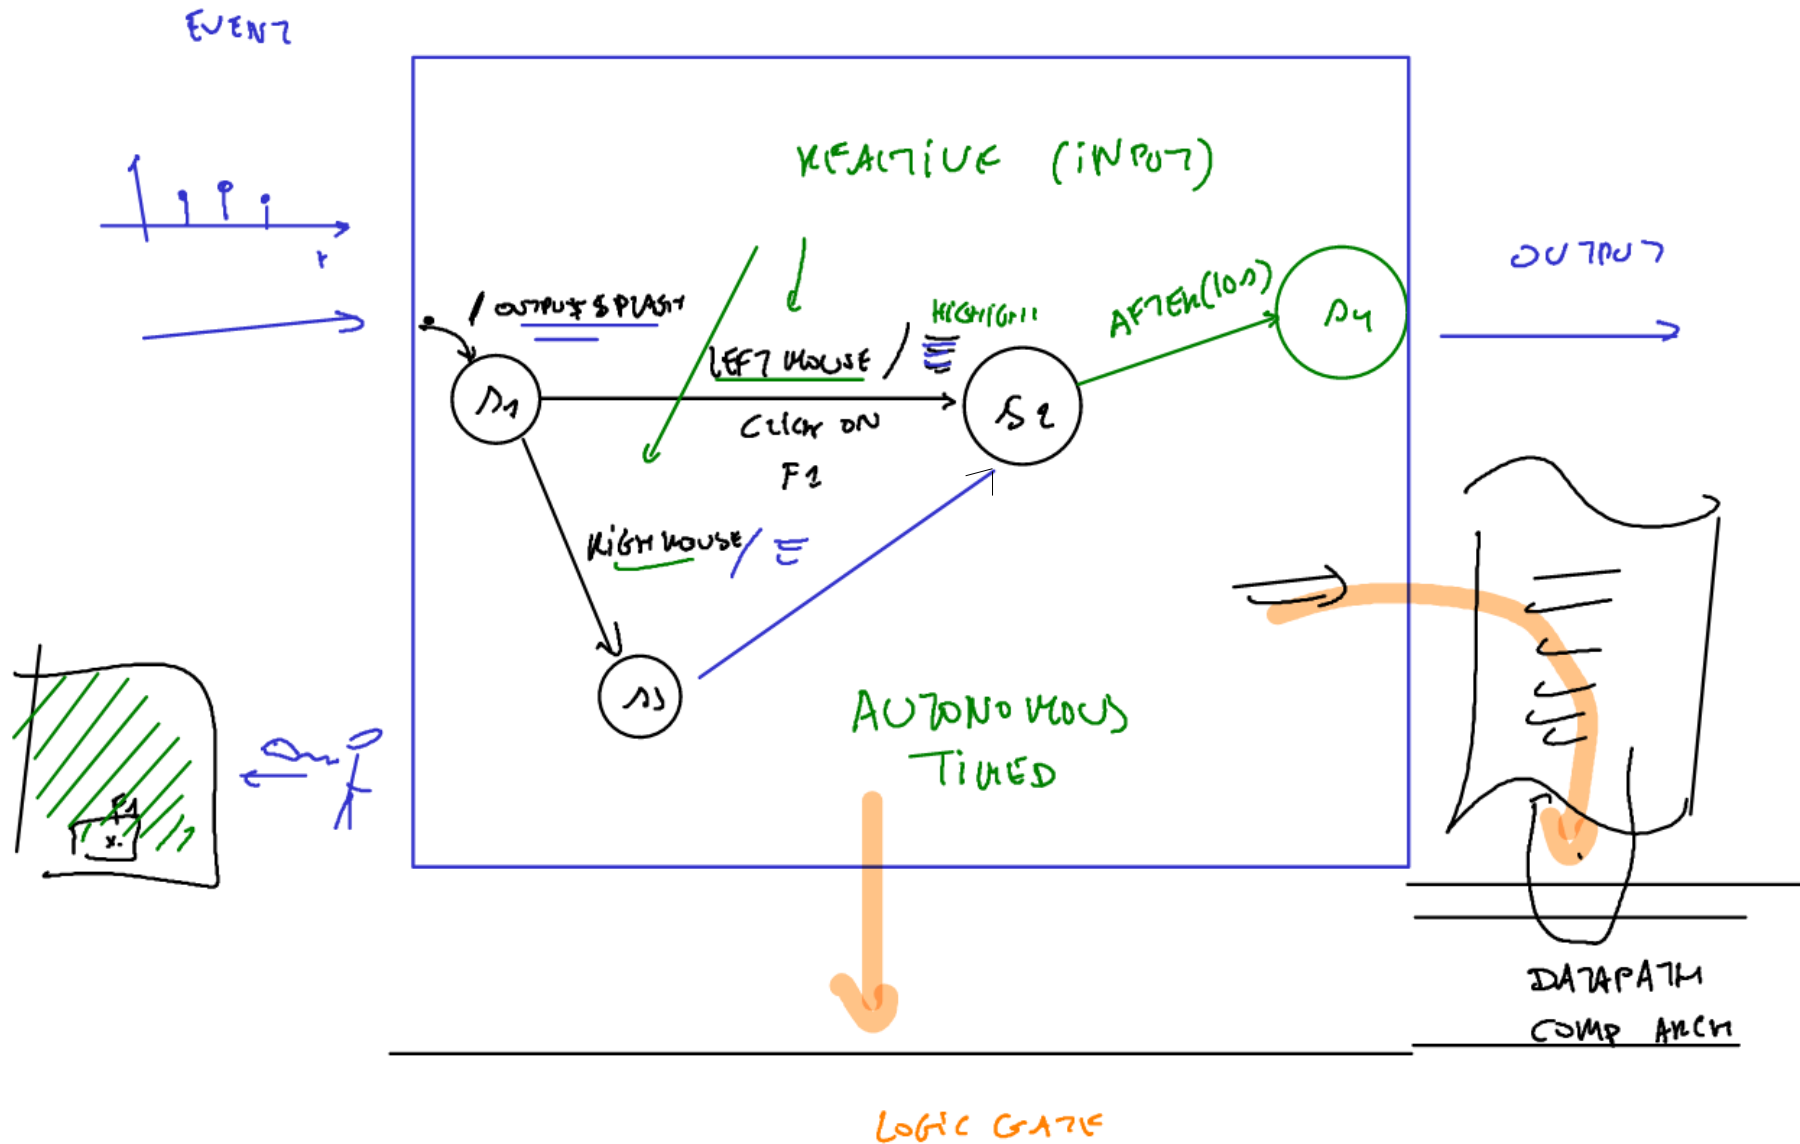
$$Acc = \{ s_3 \}$$

for "language" recognition



# Finite State Machines (Discrete Event – DE)

getimede, reactieve en autonome systemen



luistert naar input van buitenwereld en geeft output  
ook heeft de output een notie van tijd, "wacht tot ..."

for system (behaviour) specification (and synthesis)

reageren op input, outputs geven, ...




# Traffic Light (Discrete-Time – DT)

fairness (in case of conflict)


	Inputs		Next state
	NScar	EWcar	
NSgreen	0	0	NSgreen
NSgreen	0	1	EWgreen
NSgreen	1	0	NSgreen
NSgreen	1	1	EWgreen
EWgreen	0	0	EWgreen
EWgreen	0	1	EWgreen
EWgreen	1	0	NSgreen
EWgreen	1	1	NSgreen

	Outputs	
	NSlite	EWlite
NSgreen	1	0
EWgreen	0	1

# Traffic Light



# Traffic Light

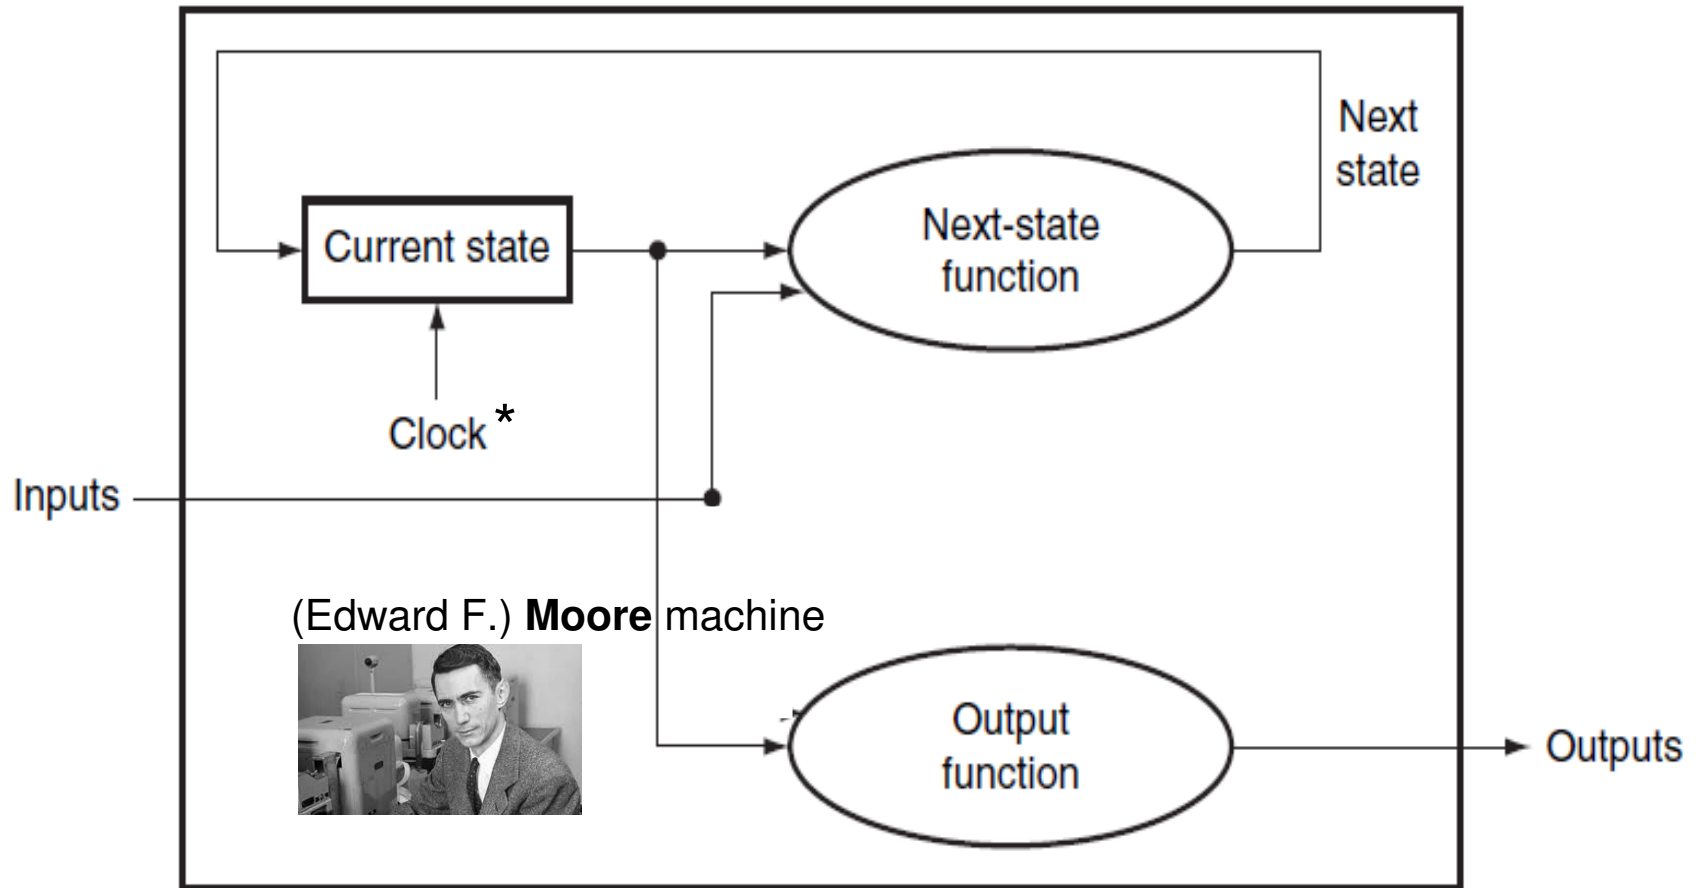


encode CurrentState (NSgreen or Ewgreen) in 1 bit

```
NextState = (~CurrentState . EWcar) + (CurrentState . ~NScar)  
NSlite     = ~CurrentState  
EWlite     = CurrentState
```

... evaluated every clock cycle ... add after(delay) ... ?


# Finite State Machines



\* Discrete-Time (DT) realization (vs. Discrete-Event (DE))

output enkel gebaseerd op huidige toestand

# Finite State Machines



\* Discrete-Time (DT) realization (vs. Discrete-Event (DE))

output gebaseerd op huidige toestand en input

# Traffic Light

