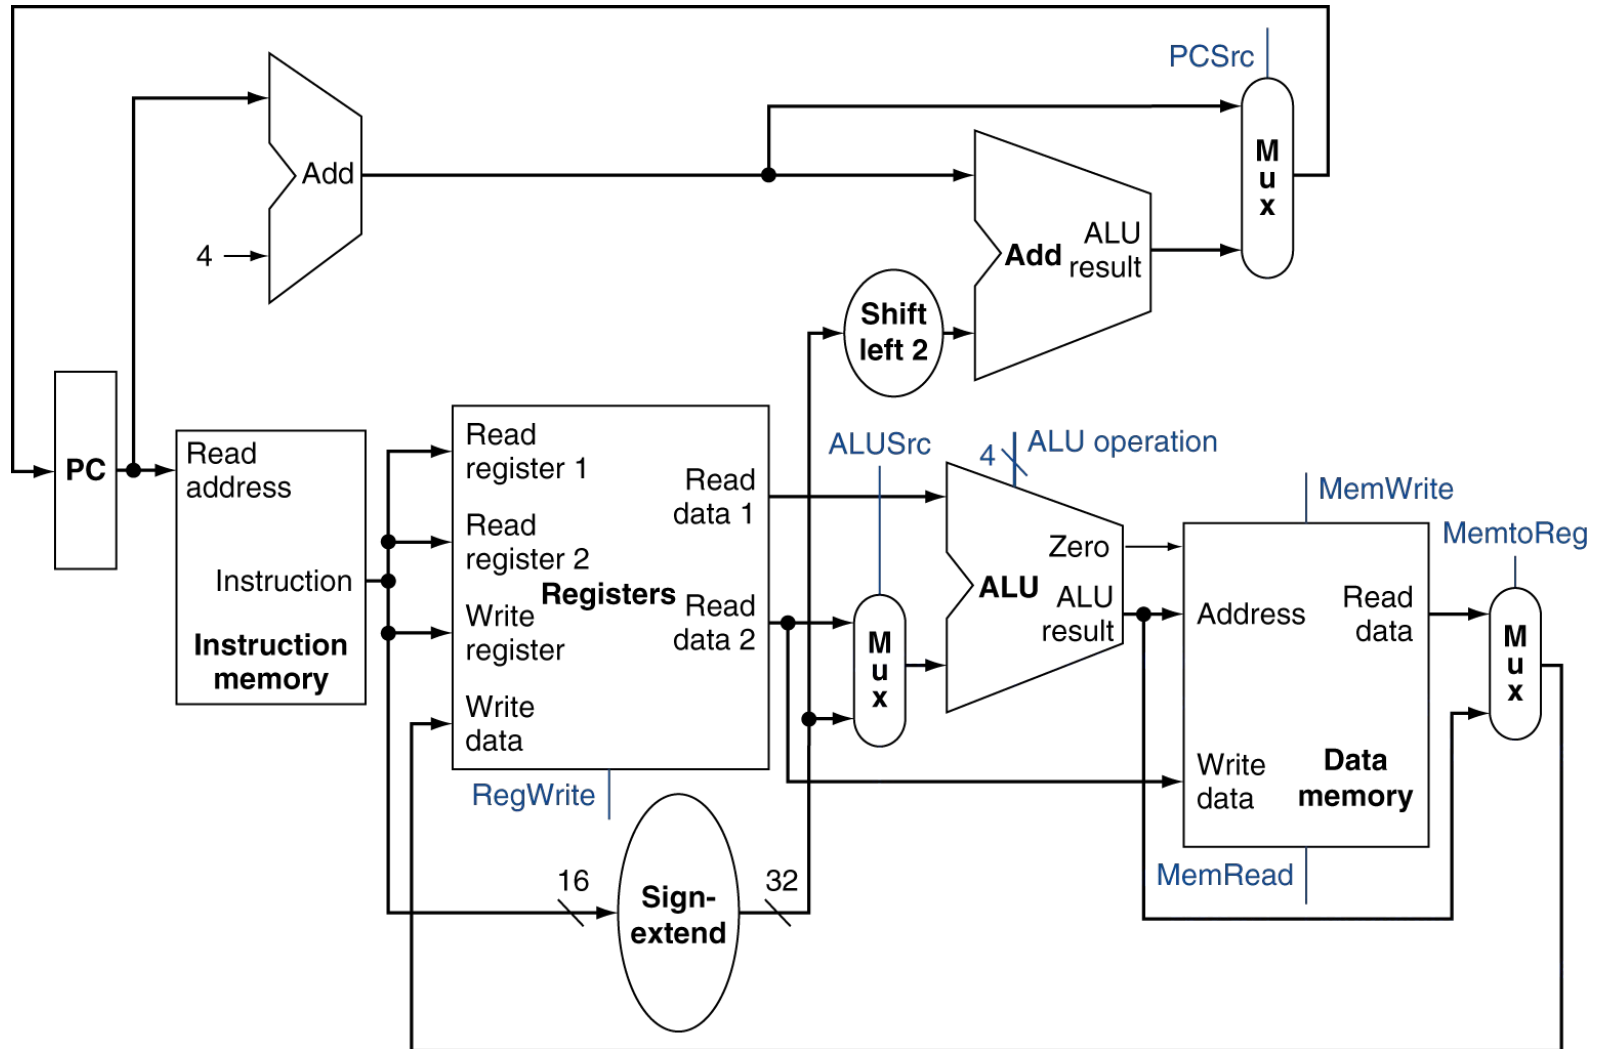


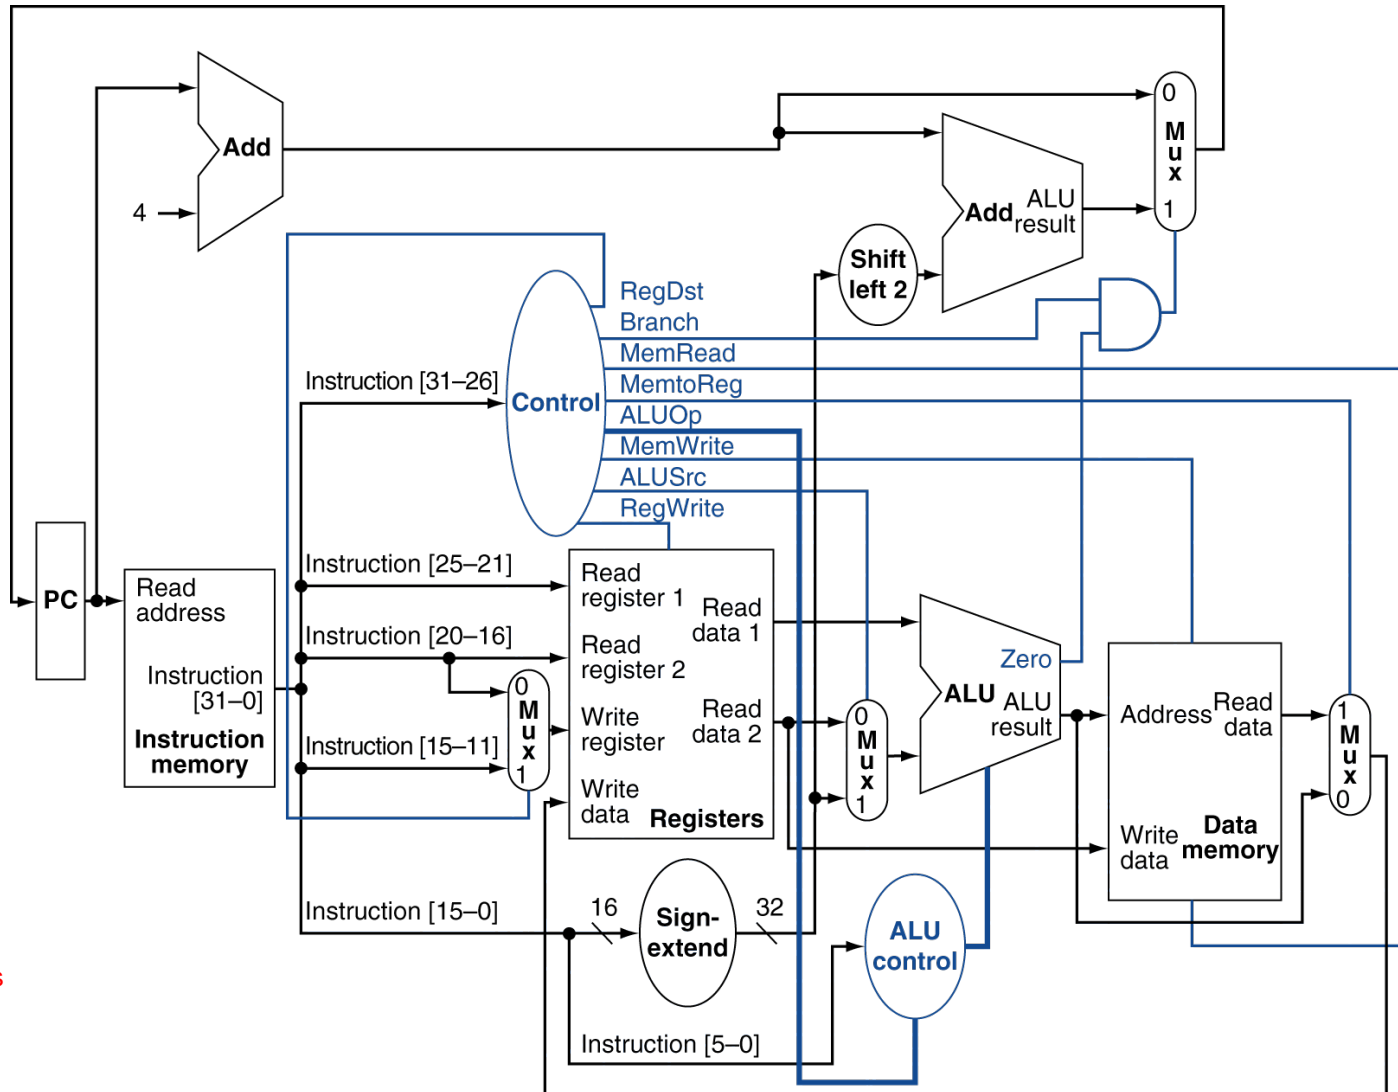
# Pipelining



# Full Datapath



# Datapath With Control



5 fases



# Performance Issues

- **Longest delay** determines **clock period**
    - **Critical path:** **load** instruction
      - Instruction memory → register file → ALU → data memory → register file
  - Not desirable to vary period for different instructions
- Violates the **regularity** design principle
- We will improve performance by **pipelining**, exploiting **instruction-level parallelism (ILP)**

# Pipelining Analogy

Pipelined laundry: **overlapping** execution

- **Parallelism** improves performance



- 4 loads:

- Speedup  
 $= 8/3.5 = 2.3$

- Non-stop:

- Speedup (for very large  $n$ )  
 $= \frac{Kn}{n + (K-1)} \approx K$   
 $K$  = number of stages

assume: all tasks take  
**same** amount of time

# MIPS Pipeline

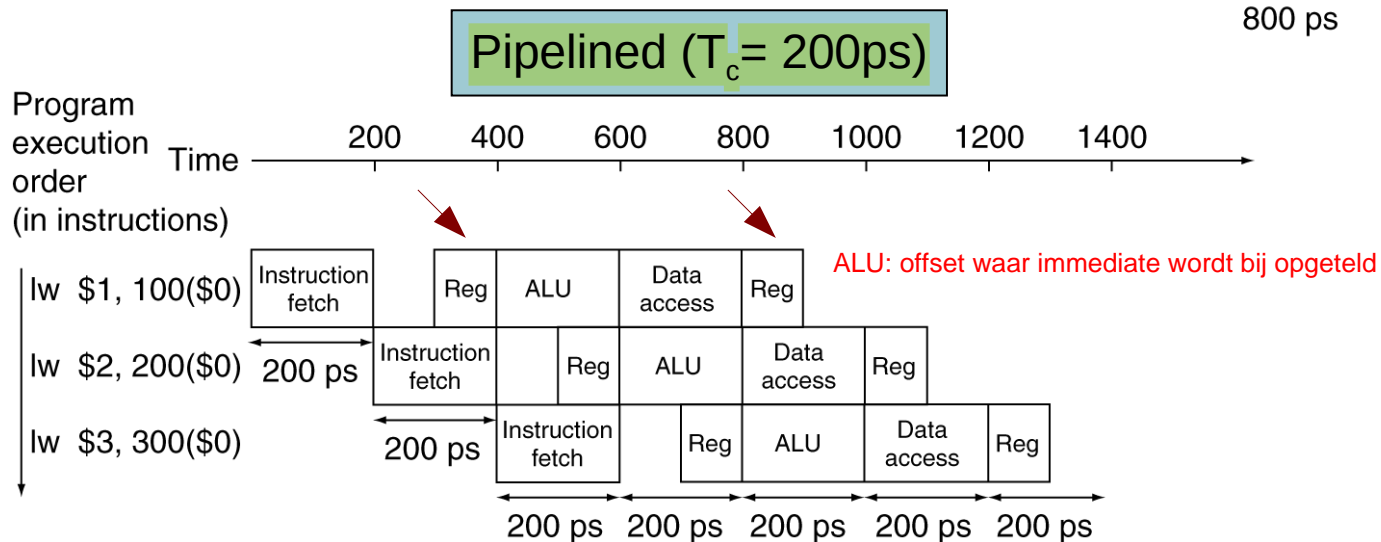
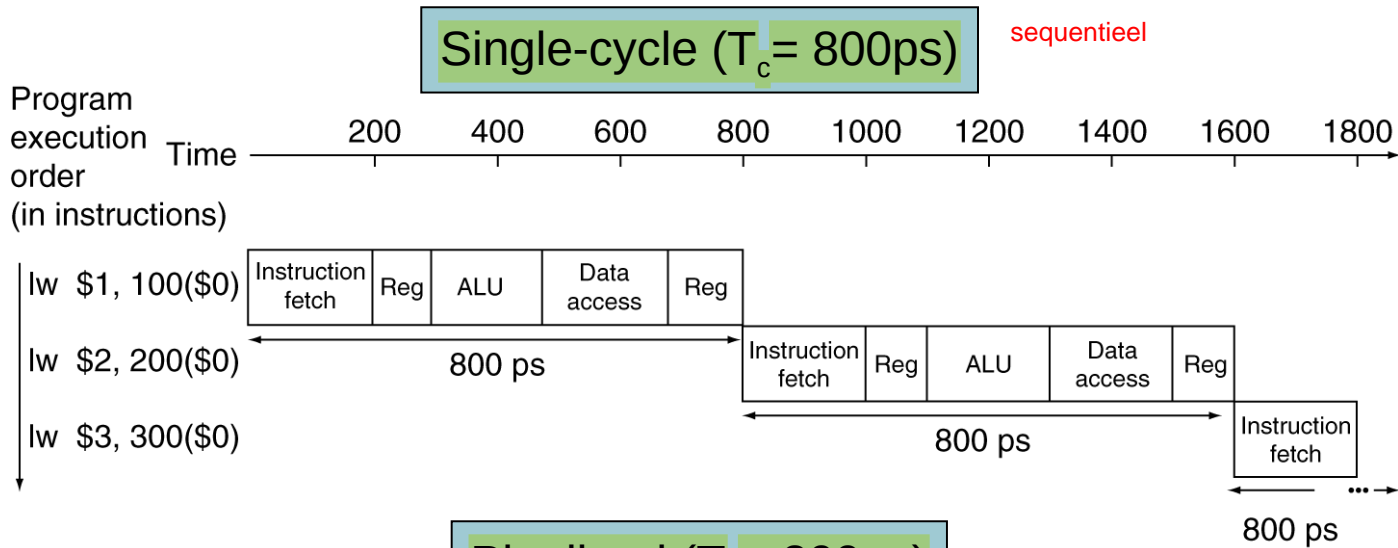
- Five **stages**, one step per stage
  1. **IF**: Instruction fetch from memory
  2. **ID**: Instruction decode and register read
  3. **EX**: Execute operation or calculate address
  4. **MEM**: Access memory operand
  5. **WB**: Write result back to register

# Pipeline Performance

- Assume time for stages is
  - 100ps for register read or write
  - 200ps for other stages

Instr	Instr fetch	Register read	ALU op	Memory access	Register write	Total time
lw	200ps	100 ps	200ps	200ps	100 ps	800ps
sw	200ps	100 ps	200ps	200ps		700ps
R-format	200ps	100 ps	200ps		100 ps	600ps
beq	200ps	100 ps	200ps			500ps

# Pipeline Performance





# Pipeline Speedup

- If all **stages are balanced**
  - *i.e.*, all take the same time
  - Time between instructions<sub>pipelined</sub> =  
$$\frac{\text{Time between instructions}_{\text{nonpipelined}}}{\text{Number of stages}}$$
- If not balanced, speedup is less
- Speedup due to increased **throughput**
  - **Latency** (time for each instruction) does **not** decrease!

# Pipelining and ISA Design

- MIPS ISA designed for **pipelining**
  - **All** instructions are **same length** (32-bits)
    - Easier to fetch and decode in **one cycle**
    - cf. x86: 1- to 17-byte instructions
  - **Few** and **regular** instruction **formats**
    - Can decode **and** read registers in one step
  - **Load/store** addressing
    - Can **calculate** address in 3<sup>rd</sup> stage, access memory in 4<sup>th</sup> stage
  - **Alignment** of memory operands
    - Memory access takes only **one cycle**

# Hazards



- Situations that **prevent** starting the next instruction in the next cycle. Need to wait!
- **Structure** hazards
  - A required resource is busy
- **Data** hazard
  - Need to wait for a previous instruction to complete its data read/write
- **Control** hazard
  - Deciding on control action depends on a previous instruction

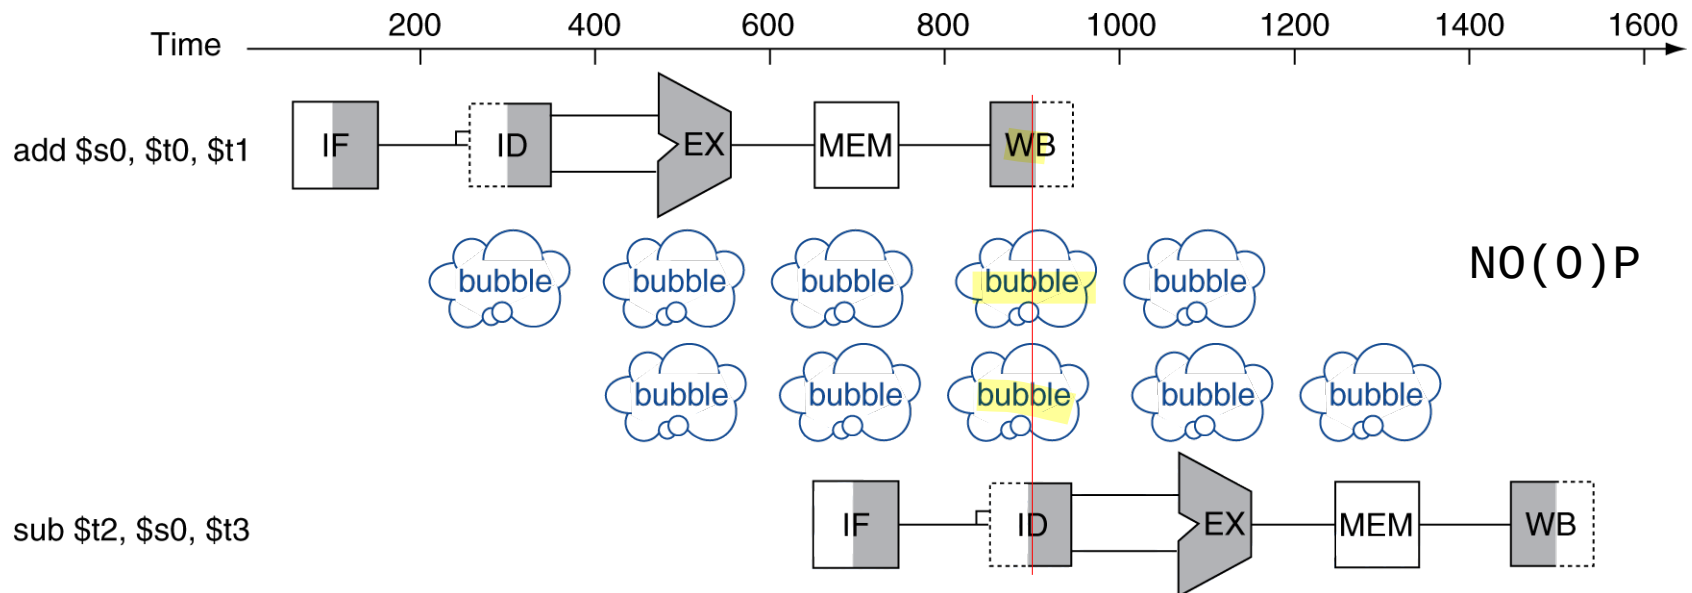
# Structure Hazards

- Conflict for **use** of a resource
- In MIPS pipeline with a **single memory**
  - Load/store requires data access
  - Instruction fetch requires memory access too → would have to **stall** for that cycle
    - Would cause a pipeline “**bubble**”
- Hence, pipelined datapaths require **separate** instruction/data memories
  - or separate instruction/data **caches**

# Data Hazards

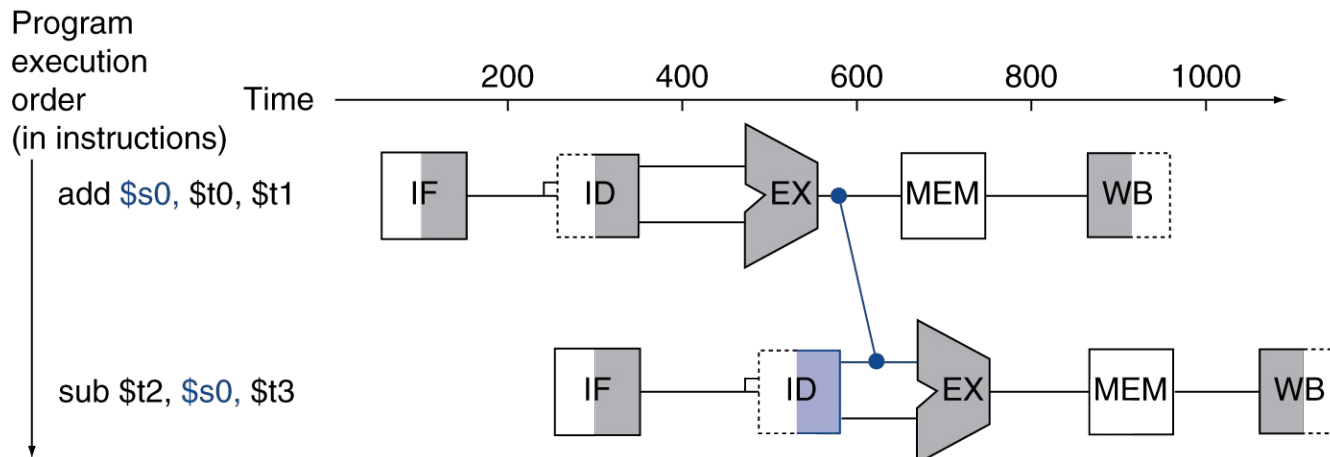
- An instruction depends on **completion** of data access by a **previous** instruction

- add **\$s0**, \$t0, \$t1
  - sub \$t2, **\$s0**, \$t3



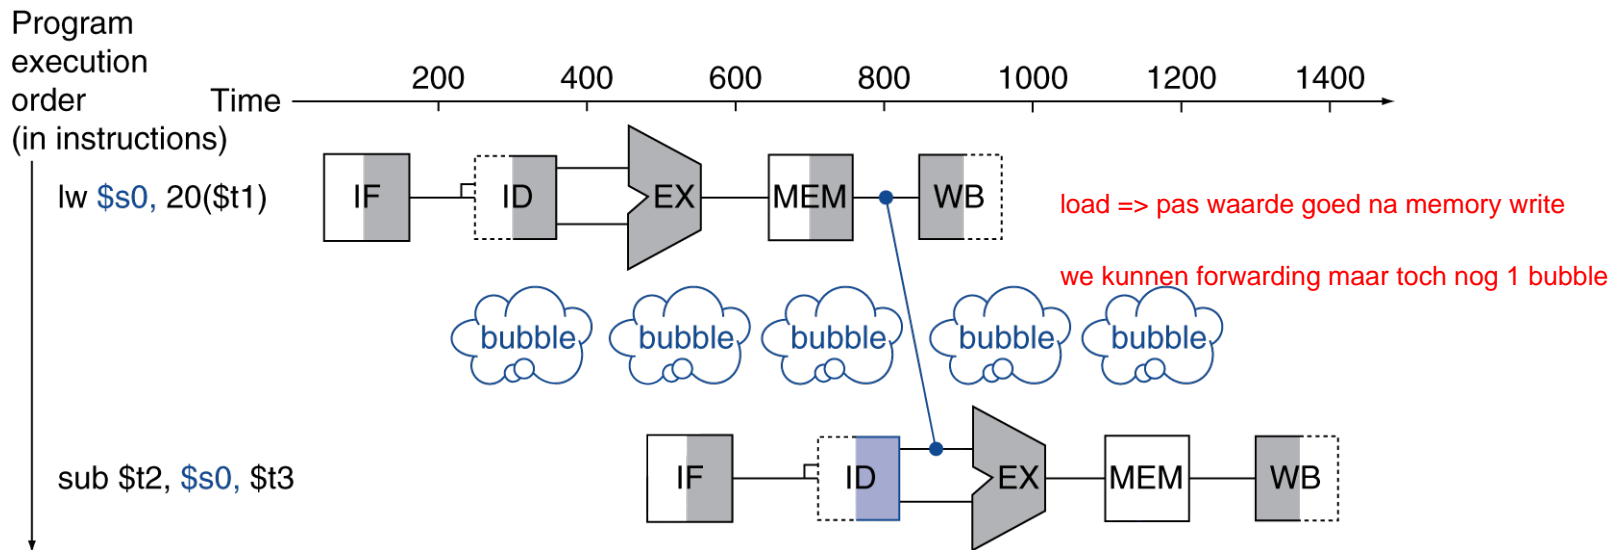
# Forwarding (aka Bypassing)

- Use result when it is computed
  - **Do not wait** for it to be **stored** in a **register**
  - Requires **extra connections** in the datapath (**extra hardware to detect and fix hazard**)



# Load-Use Data Hazard

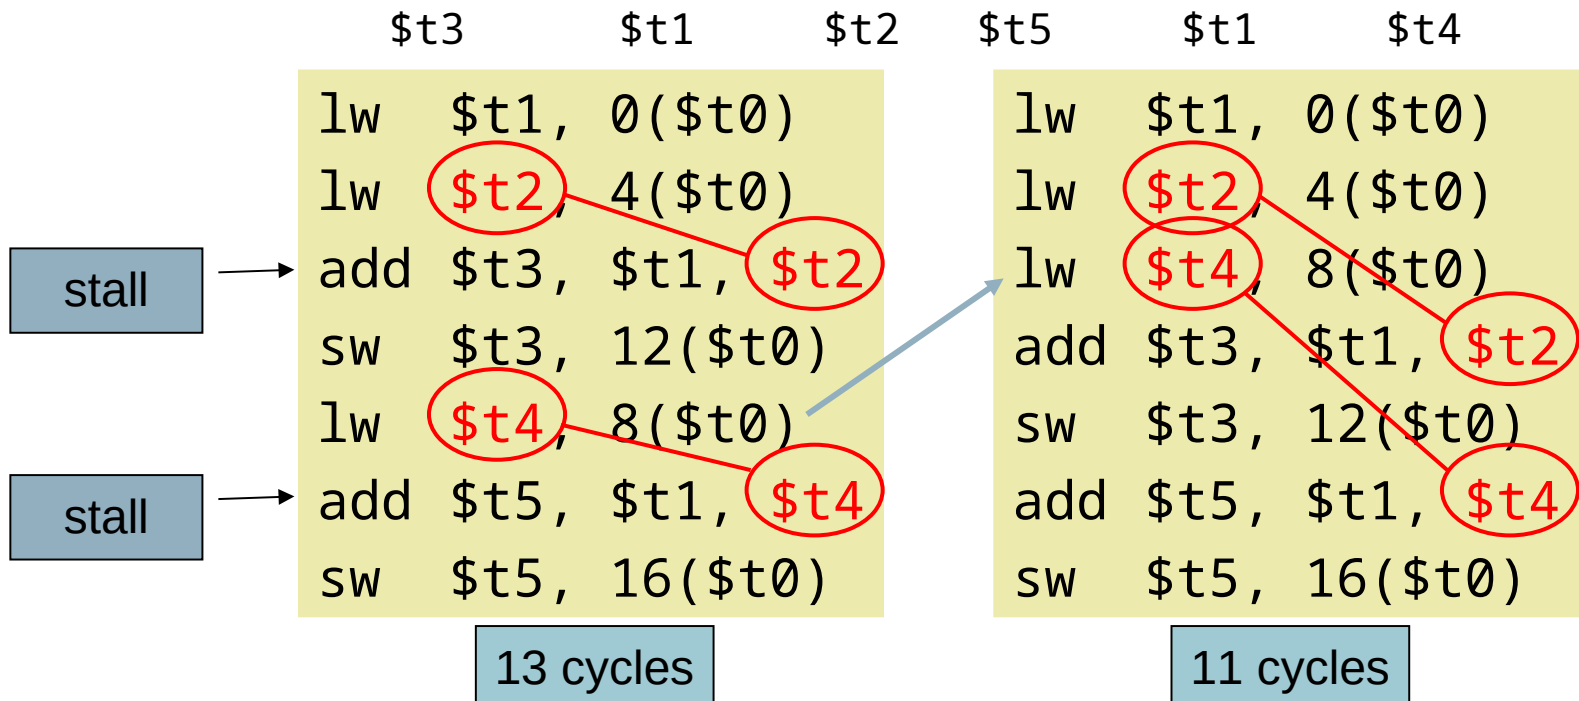
- **Can't always avoid** stalls by forwarding
  - If value not computed when needed
  - Can't forward backward in time! (**causality**)



load gevolgd door add dus 1x 200ps delay met forwarding  
anders verkeerde waarde

# Code Scheduling to Avoid Stalls

- **Re-order** code to avoid use of load result in the next instruction
- code for  $A = B + E$ ;  $C = B + F$ ;



$n + (K-1)$   
K stages, n instructions  $\Rightarrow 7 + (5-1) = 11$

ipv tijd te waste, plaats andere instructie tussen de tijd dat we moeten wachten zodat er minder tijd/cycles zijn **16**

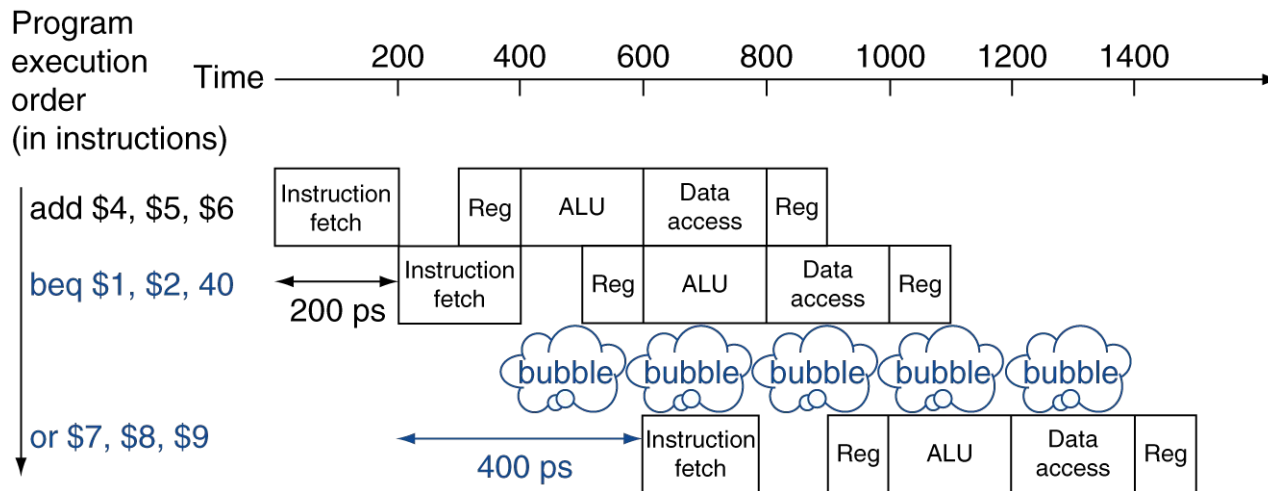


# Control Hazards

- **Branch** determines **flow of control**
  - **Fetching next instruction**
    - depends on test outcome
  - Pipeline **can't** always **fetch correct instruction**
    - still working on **ID** stage of branch
- In MIPS pipeline:
  - Need to **compare** registers and **compute target** **early** in the pipeline
    - add **hardware** to do it **in ID stage**  
(ID = Instruction Decode and Register Read)

# Stall on Branch

- Wait until branch outcome determined before fetching next instruction

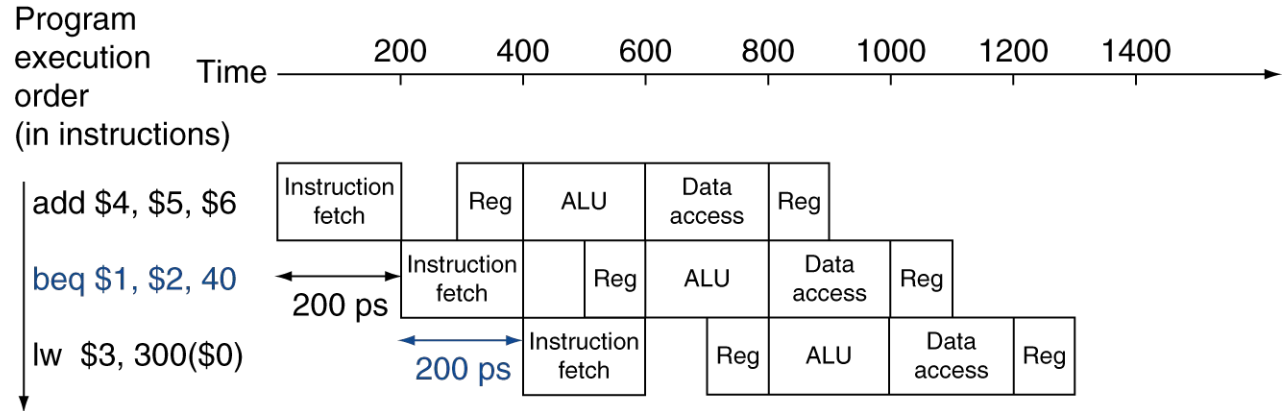


# Branch Prediction

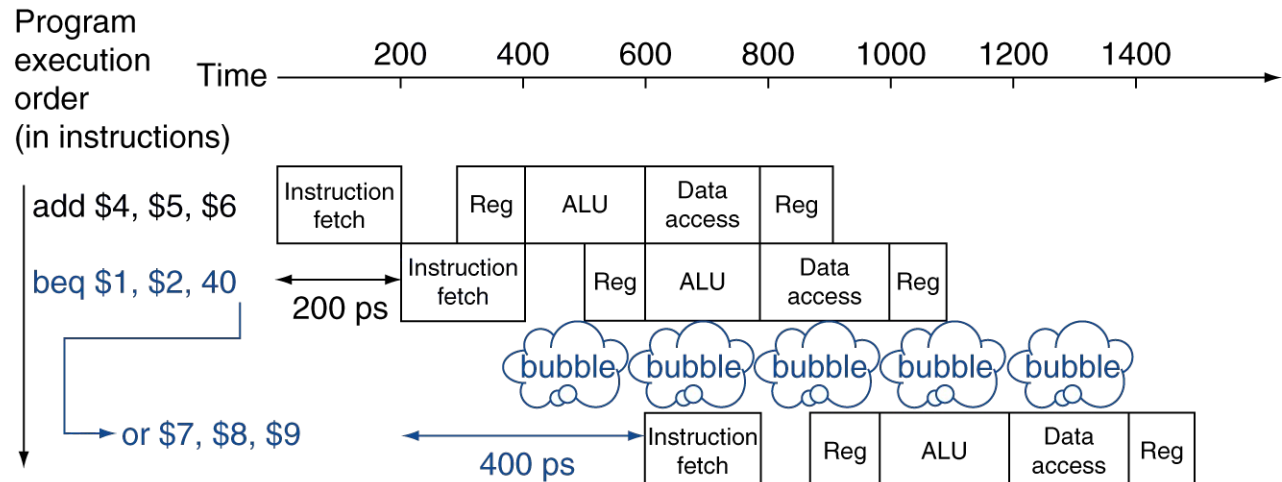
- Longer pipelines **can't** readily **determine** branch outcome **early**
  - Stall penalty becomes unacceptable
- **Predict** outcome of branch
  - **Only** stall if **prediction is wrong**
- In MIPS pipeline
  - Can predict **branches not taken**
  - **Fetch** instruction **after branch**, with **no delay**

# MIPS with Predict Not Taken

Prediction  
correct



Prediction  
incorrect



# More-Realistic Branch Prediction

- **Static** branch prediction
  - Based on **typical** branch behavior
  - Example: **loop** and **if**-statement branches
    - Predict backward branches taken
    - Predict forward branches not taken
- **Dynamic** branch prediction
  - Hardware **measures actual** branch behavior
    - e.g., **record** recent **history** of each branch
  - **Extrapolate:**  
assume future behavior will continue the trend
    - When wrong, stall while **re-fetching**, and **update history**

bubble

# Pipeline Summary

## The BIG Picture

- Pipelining improves performance by increasing instruction **throughput**
  - Executes multiple instructions in **parallel**
  - Each instruction has the **same latency**
- Subject to **hazards**
  - structure, data, control
- Instruction set design (**ISA**) affects **complexity** of pipeline implementation

# MIPS Pipelined Datapath

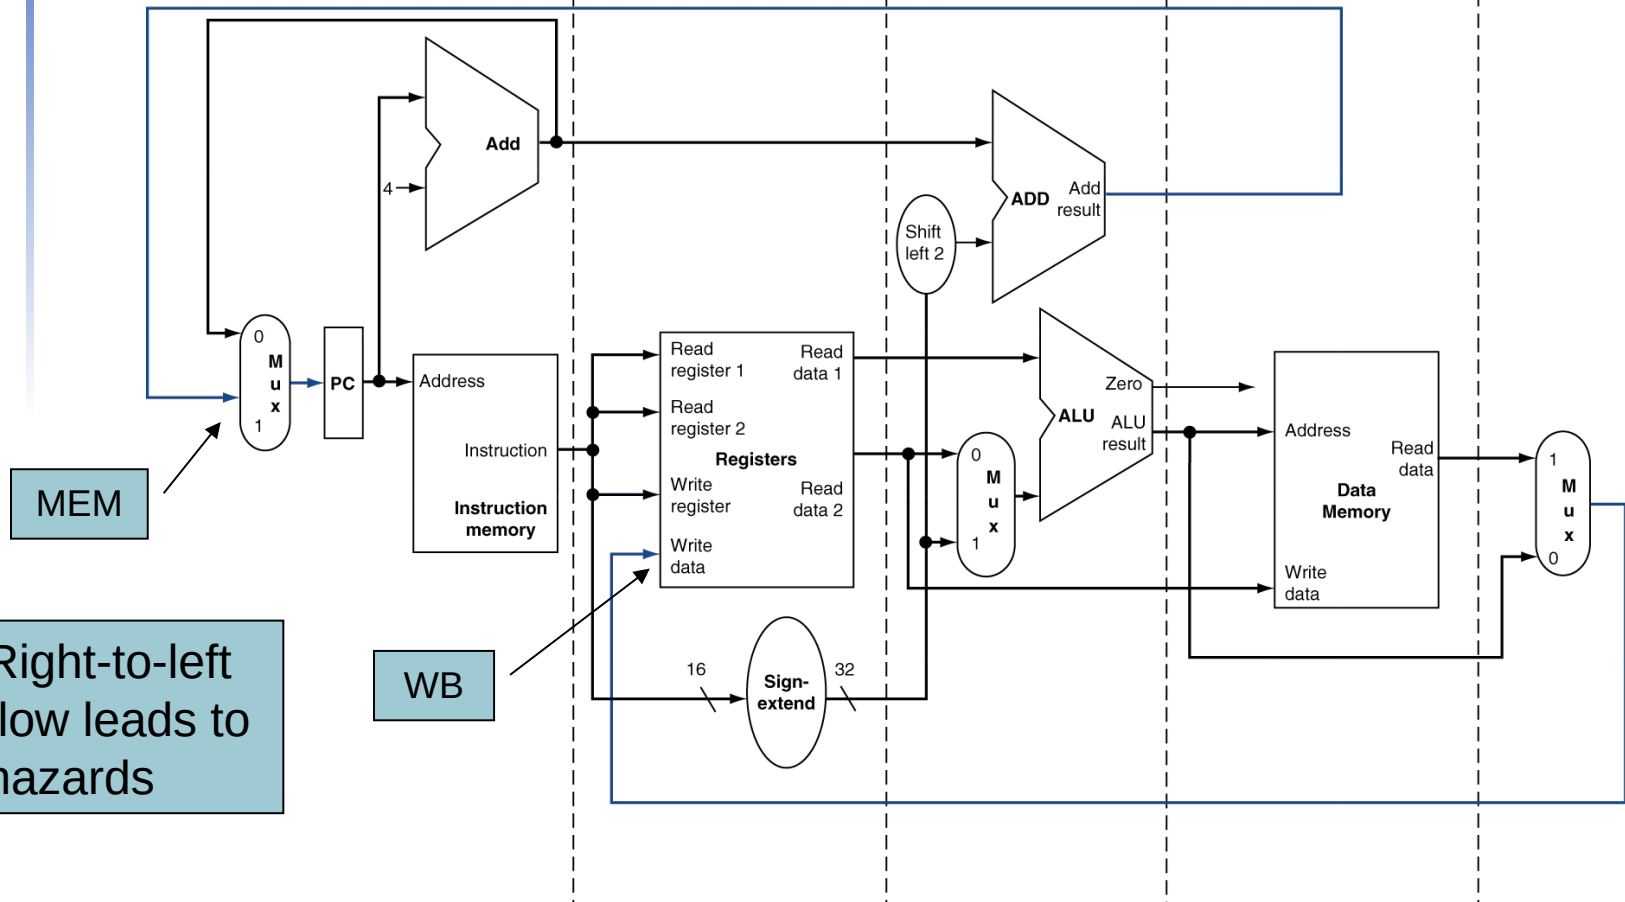
IF: Instruction fetch

ID: Instruction decode/  
register file read

EX: Execute/  
address calculation

MEM: Memory access

WB: Write back

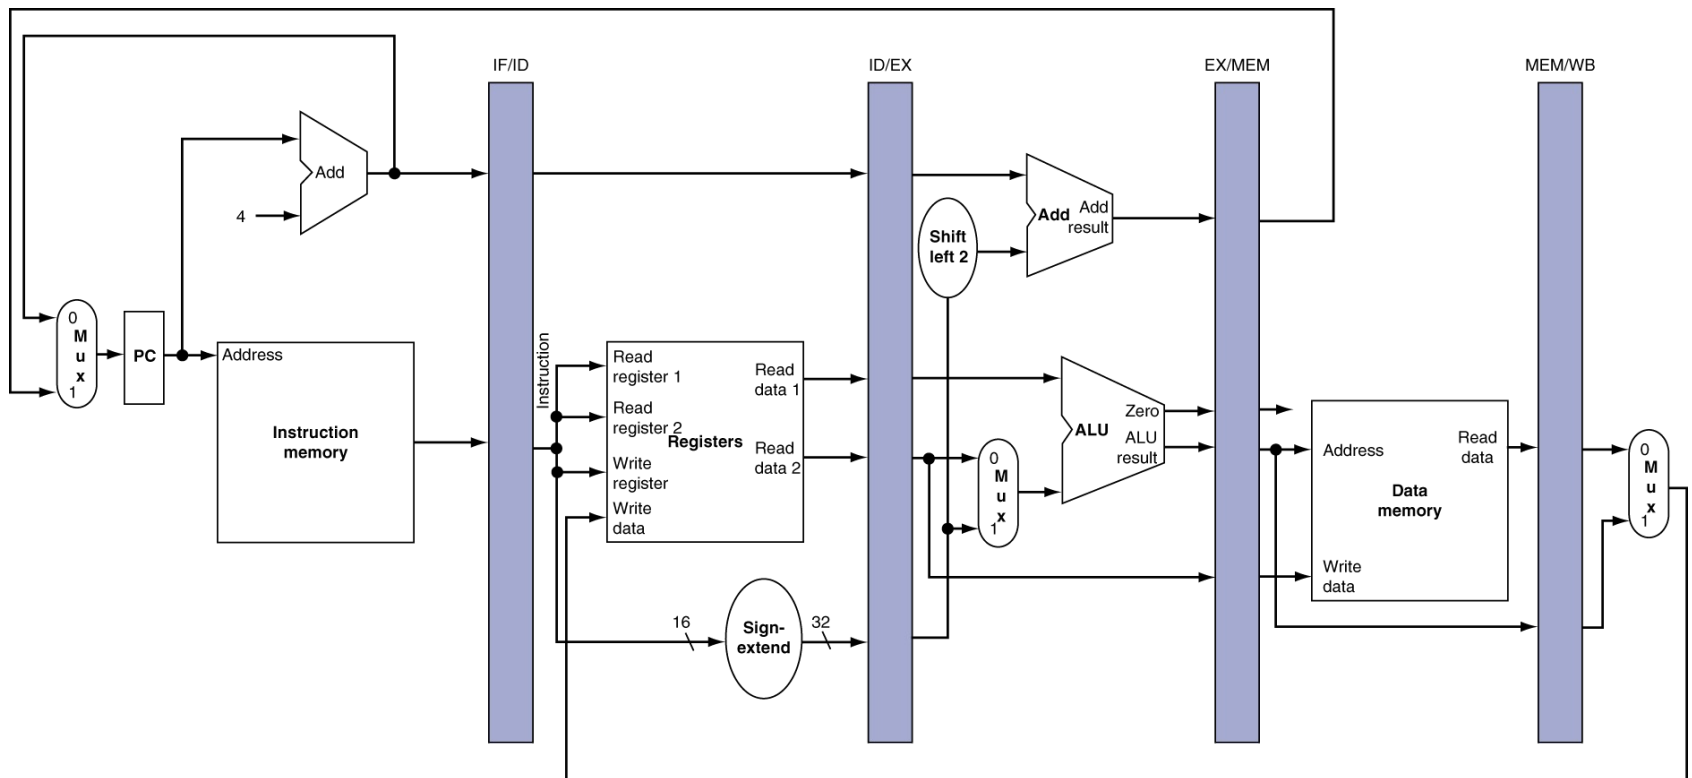


Right-to-left  
flow leads to  
hazards

WB

# Pipeline registers

- Need registers between stages to hold information produced in previous cycle

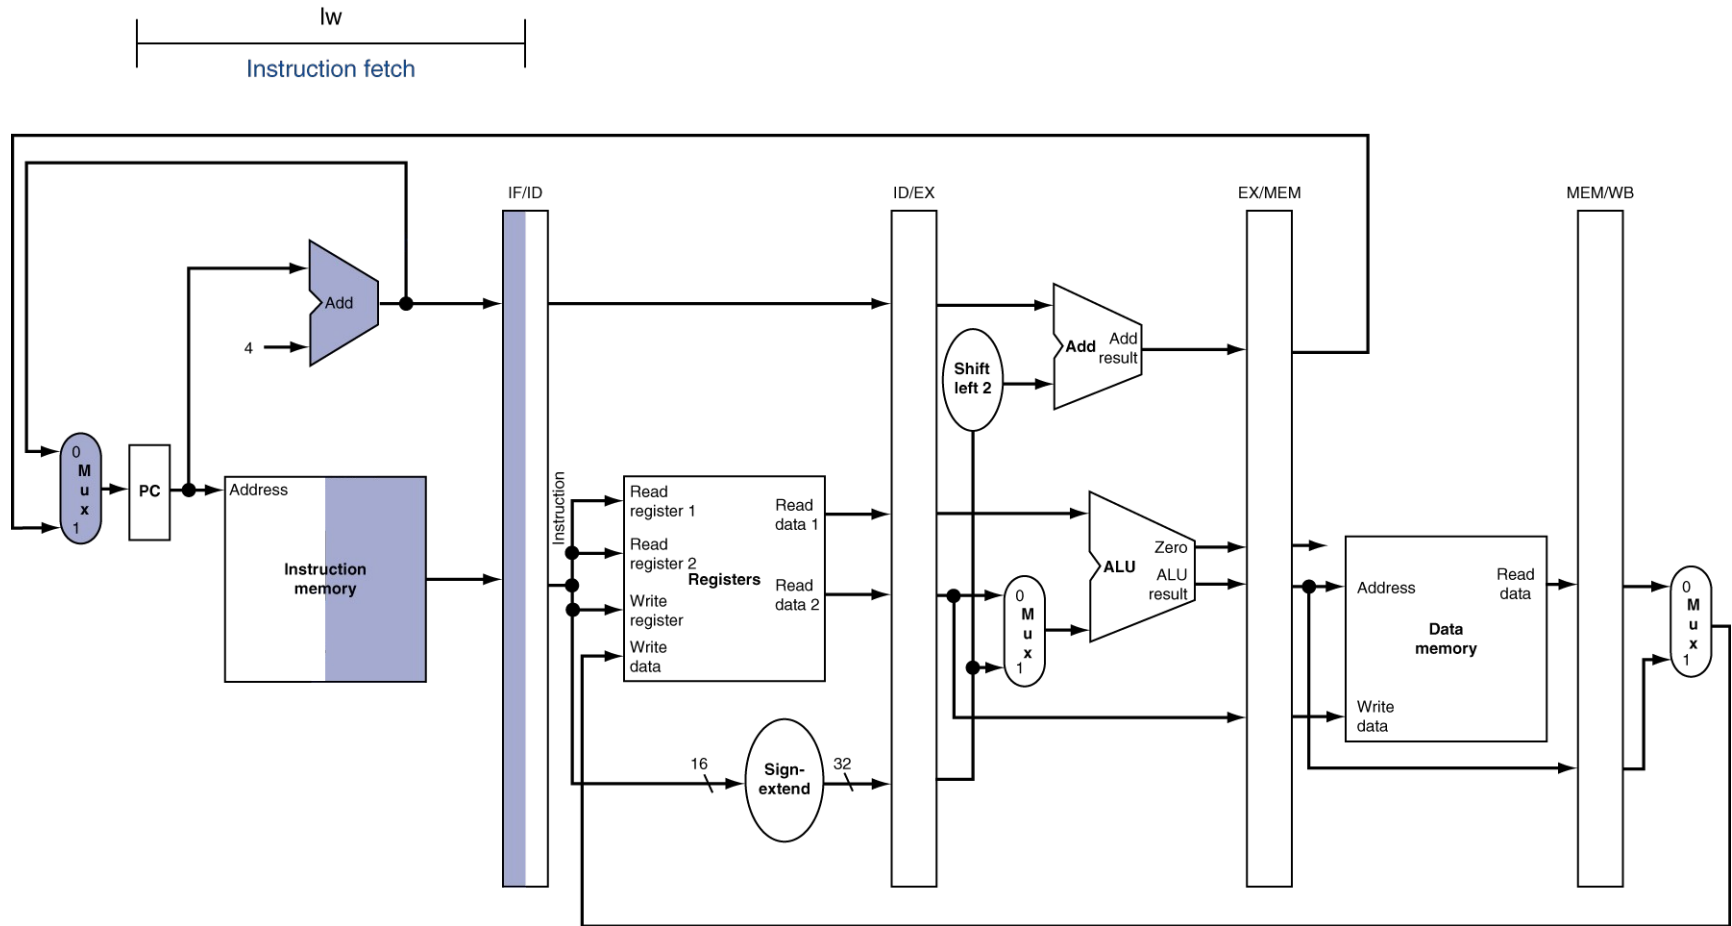




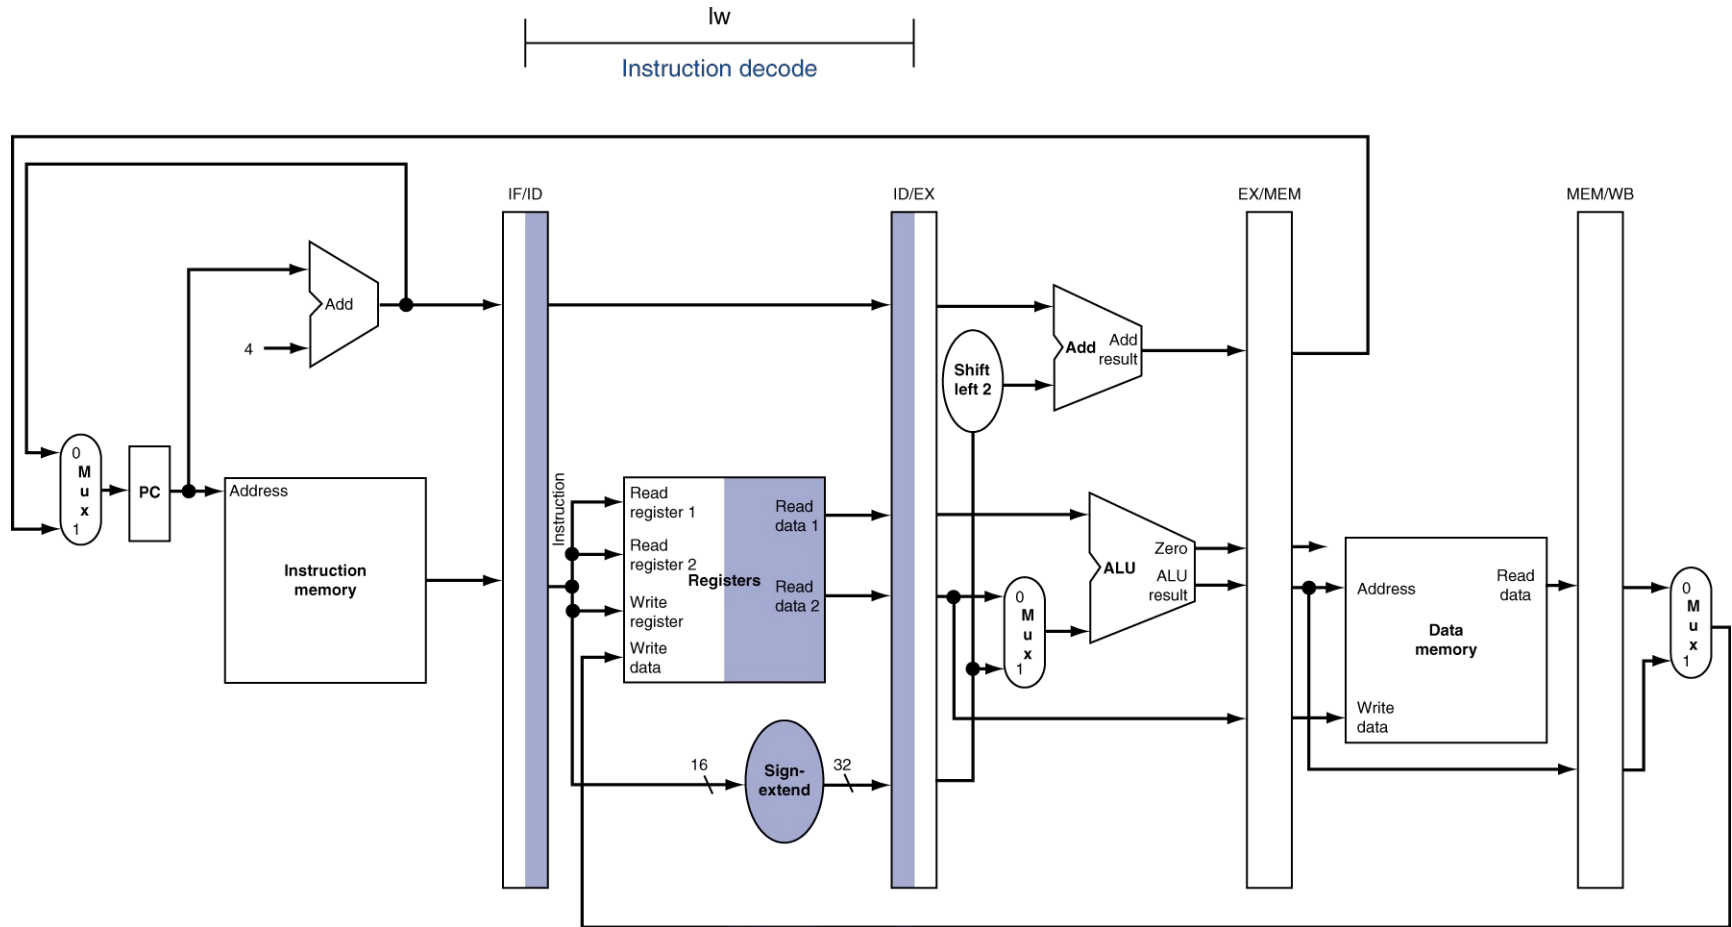
# Pipeline Operation

- Cycle-by-cycle flow of instructions through the pipelined datapath
  - “single-clock-cycle” pipeline diagram
    - shows pipeline usage in a single cycle: “snapshot”
    - highlight **resources** used
  - vs. “multi-clock-cycle” diagram
    - shows operation **over time**
- next: look at “single-clock-cycle” diagrams for load and store

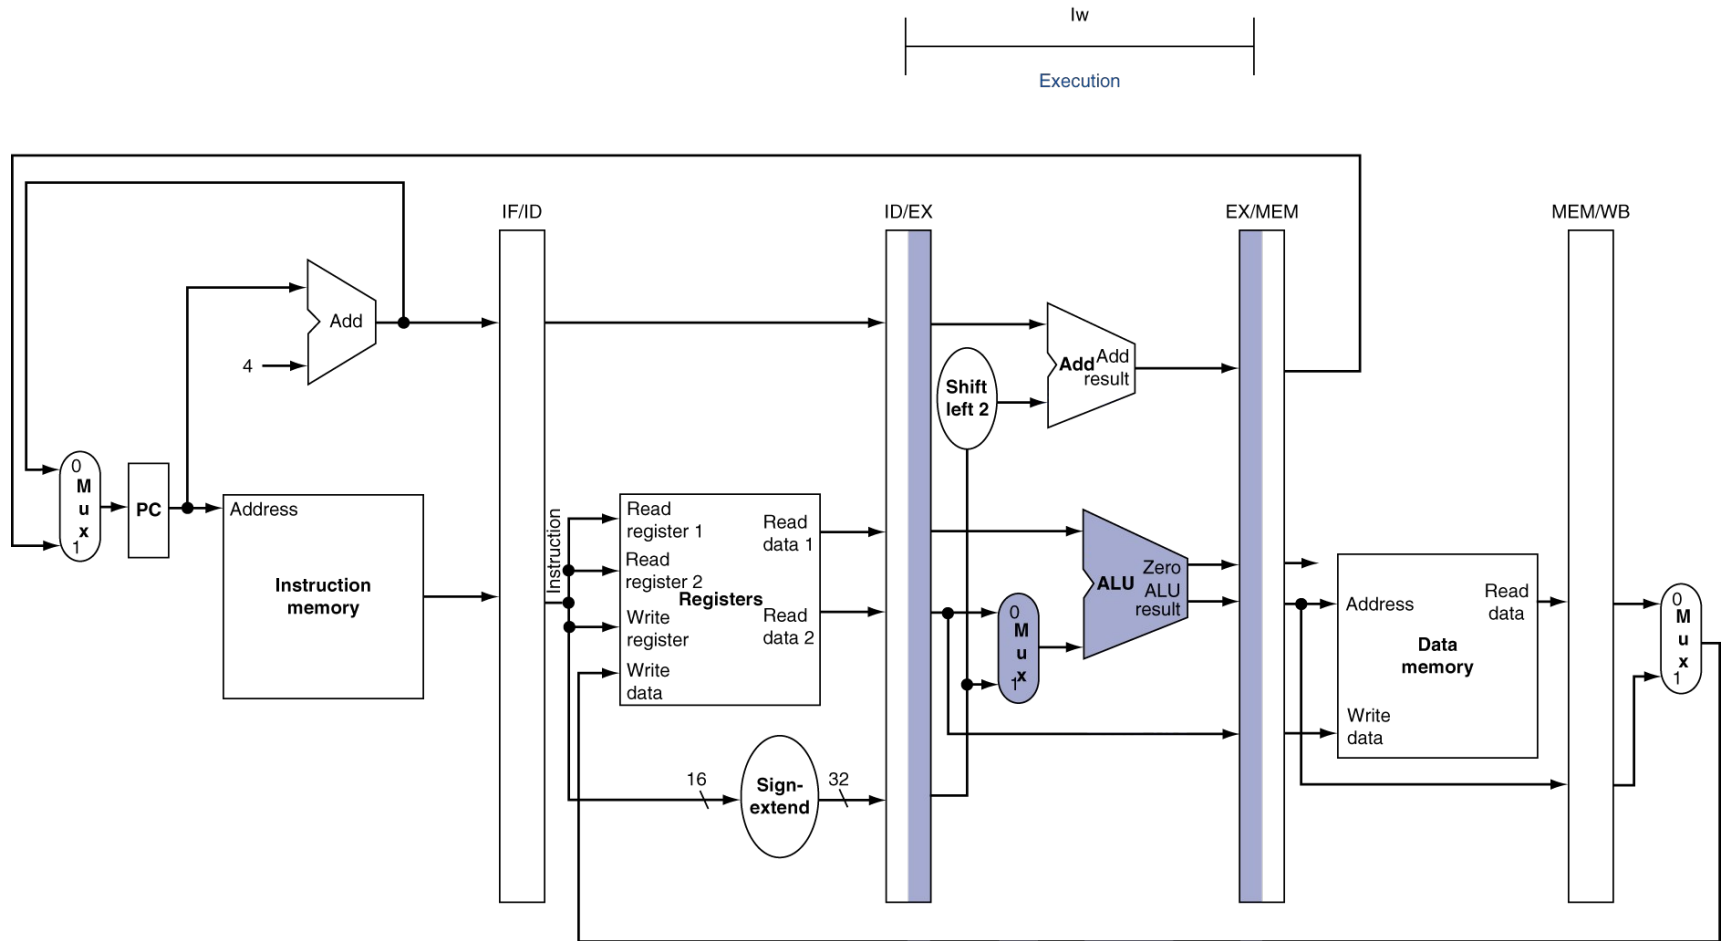
# IF for Load, Store, ...



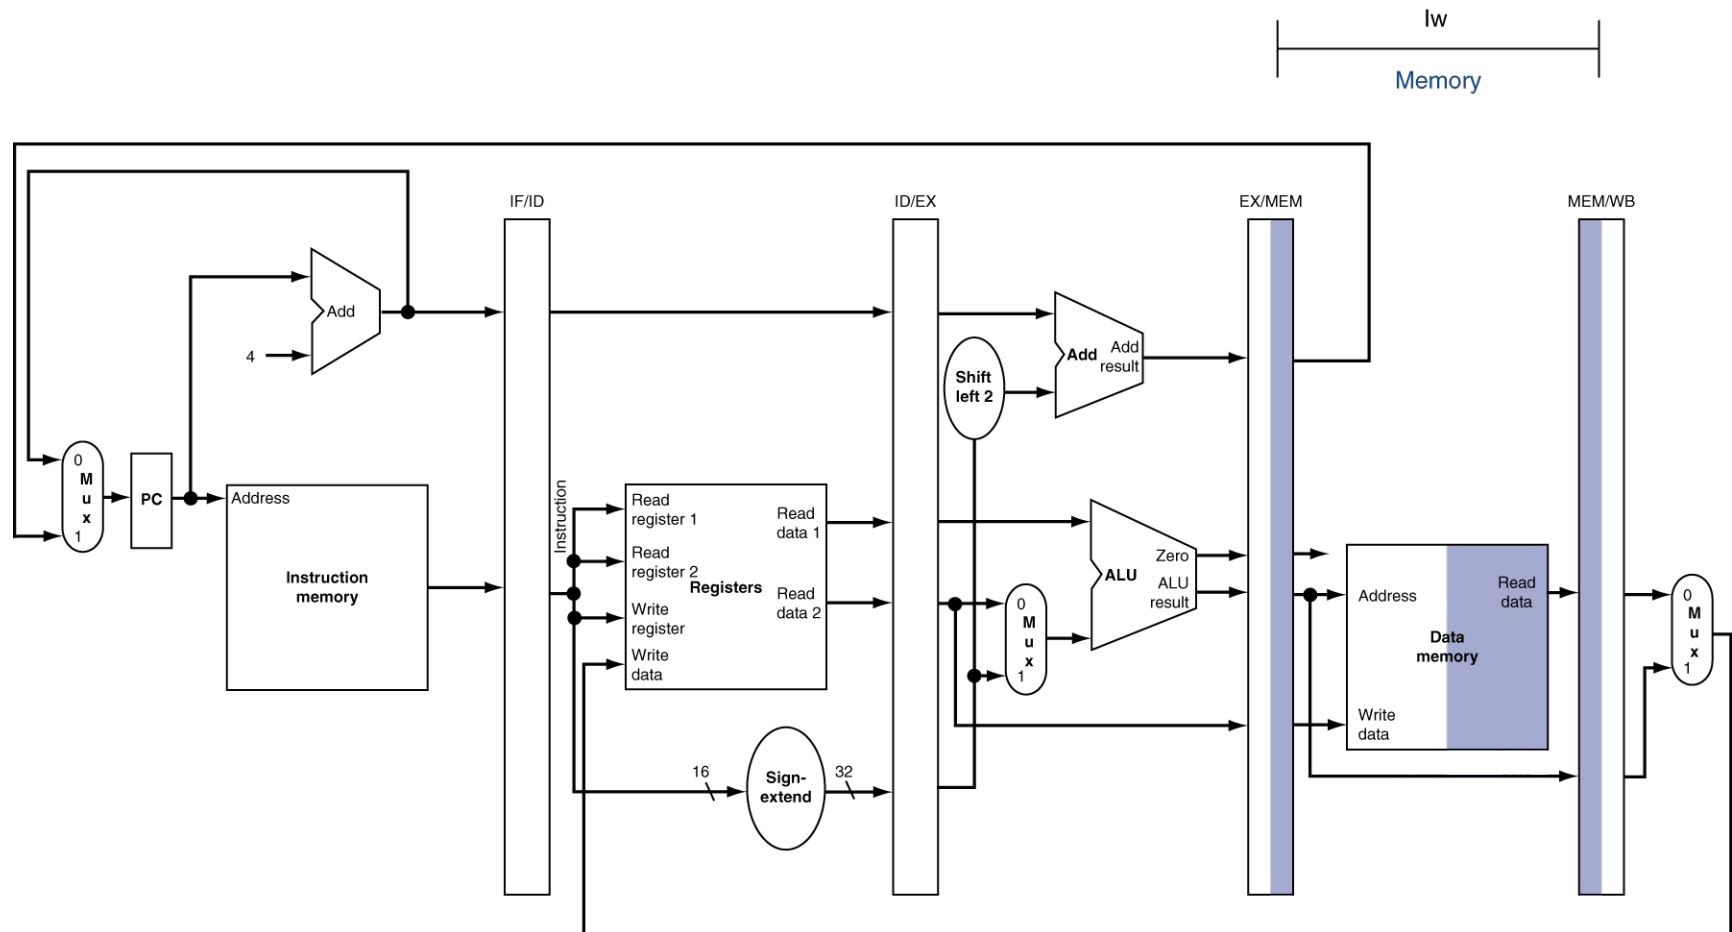
# ID for Load, Store, ...



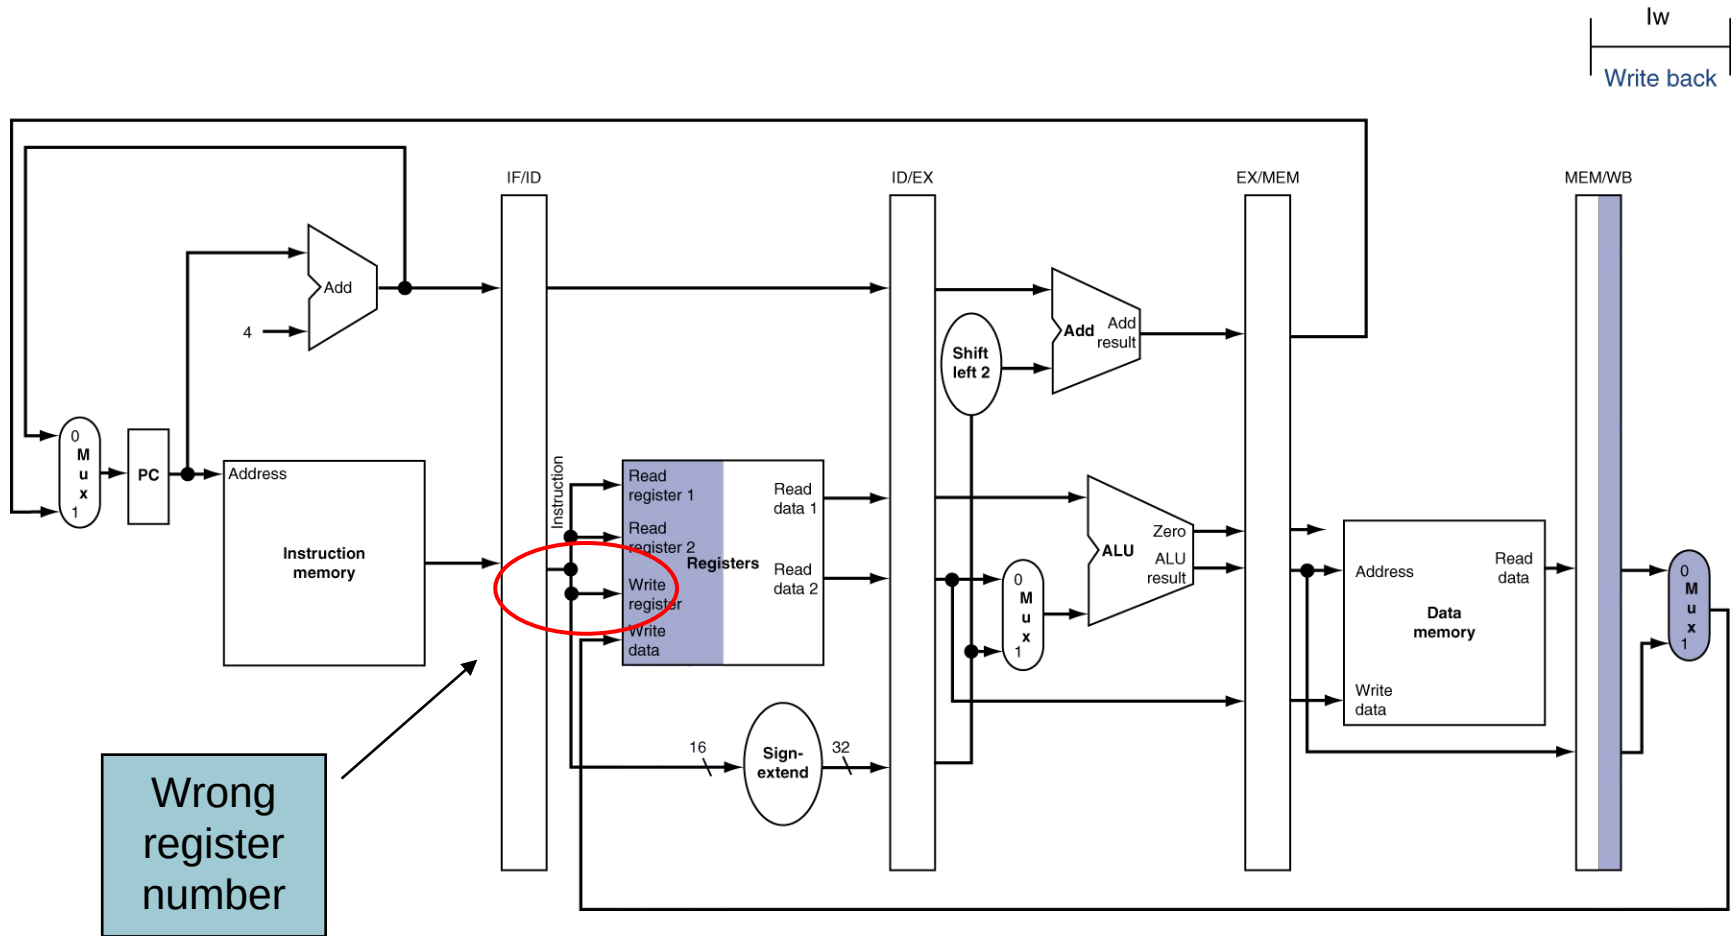
# EX for Load



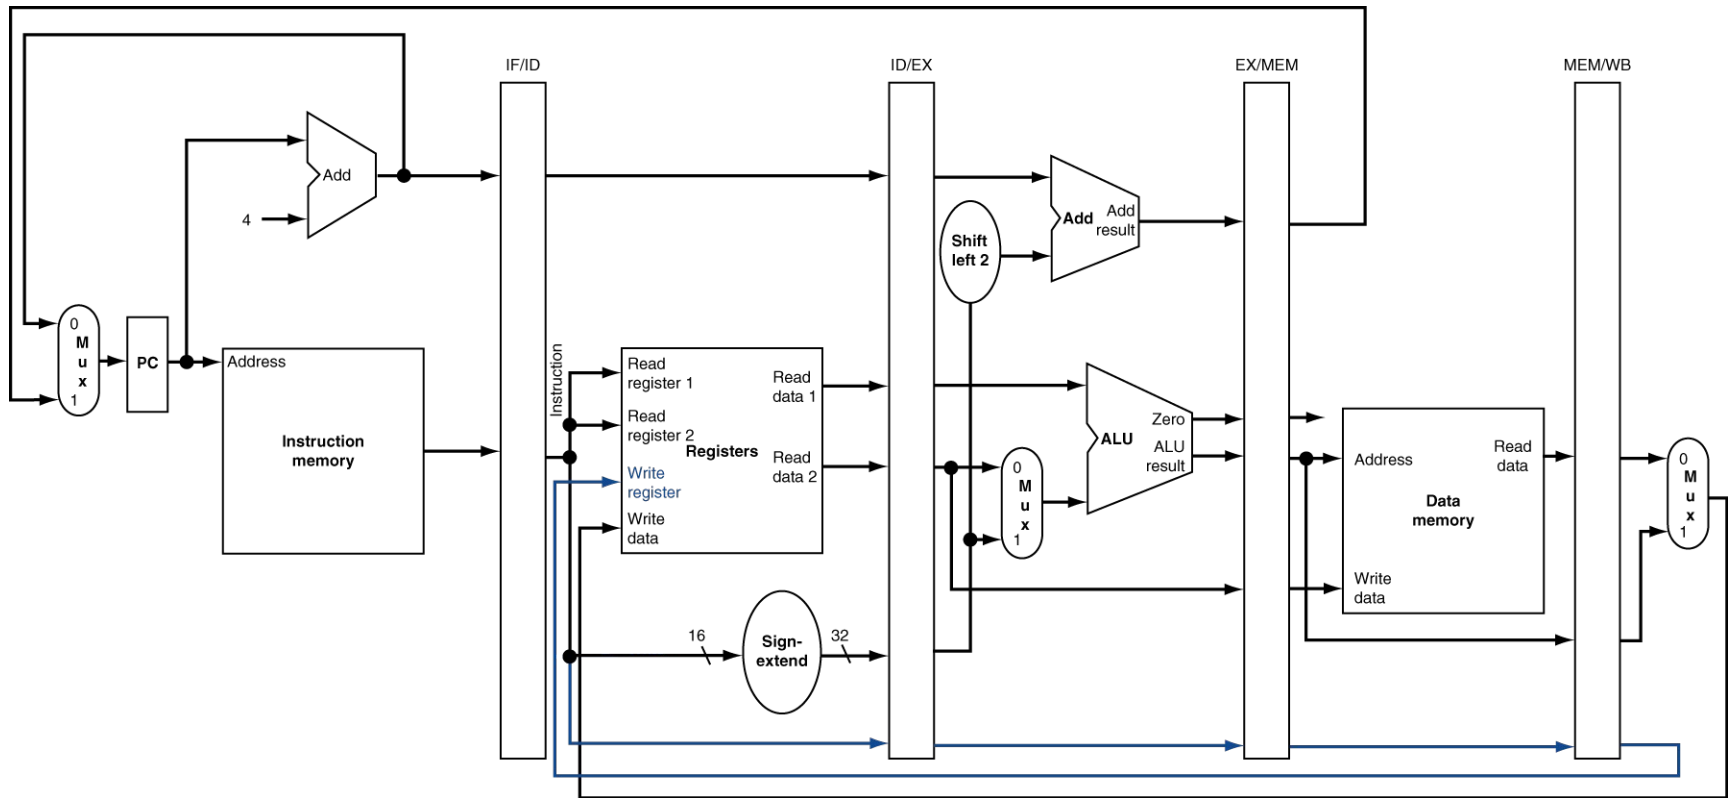
# MEM for Load



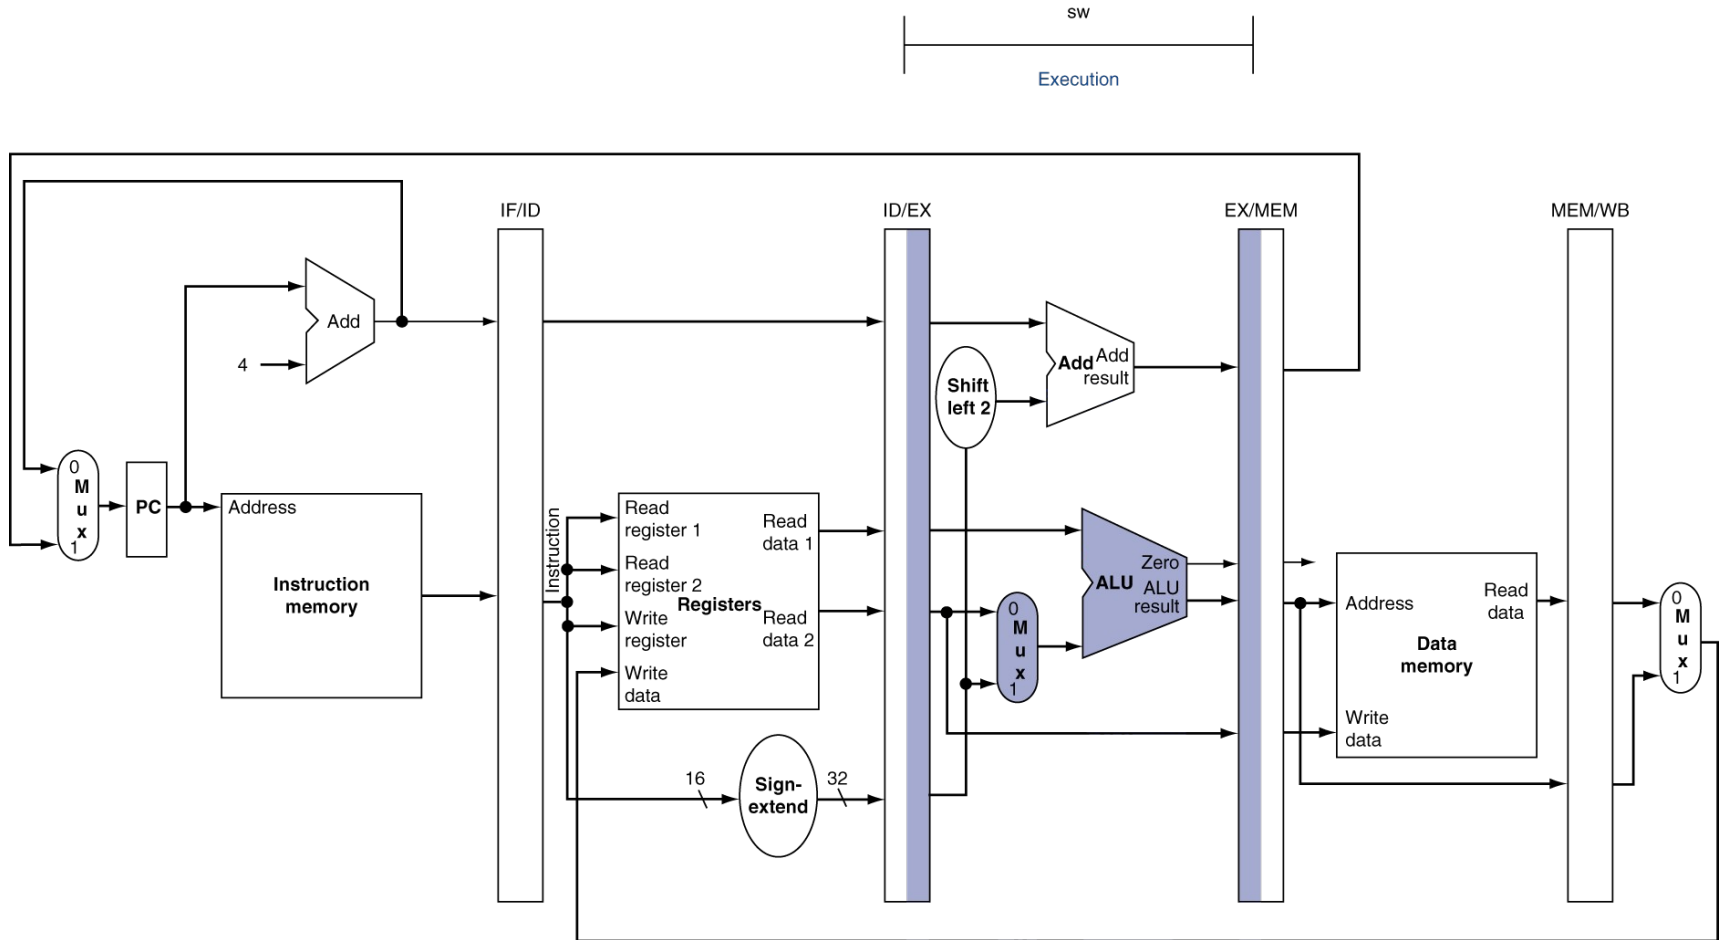
# WB for Load



# Corrected Datapath for Load

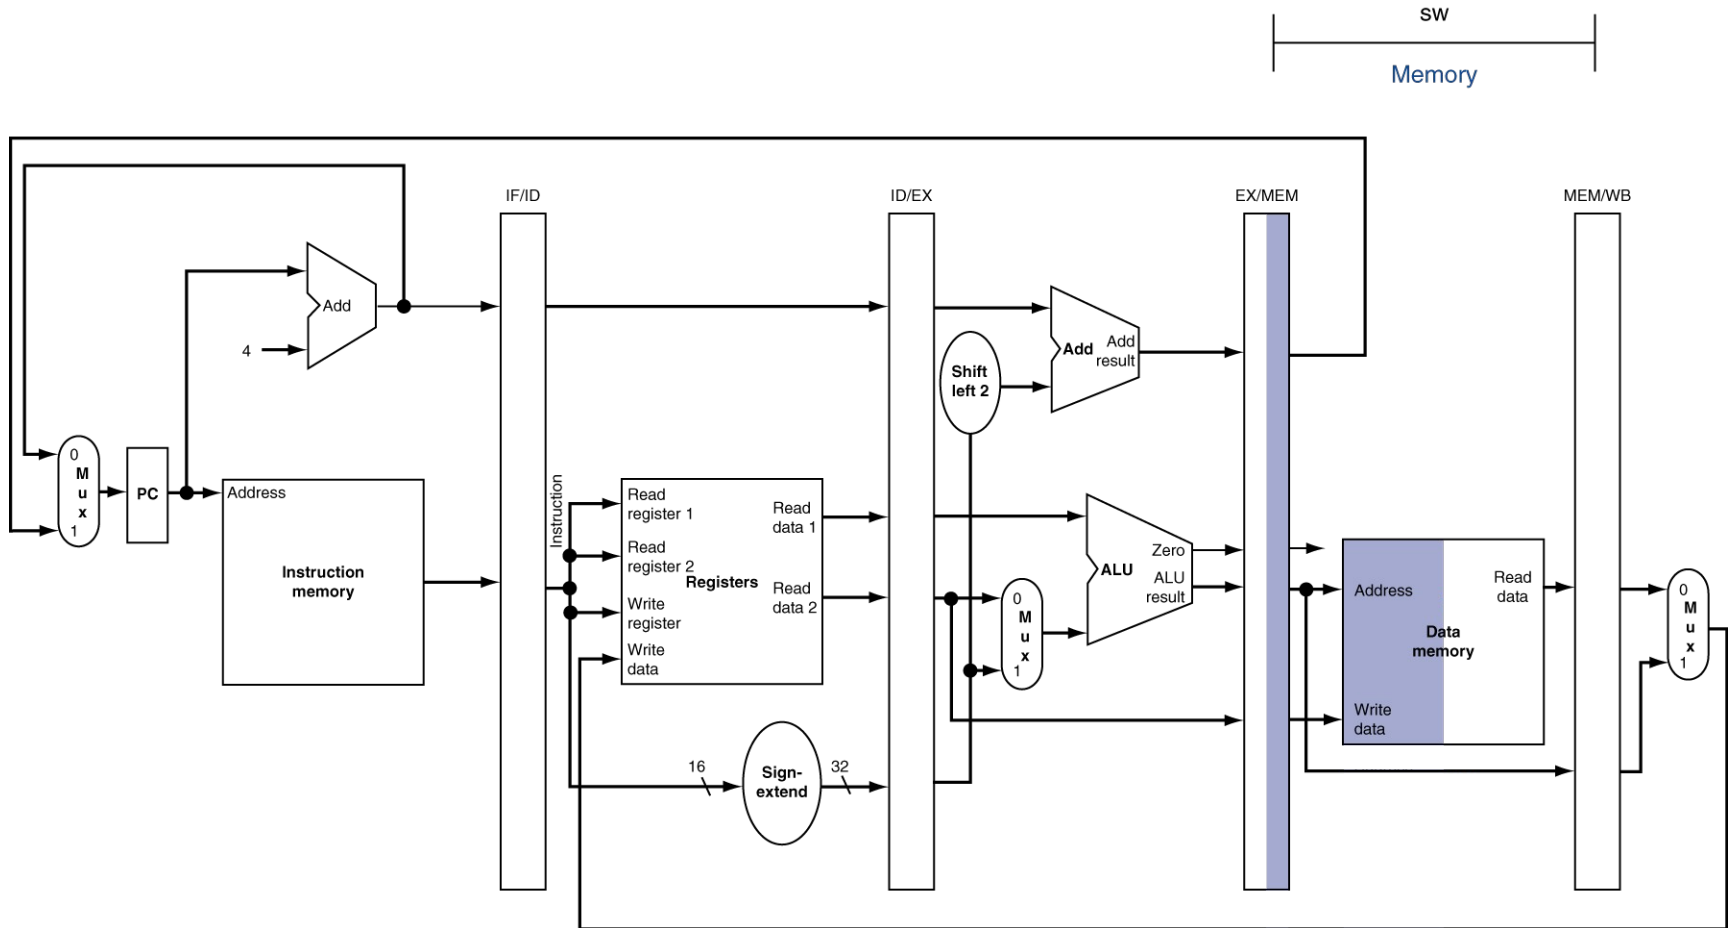


# EX for Store

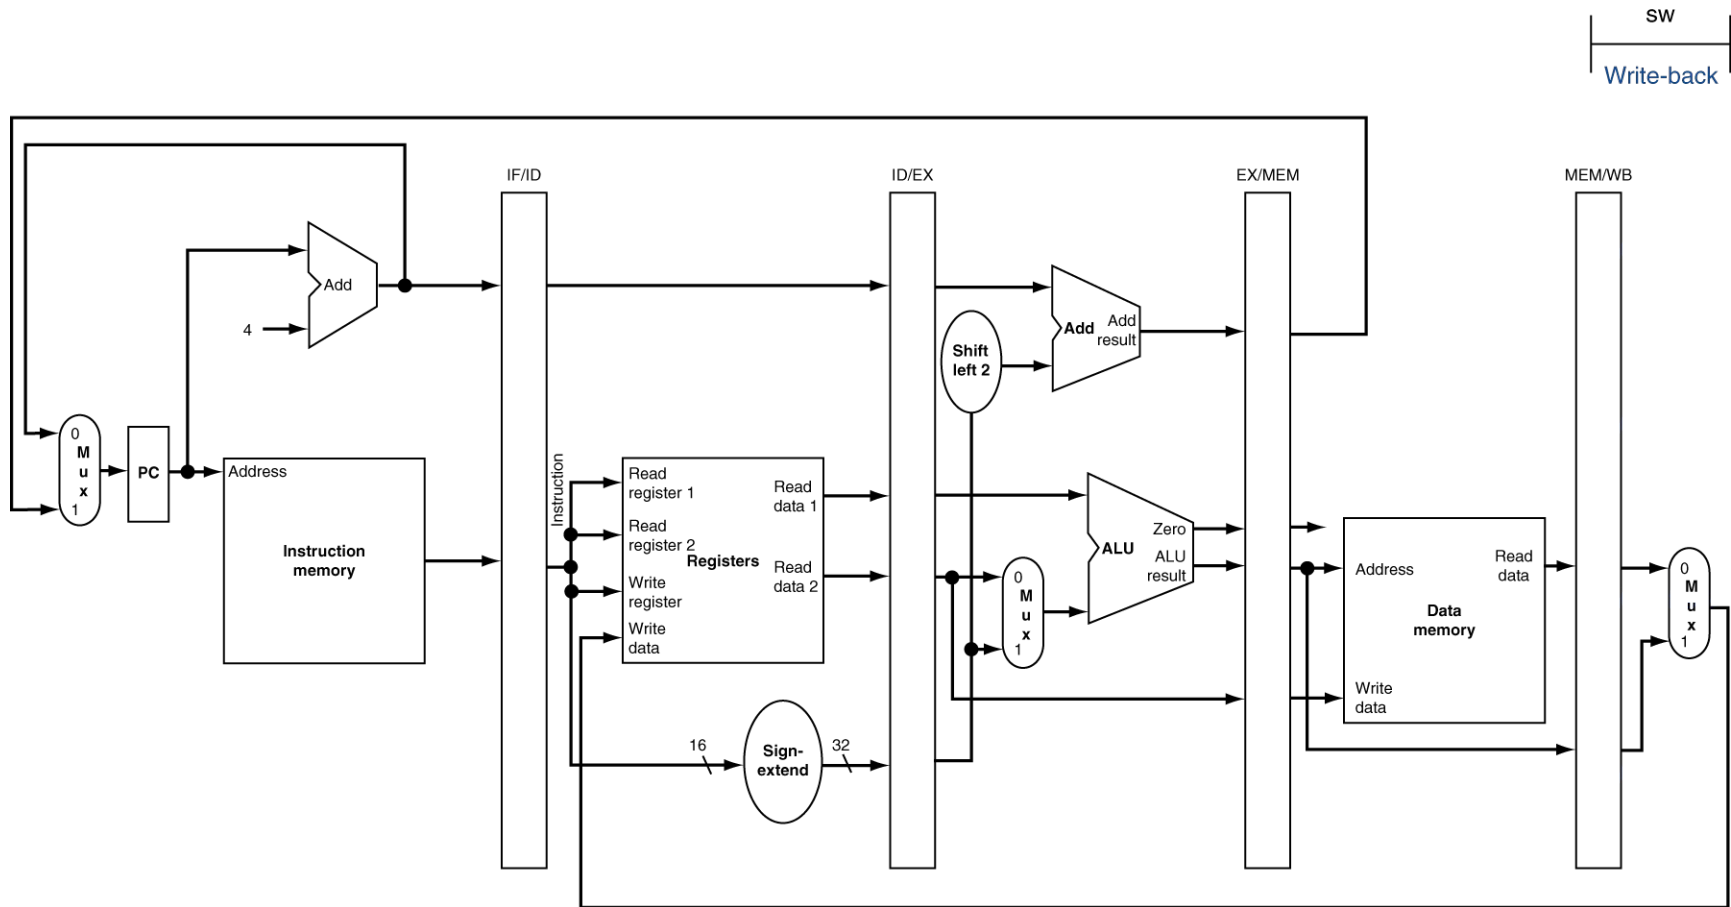




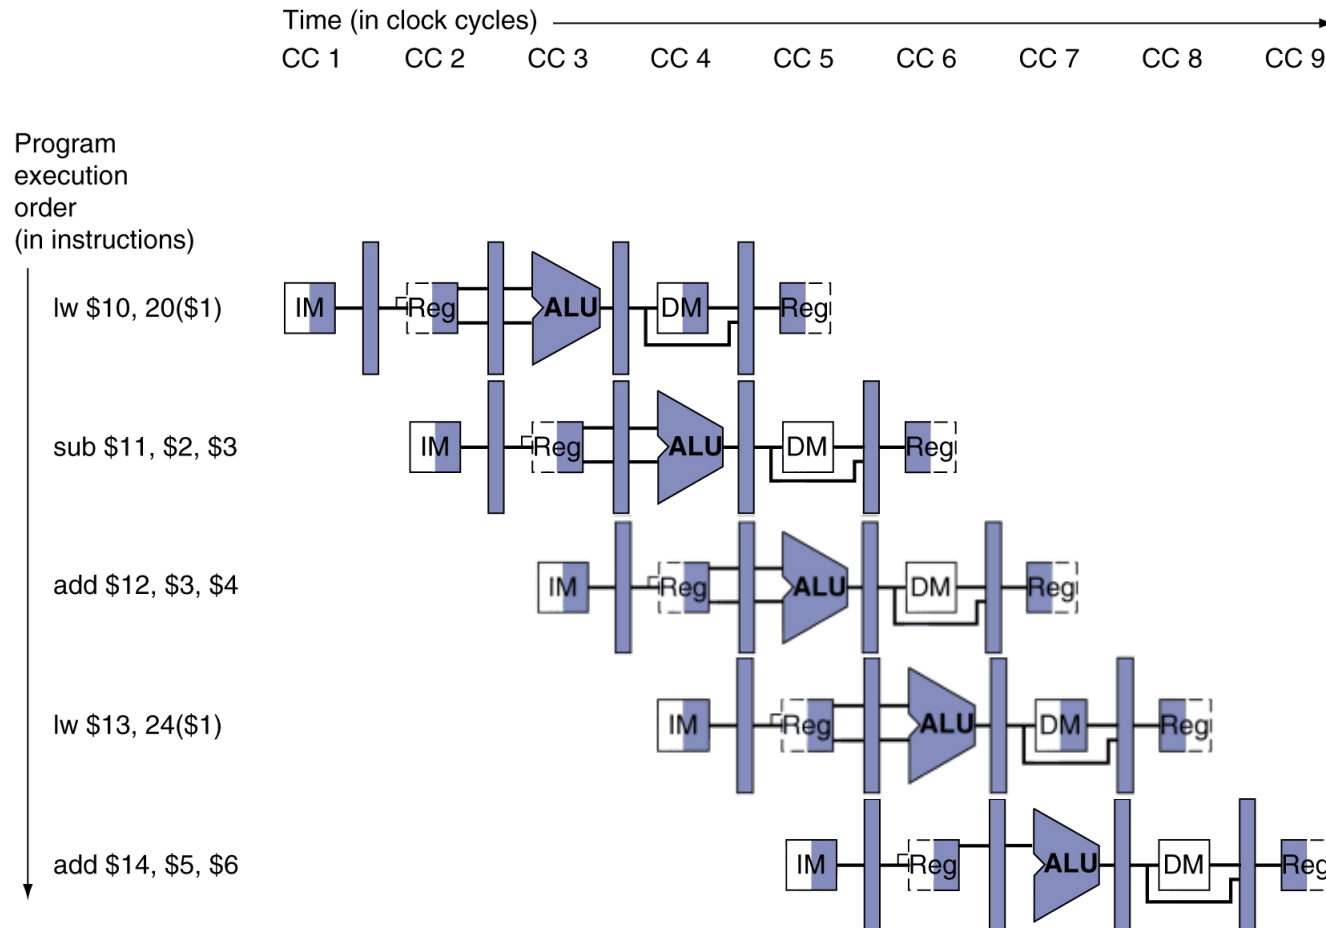
# MEM for Store



# WB for Store




# Multi-Cycle Pipeline Diagram



# Multi-Cycle Pipeline Diagram

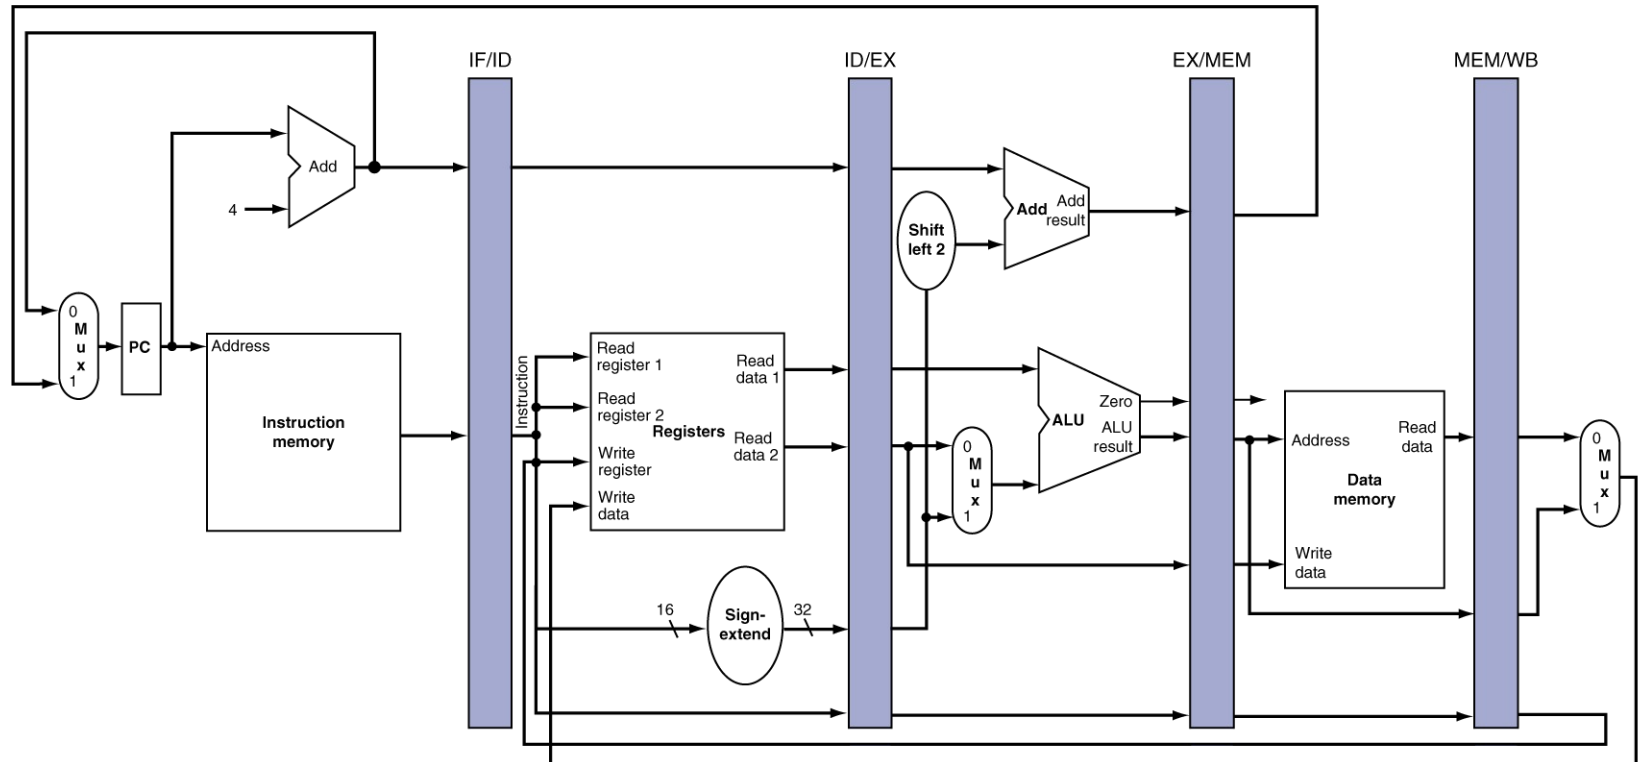
traditional form



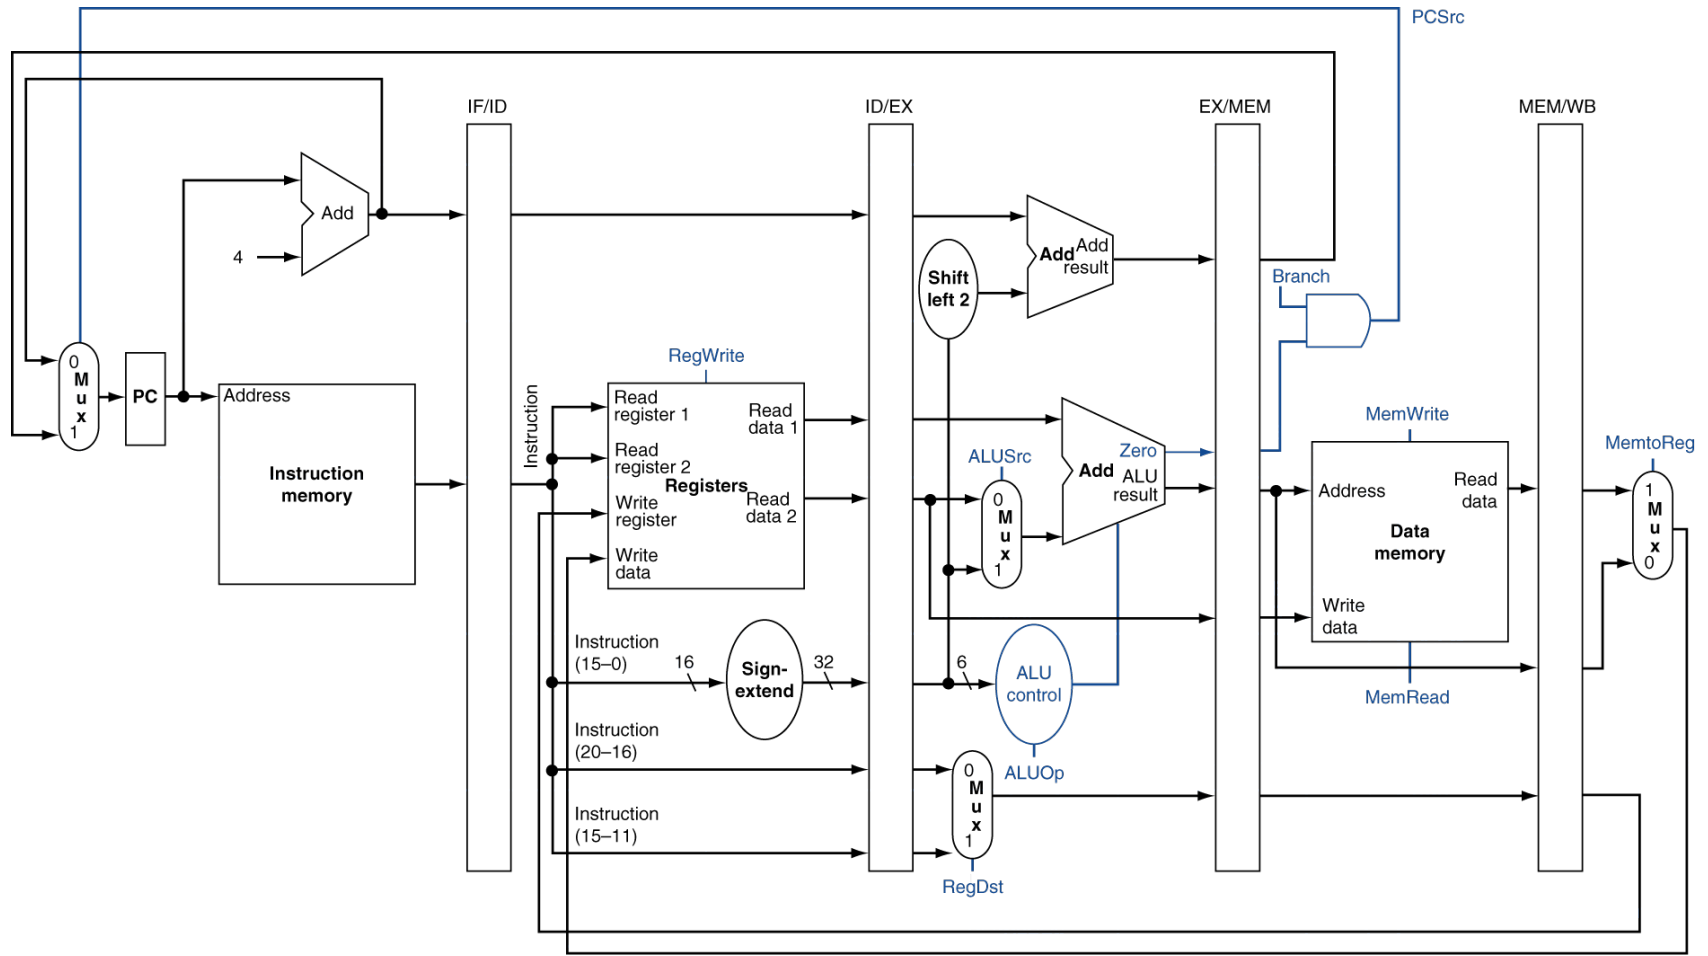
# Single-Cycle Pipeline Diagram

state of pipeline in a given cycle

add \$14, \$5, \$6	lw \$13, 24 (\$1)	add \$12, \$3, \$4	sub \$11, \$2, \$3	lw \$10, 20(\$1)
Instruction fetch	Instruction decode	Execution	Memory	Write-back




# Pipelined Control (simplified)

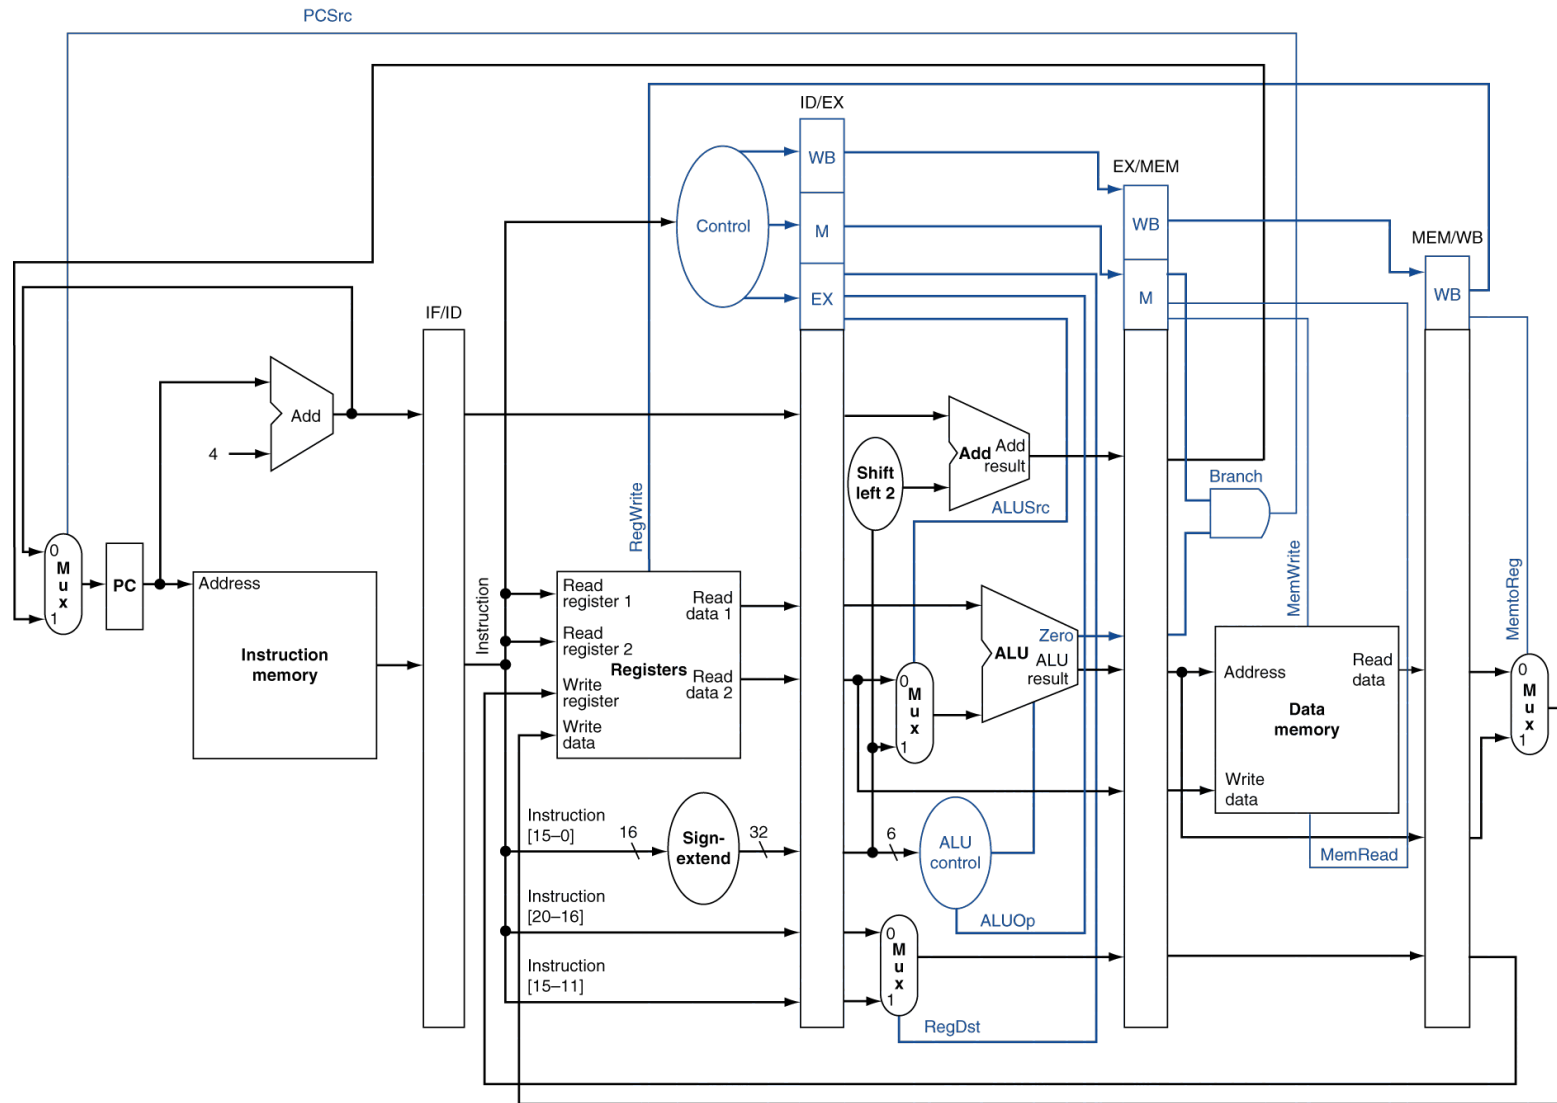


# Pipelined Control

control signals derived from instruction  
as in single-cycle implementation



# Pipelined Control





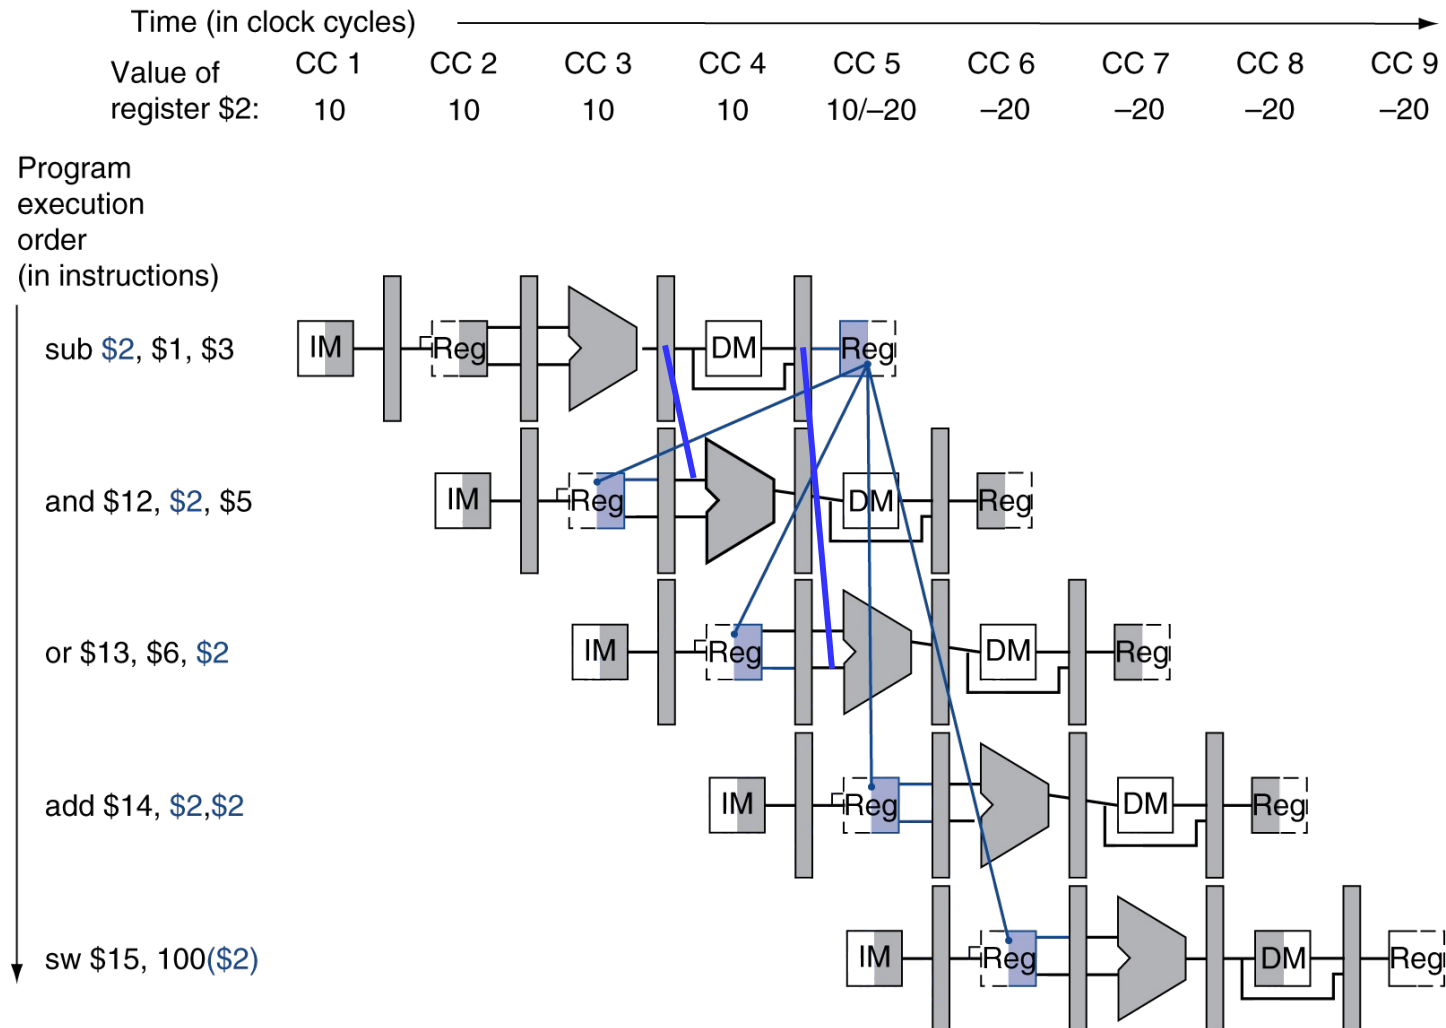
# Data Hazards in ALU Instructions

- Consider this sequence:

```
sub  $2, $1, $3
and  $12, $2, $5
or   $13, $6, $2
add  $14, $2, $2
sw   $15, 100($2)
```

- We can **resolve** hazards with **forwarding**
  - How do we **detect when** to forward?

# Dependencies & Forwarding



# Detecting the Need to Forward

- Pass register numbers along pipeline
  - *e.g.*, `ID/EX.RegisterRs` = register number for Rs sitting in ID/EX pipeline register
- ALU operand register numbers in EX stage are given by
  - `ID/EX.RegisterRs`, `ID/EX.RegisterRt`

- Data hazards when

1a. `EX/MEM.RegisterRd` = `ID/EX.RegisterRs`

1b. `EX/MEM.RegisterRd` = `ID/EX.RegisterRt`

2a. `MEM/WB.RegisterRd` = `ID/EX.RegisterRs`

2b. `MEM/WB.RegisterRd` = `ID/EX.RegisterRt`

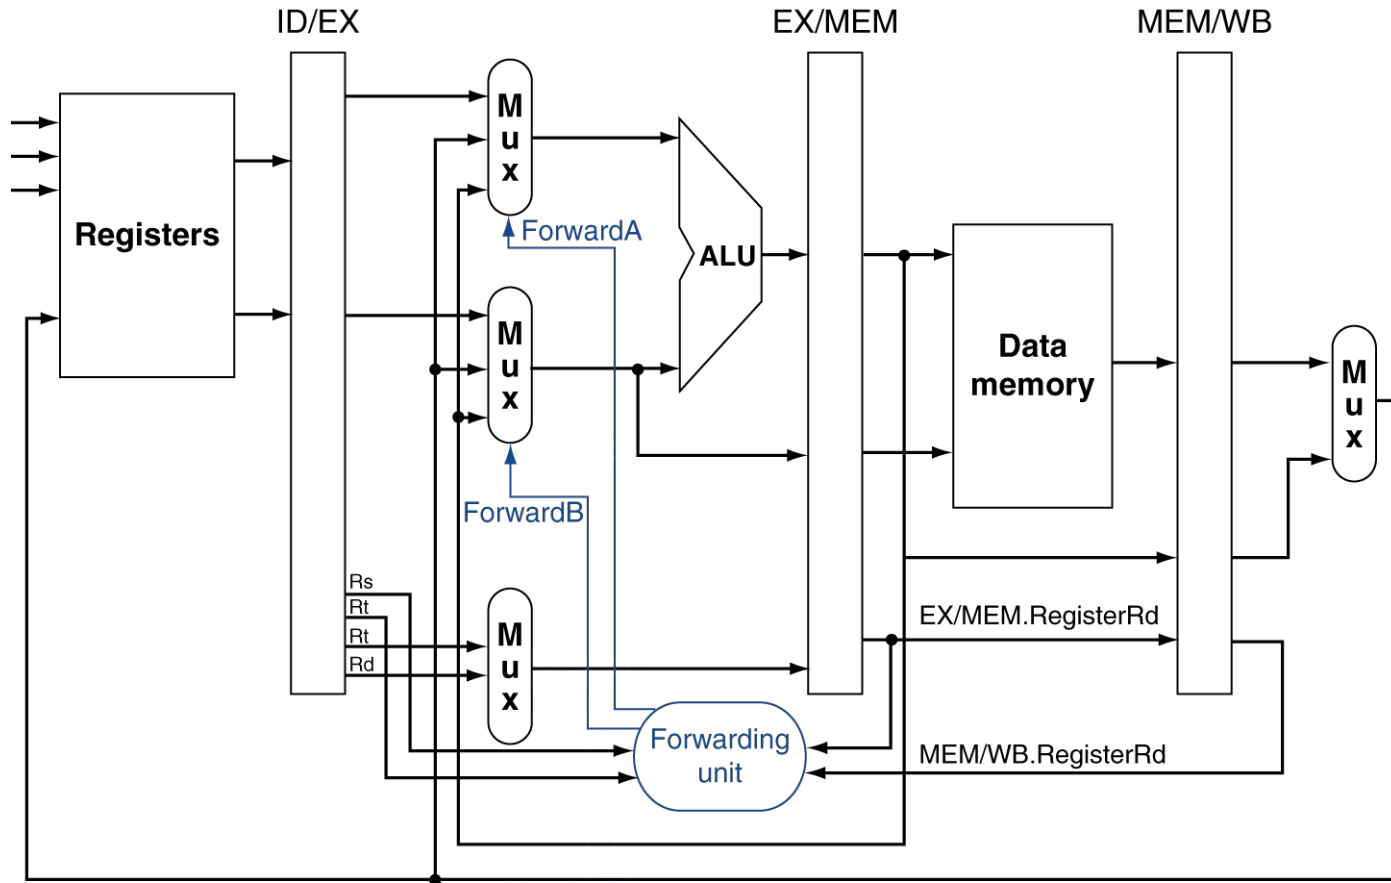
Fwd from  
EX/MEM  
pipeline reg

Fwd from  
MEM/WB  
pipeline reg

# Detecting the Need to Forward

- but **only** if forwarding instruction will **write** to a **register** (e.g., not for `sw`)!
  - `EX/MEM.RegWrite, MEM/WB.RegWrite`
- and only if **Rd** for that instruction is **not \$zero**
  - `EX/MEM.RegisterRd ≠ 0,`  
`MEM/WB.RegisterRd ≠ 0`

# Forwarding Paths



# Forwarding Conditions

- EX hazard
  - if (EX/MEM.RegWrite and (EX/MEM.RegisterRd  $\neq$  0)  
and (EX/MEM.RegisterRd = ID/EX.RegisterRs))  
ForwardA = 10
  - if (EX/MEM.RegWrite and (EX/MEM.RegisterRd  $\neq$  0)  
and (EX/MEM.RegisterRd = ID/EX.RegisterRt))  
ForwardB = 10
- MEM hazard
  - if (MEM/WB.RegWrite and (MEM/WB.RegisterRd  $\neq$  0)  
and (MEM/WB.RegisterRd = ID/EX.RegisterRs))  
ForwardA = 01
  - if (MEM/WB.RegWrite and (MEM/WB.RegisterRd  $\neq$  0)  
and (MEM/WB.RegisterRd = ID/EX.RegisterRt))  
ForwardB = 01

# Double Data Hazard

- Consider the sequence:

add \$1, \$1, \$2

add \$1, \$1, \$3

add \$1, \$1, \$4

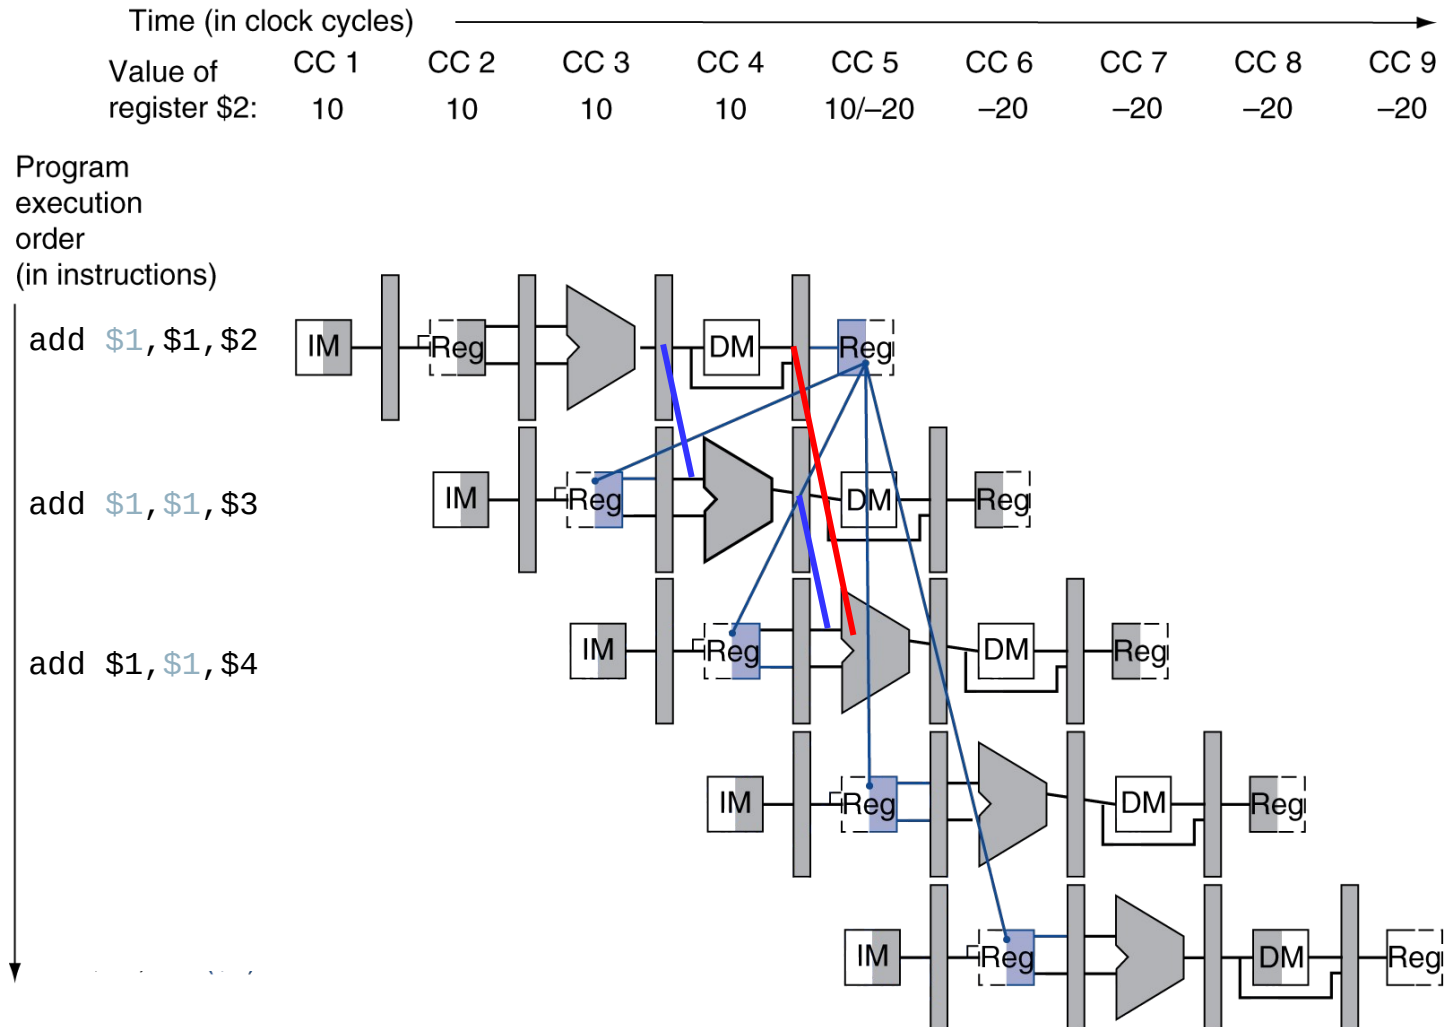
- Both hazards occur

- Want to use the **most recent**

- Revise MEM hazard condition

- Only fwd if EX hazard condition is not true

# Double Data Hazard



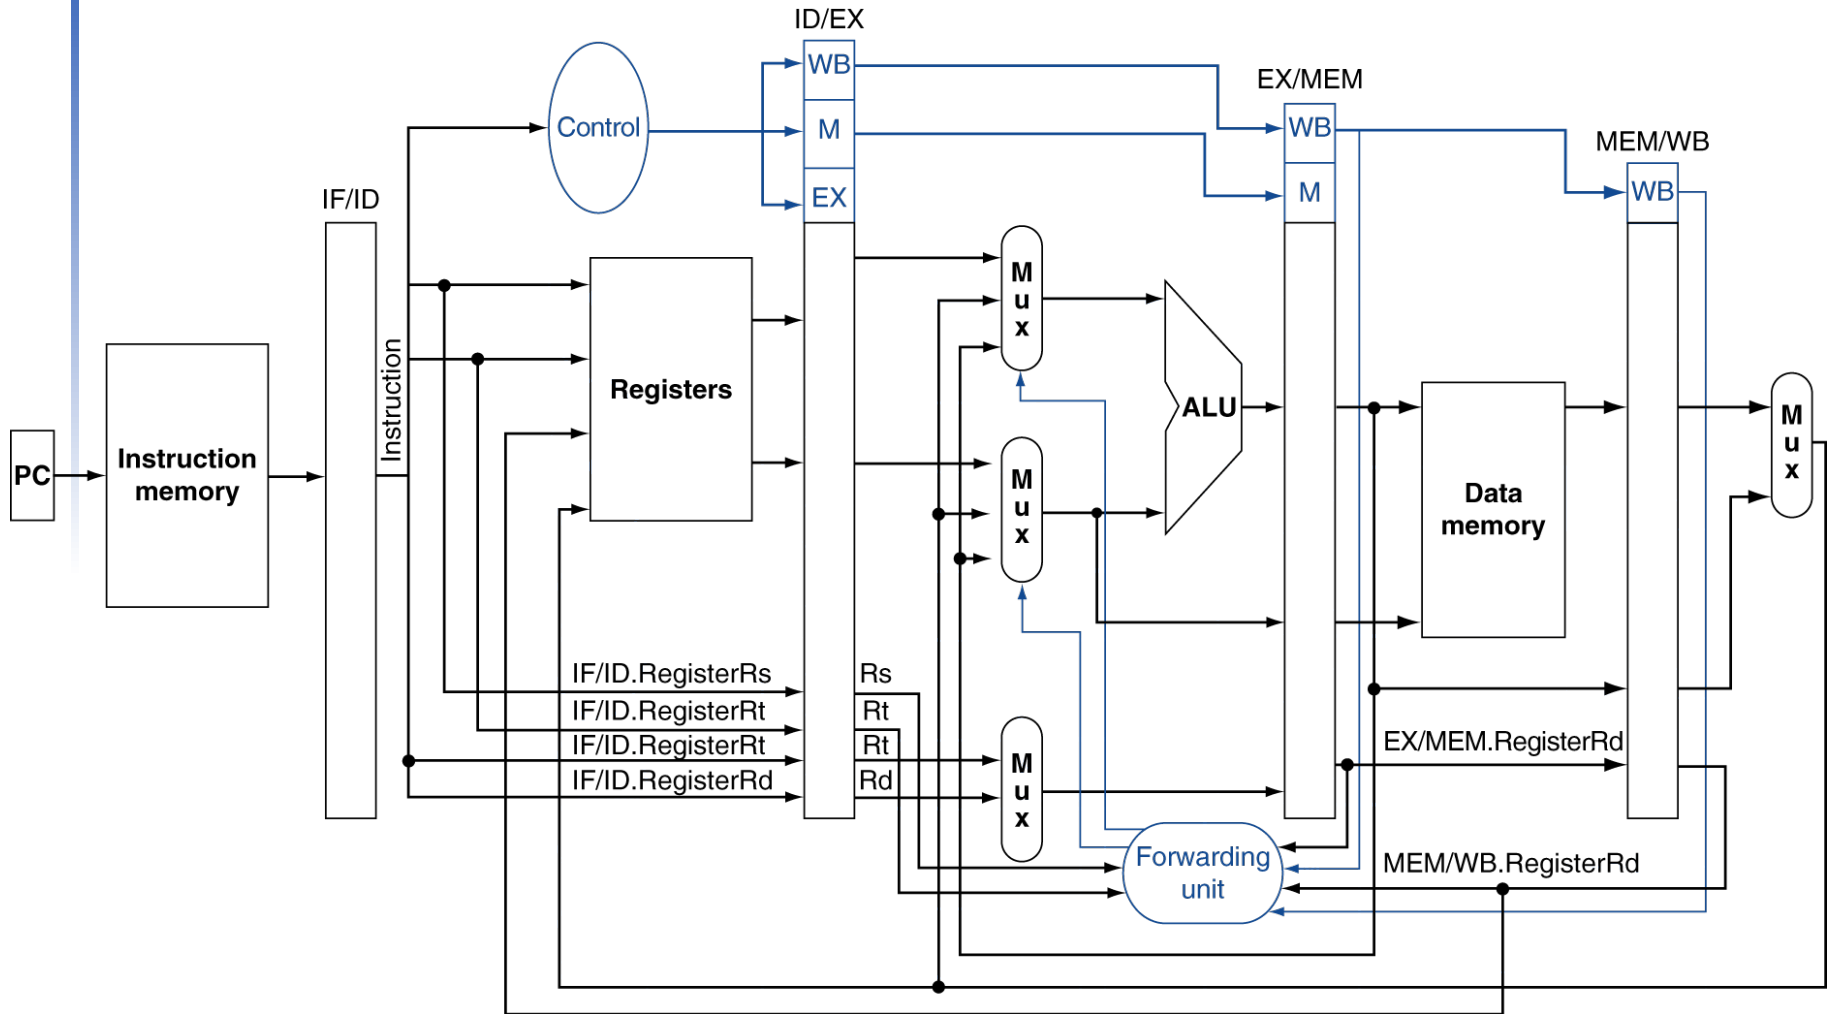


# Revised Forwarding Condition

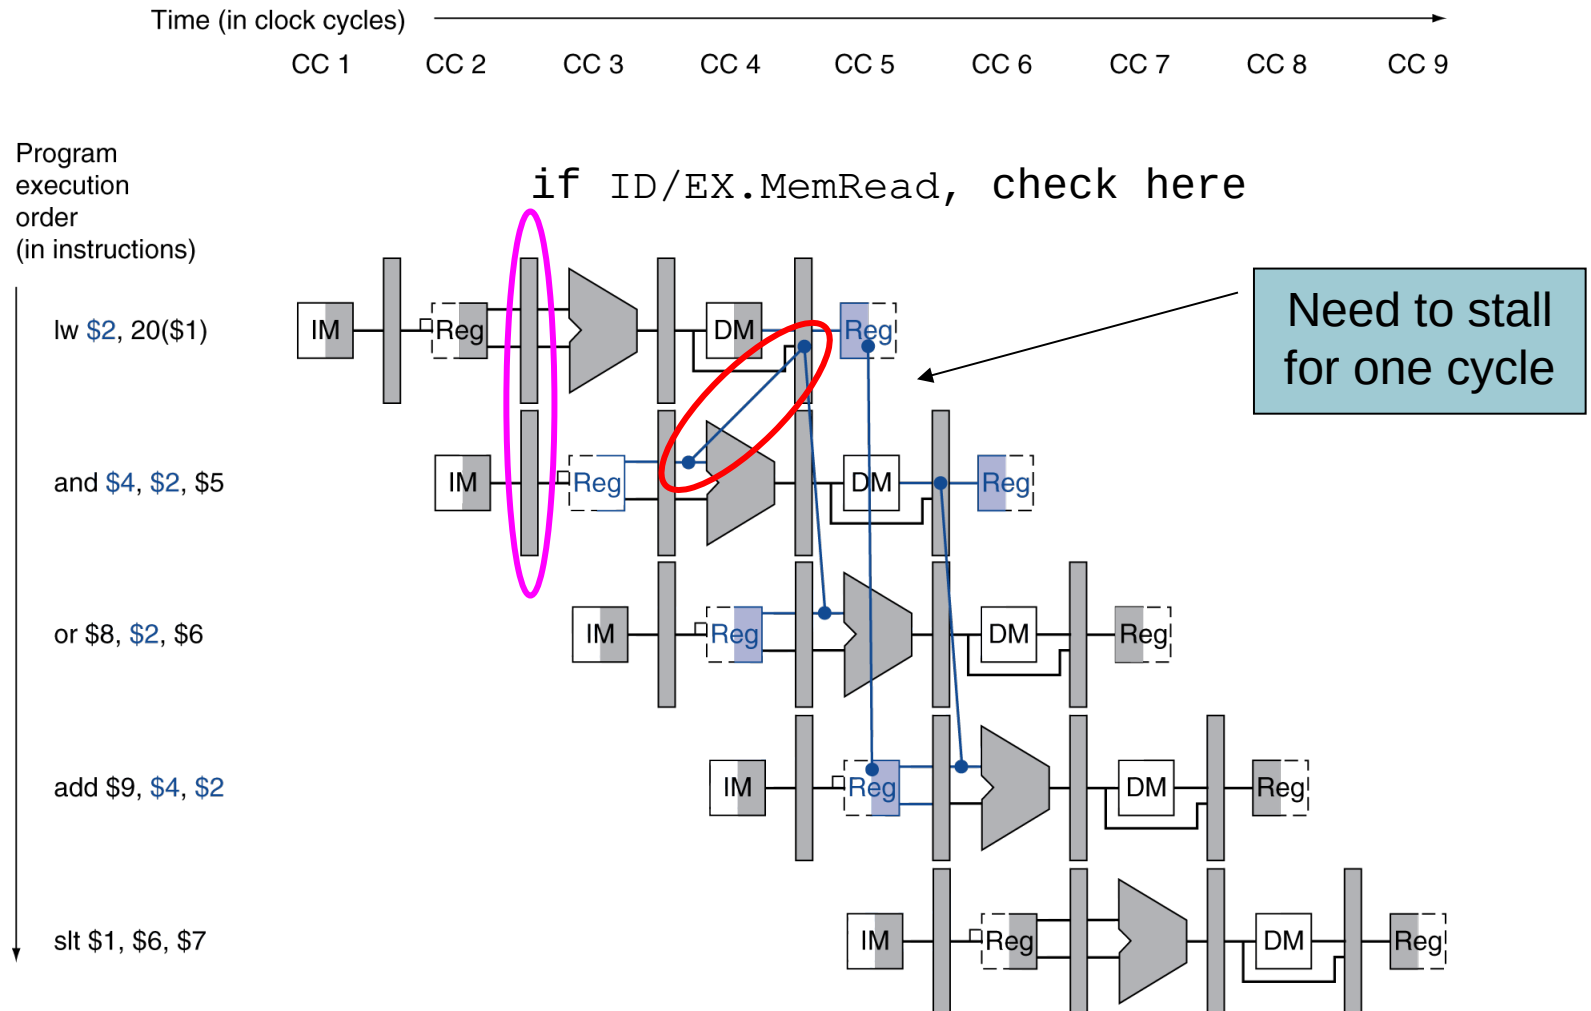
## MEM hazard

- if (MEM/WB.RegWrite and (MEM/WB.RegisterRd  $\neq$  0)  
    and not (EX/MEM.RegWrite and (EX/MEM.RegisterRd  $\neq$  0)  
        and (EX/MEM.RegisterRd = ID/EX.RegisterRs))  
    and (MEM/WB.RegisterRd = ID/EX.RegisterRs))  
    ForwardA = 01
- if (MEM/WB.RegWrite and (MEM/WB.RegisterRd  $\neq$  0)  
    and not (EX/MEM.RegWrite and (EX/MEM.RegisterRd  $\neq$  0)  
        and (EX/MEM.RegisterRd = ID/EX.RegisterRt))  
    and (MEM/WB.RegisterRd = ID/EX.RegisterRt))  
    ForwardB = 01

# Datapath with Forwarding



# Load-Use Data Hazard



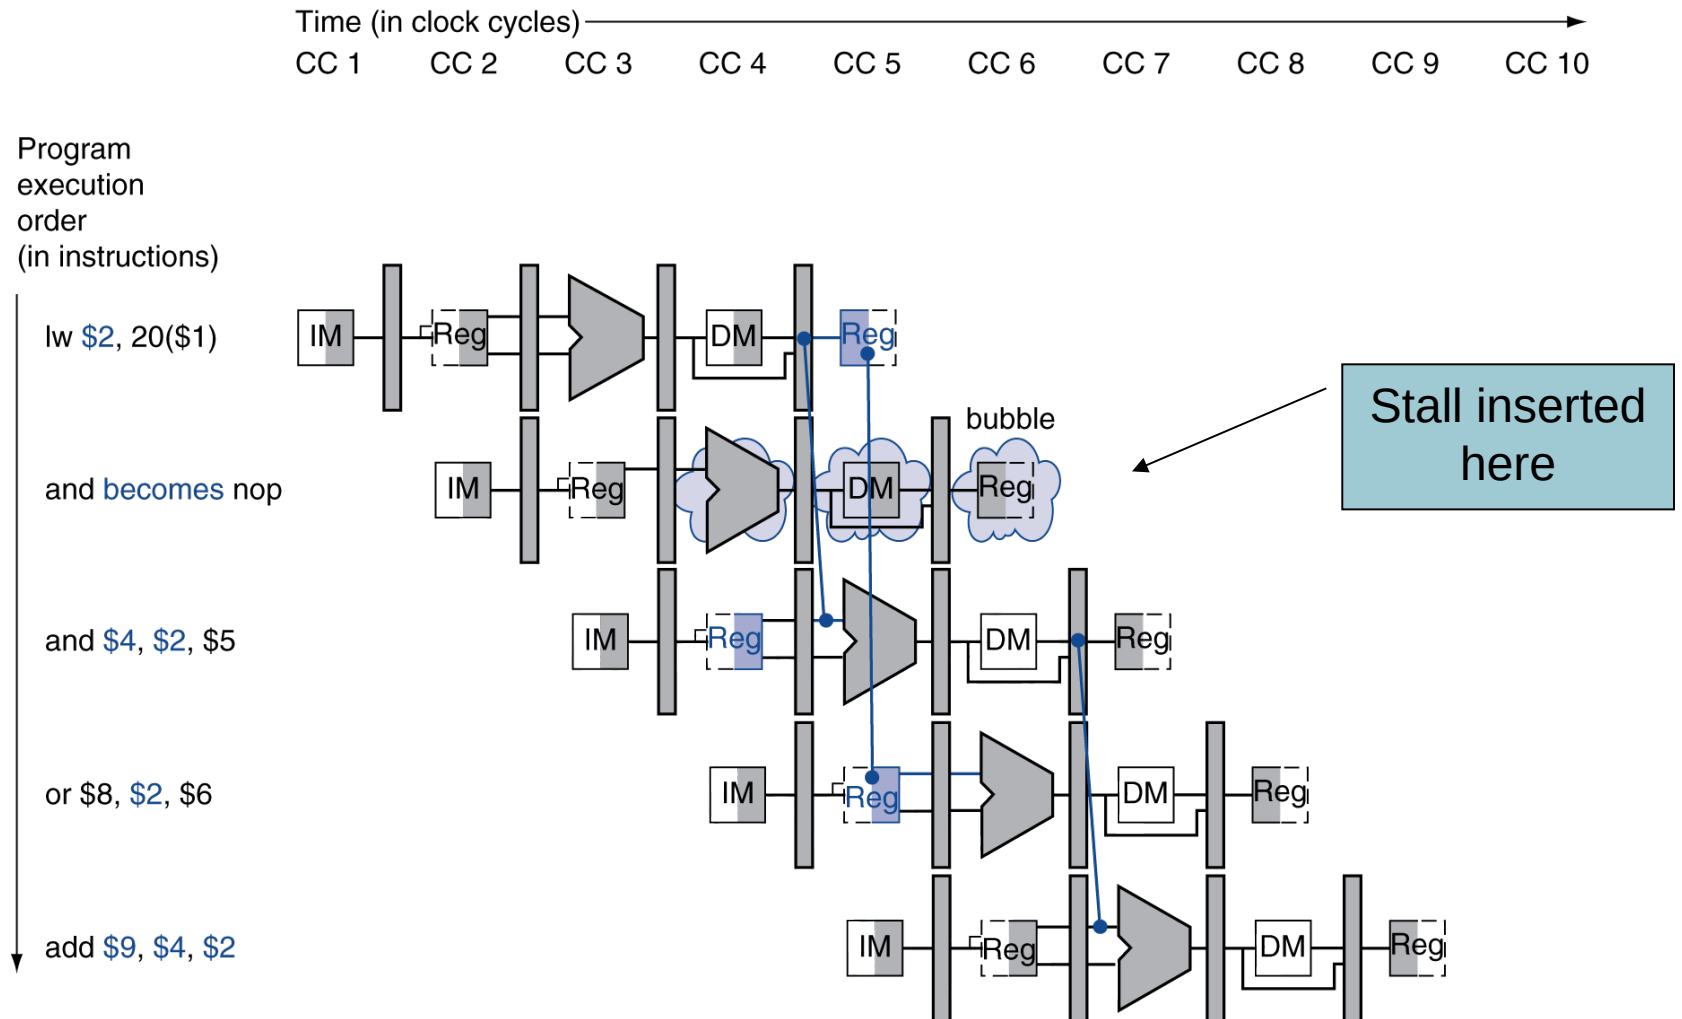
# Load-Use Hazard Detection

- **Check** when using instruction is **decoded** in **ID** stage
- ALU operand register numbers in ID stage are given by
  - `IF/ID.RegisterRs, IF/ID.RegisterRt`
- Load-use hazard when target of load = input of computation
  - `ID/EX.MemRead and  
((ID/EX.RegisterRt = IF/ID.RegisterRs) or  
(ID/EX.RegisterRt = IF/ID.RegisterRt))`
- If detected, stall and insert bubble

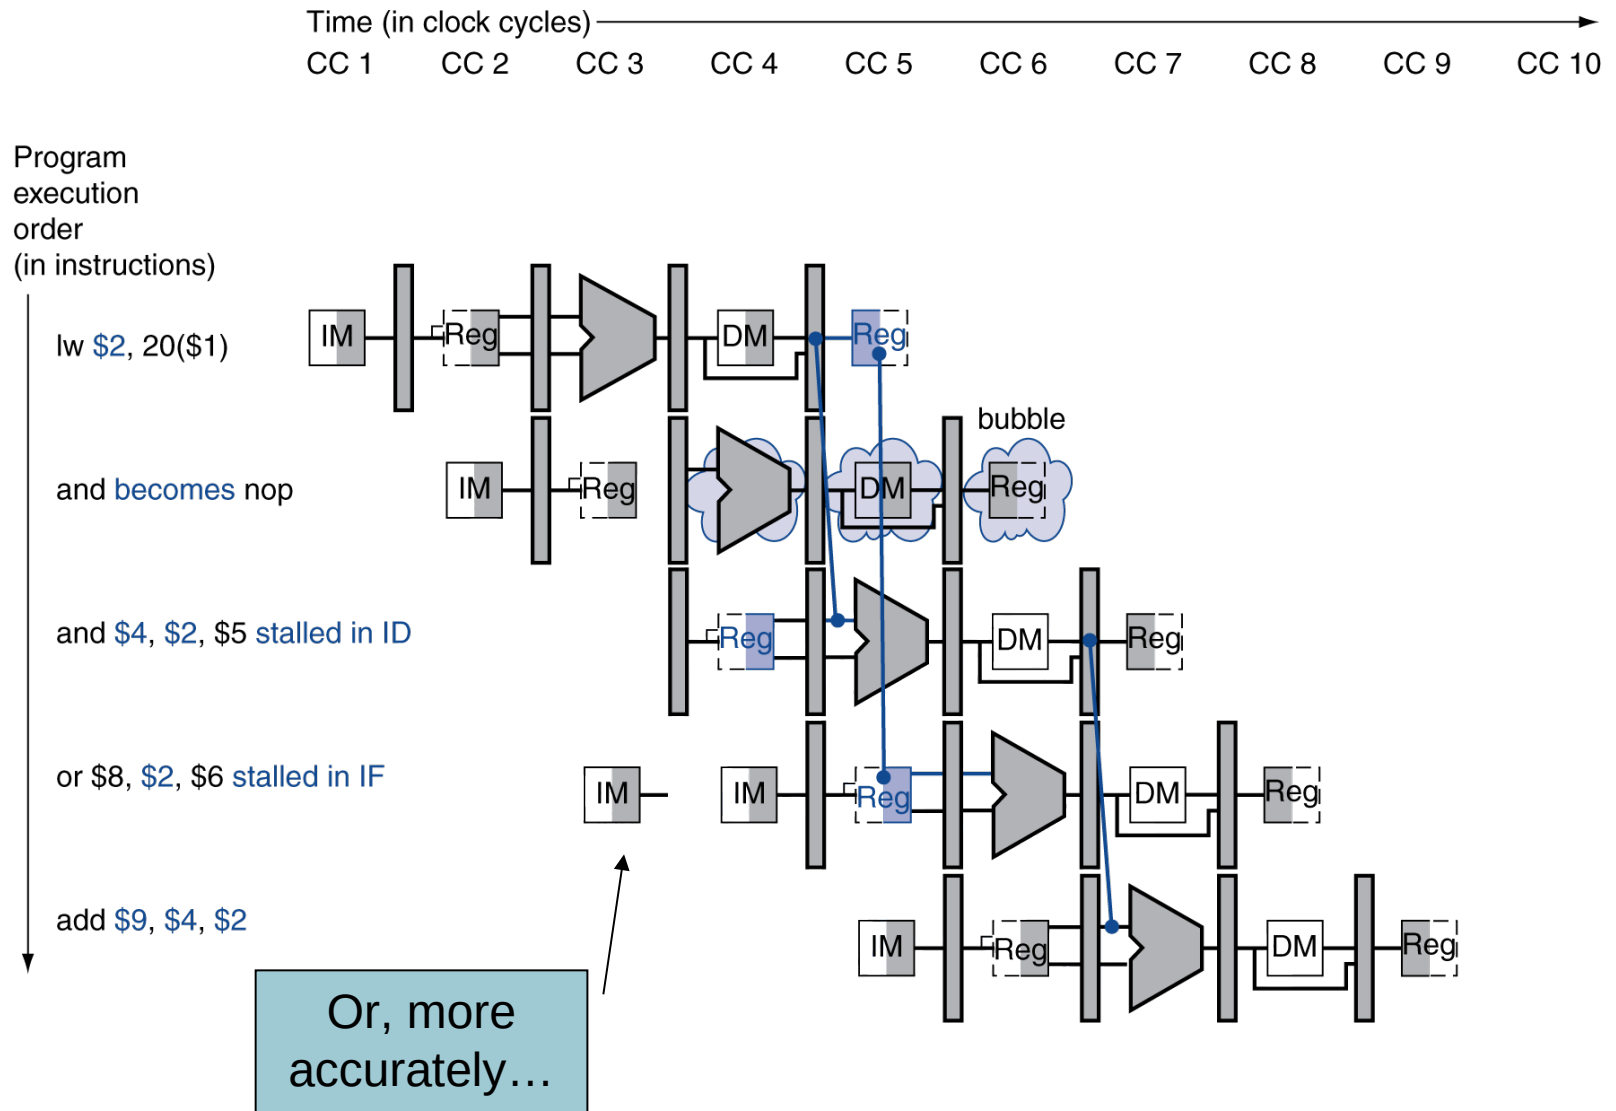
# How to Stall the Pipeline?

- Force **control values** in **ID/EX** register to **0**  
→ EX, MEM and WB do **NOOP** (no-operation)
- **Prevent update** of **PC** and **IF/ID** register
  - Current instruction is decoded again
  - Following instruction is fetched again
  - 1-cycle stall allows MEM to read data for  $1_w$ 
    - Can subsequently forward to EX stage

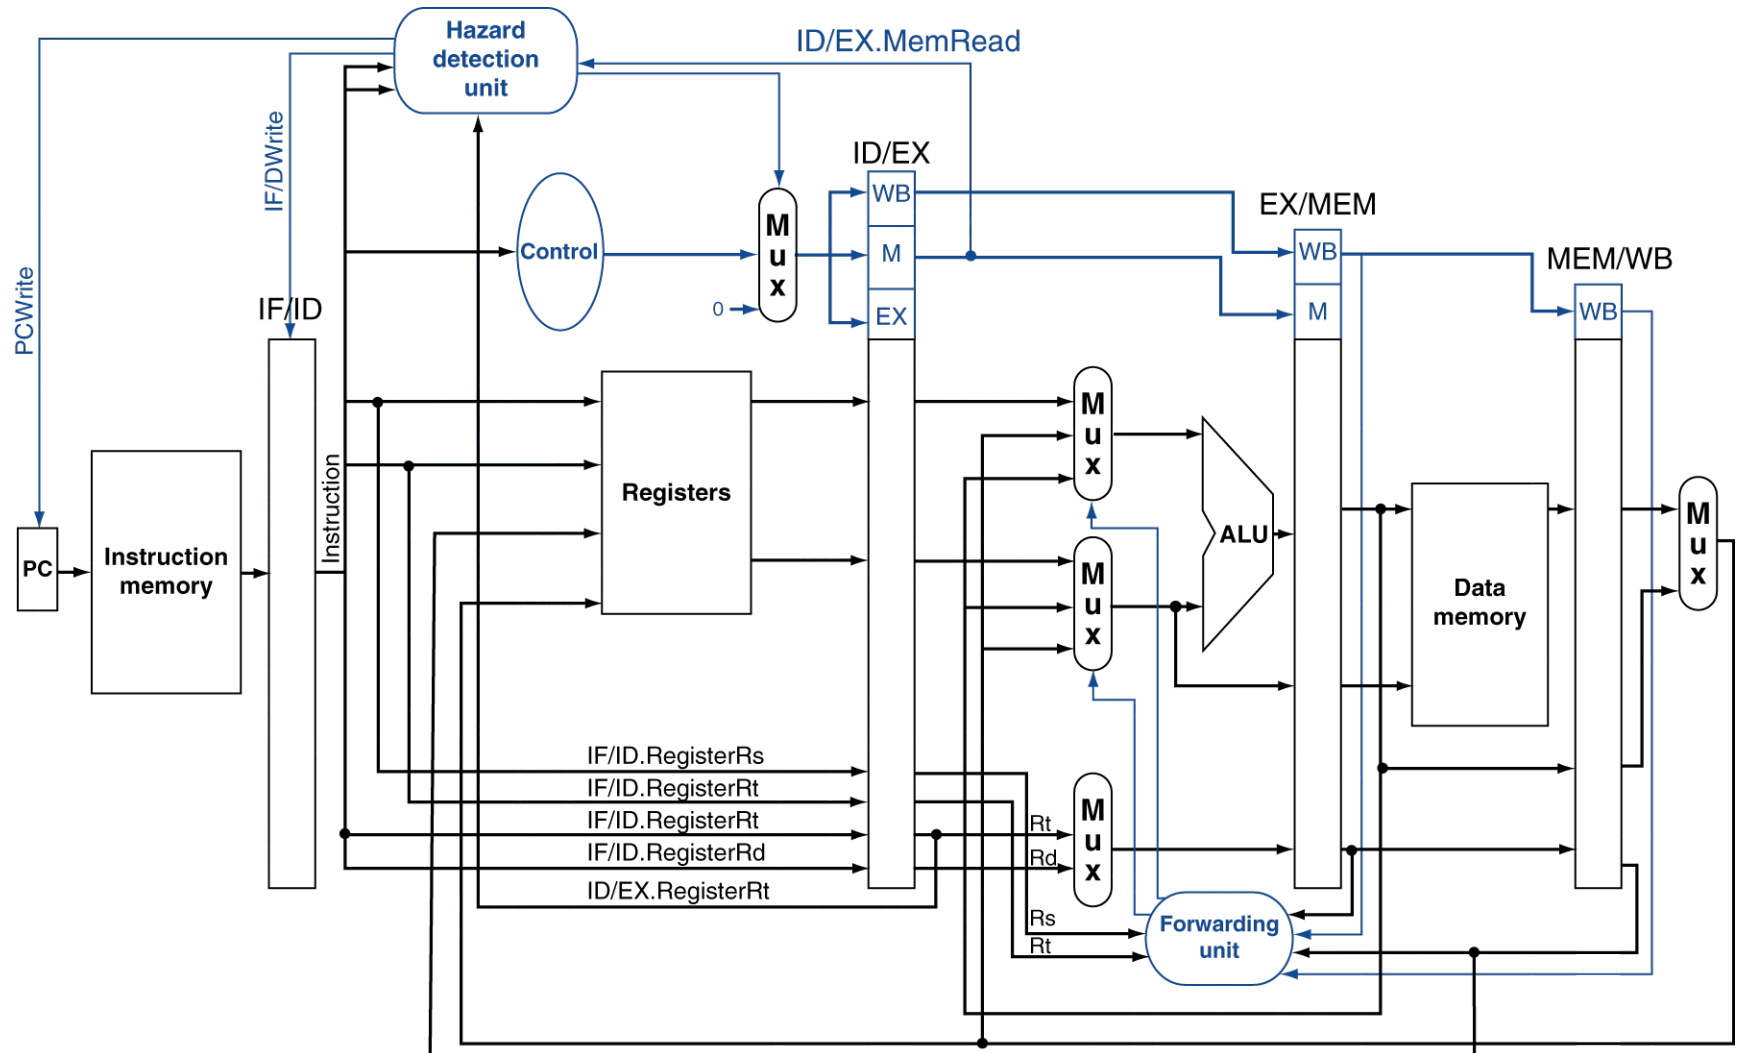
# Stall/Bubble in the Pipeline



# Stall/Bubble in the Pipeline



# Datapath with Hazard Detection





# Stalls and Performance

## The BIG Picture

- **Stalls reduce performance**
  - But are required to get **correct results**
- **Compiler** can **arrange** code to **avoid hazards and stalls**
  - This requires knowledge of the pipeline structure

# Exceptions and Interrupts

- “Unexpected” events requiring change in flow of control
  - Different ISAs use the terms differently
- Exception
  - Arises **within** the CPU
    - e.g., undefined opcode, overflow, syscall, ...
- Interrupt
  - From an **external** I/O controller
- Dealing with them without sacrificing performance is hard



# Pipeline with Exceptions

